

[!\[\]\(919a2cb85b99741a73c0c31a427236a8_img.jpg\) Open in Colab](#)

Part 1: Fundamentals & Classical Statistical Methods

Time Series Analysis in Python

Setup and Imports

Before we begin, install the required packages:

```
pip install pandas numpy matplotlib statsmodels scipy scikit-learn
```

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from statsmodels.tsa.seasonal import seasonal_decompose
5 from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
6 from statsmodels.tsa.stattools import adfuller, kpss
7 from statsmodels.tsa.arima.model import ARIMA
8 from statsmodels.tsa.statespace.sarimax import SARIMAX
9 from statsmodels.stats.diagnostic import acorr_ljungbox
10 from scipy import stats
11 from sklearn.metrics import mean_absolute_error, mean_squared_error
12 import warnings
13 warnings.filterwarnings('ignore')
14
15 plt.style.use('seaborn-v0_8-darkgrid')
```

1. What Is a Time Series?

A **time series** is a sequence of observations collected **over time**, usually at regular intervals.

Examples:

- Daily stock prices
- Hourly electricity consumption
- Monthly sales revenue
- Sensor measurements every second

Why Time Series Are Different

- Observations are **dependent**
- The order of data points **cannot be shuffled**
- Past values influence future values

This dependency structure is the core challenge of time series analysis.

Key Components

Time series typically contain four components:

1. **Trend (T)**: Long-term increase or decrease in the data
2. **Seasonality (S)**: Regular, predictable patterns that repeat over fixed periods (e.g., yearly, monthly)
3. **Cyclical (C)**: Patterns that repeat but not at fixed intervals (e.g., economic cycles)
4. **Noise/Irregular (I)**: Random variation that cannot be attributed to trend, seasonality, or cyclical

Mathematical Representation

- **Additive Model**: $Y(t) = T(t) + S(t) + C(t) + I(t)$
 - Use when seasonal variation is roughly constant over time
- **Multiplicative Model**: $Y(t) = T(t) \times S(t) \times C(t) \times I(t)$
 - Use when seasonal variation increases with the level of the series

Creating a Sample Time Series

Let's create a synthetic time series that contains trend, seasonality, and noise:

```
1 # Set random seed for reproducibility
2 np.random.seed(42)
3
4 # Generate date range
5 date_range = pd.date_range(start='2020-01-01', end='2023-12-31', freq='D')
6 n = len(date_range)
7
8 # Components
9 trend = np.linspace(100, 150, n) # Linear trend from 100 to 150
10 seasonality = 10 * np.sin(2 * np.pi * np.arange(n) / 365.25) # Yearly se
```

```
11 noise = np.random.normal(0, 5, n) # Random noise
12
13 # Combine components (additive model)
14 ts_data = trend + seasonality + noise
15
16 # Create pandas Series with datetime index
17 ts = pd.Series(ts_data, index=date_range, name='Value')
18
19 print("Sample Time Series:")
20 print(ts.head(10))
21 print(f"\nShape: {ts.shape}")
22 print(f"Period: {ts.index.min()} to {ts.index.max()}")
```

Sample Time Series:

```
2020-01-01    102.483571
2020-01-02    99.514941
2020-01-03    103.650916
2020-01-04    108.233733
2020-01-05    99.653774
2020-01-06    99.859609
2020-01-07    109.131857
2020-01-08    105.278161
2020-01-09    99.298455
2020-01-10    104.563060
```

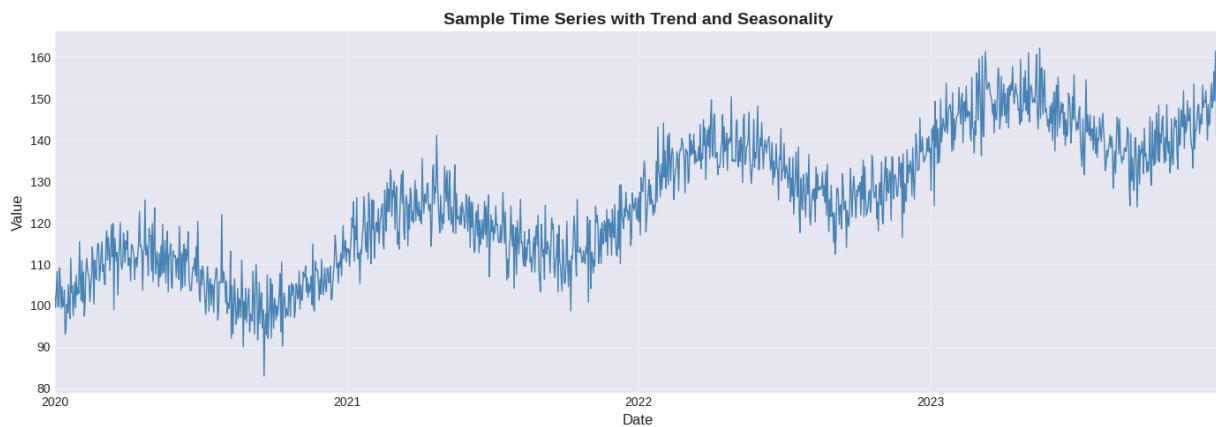
Freq: D, Name: Value, dtype: float64

Shape: (1461,)

Period: 2020-01-01 00:00:00 to 2023-12-31 00:00:00

▼ Visualizing the Time Series

```
1 fig, ax = plt.subplots(figsize=(14, 5))
2 ts.plot(ax=ax, linewidth=1, color='steelblue')
3 ax.set_title('Sample Time Series with Trend and Seasonality', fontsize=14,
4 ax.set_xlabel('Date', fontsize=12)
5 ax.set_ylabel('Value', fontsize=12)
6 ax.grid(True, alpha=0.3)
7 plt.tight_layout()
8 plt.show()
```



2. Time Series Decomposition

▼ Theory

Decomposition is the process of separating a time series into its constituent components. This helps us:

- Understand the underlying patterns
- Remove seasonality for better modeling
- Identify anomalies
- Choose appropriate forecasting methods

Classical Decomposition Methods

The most common method is **moving average decomposition**:

1. Estimate trend using moving average
2. Detrend the series
3. Estimate seasonal component by averaging detrended values for each season

4. Calculate residuals: Residual = Original - Trend - Seasonal

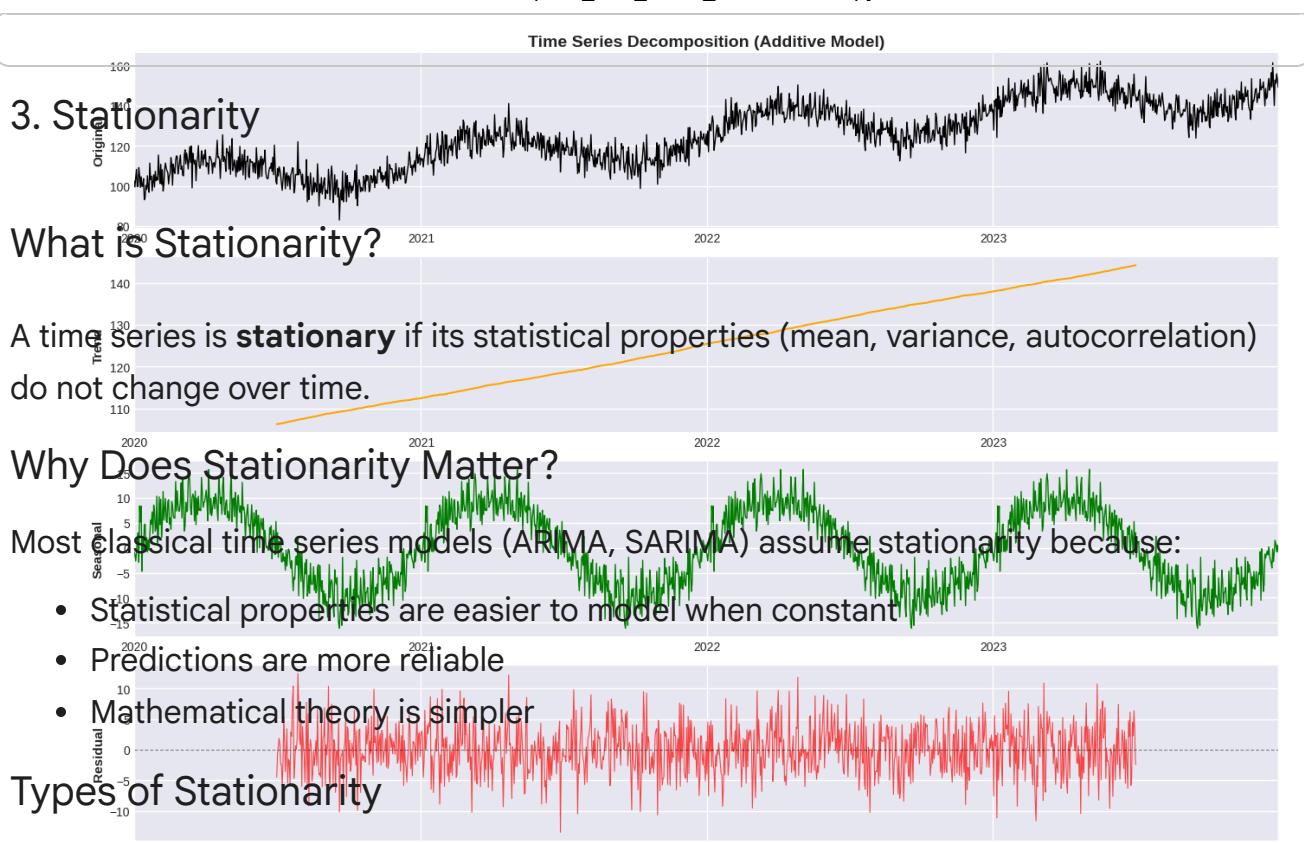
Performing Decomposition

```
1 # Perform additive decomposition
2 # Period = 365 because we have daily data with yearly seasonality
3 decomposition = seasonal_decompose(ts, model='additive', period=365)
4
5 # Extract components
6 trend_component = decomposition.trend
7 seasonal_component = decomposition.seasonal
8 residual_component = decomposition.resid
9
10 print("Decomposition Components:")
11 print(f"Trend: {trend_component.shape}")
12 print(f"Seasonal: {seasonal_component.shape}")
13 print(f"Residual: {residual_component.shape}")
```

```
Decomposition Components:
Trend: (1461,)
Seasonal: (1461,)
Residual: (1461,)
```

▼ Visualizing Decomposition

```
1 fig, axes = plt.subplots(4, 1, figsize=(14, 10))
2
3 # Original
4 ts.plot(ax=axes[0], linewidth=1, color='black')
5 axes[0].set_ylabel('Original', fontsize=11, fontweight='bold')
6 axes[0].set_title('Time Series Decomposition (Additive Model)', fontsize=14)
7
8 # Trend
9 trend_component.plot(ax=axes[1], linewidth=1.5, color='orange')
10 axes[1].set_ylabel('Trend', fontsize=11, fontweight='bold')
11
12 # Seasonal
13 seasonal_component.plot(ax=axes[2], linewidth=1, color='green')
14 axes[2].set_ylabel('Seasonal', fontsize=11, fontweight='bold')
15
16 # Residual
17 residual_component.plot(ax=axes[3], linewidth=0.8, color='red', alpha=0.7)
18 axes[3].set_ylabel('Residual', fontsize=11, fontweight='bold')
19 axes[3].set_xlabel('Date', fontsize=12)
20 axes[3].axhline(y=0, color='black', linestyle='--', linewidth=0.8, alpha=0.7)
21
22 plt.tight_layout()
23 plt.show()
```

1. **Strict Stationarity:** Joint distribution is time-invariant (very restrictive)
2. **Weak/Covariance Stationarity:** Only mean, variance, and autocorrelation are constant (commonly used)

Common Non-Stationary Patterns

- **Trend:** Mean changes over time
- **Seasonality:** Pattern repeats at regular intervals
- **Heteroscedasticity:** Variance changes over time

Augmented Dickey-Fuller (ADF) Test

The ADF test checks the null hypothesis that the series has a unit root (non-stationary).

- H_0 : Series has a unit root (non-stationary)
- H_1 : Series is stationary

Decision Rule: If $p\text{-value} < 0.05$, reject H_0 (series is stationary)

```

1 def check_stationarity(timeseries, name='Series'):
2     """
3     Perform Augmented Dickey-Fuller test
4     """
5     print(f"\n{'='*60}")
6     print(f"Stationarity Test: {name}")

```

```

7     print('*60)
8
9     # Remove NaN values
10    ts_clean = timeseries.dropna()
11
12    # Perform ADF test
13    result = adfuller(ts_clean, autolag='AIC')
14
15    print(f'ADF Statistic:      {result[0]:.6f}')
16    print(f'p-value:            {result[1]:.6f}')
17    print(f'# Lags Used:        {result[2]}')
18    print(f'# Observations:     {result[3]}')
19    print('\nCritical Values:')
20    for key, value in result[4].items():
21        print(f'  {key}: {value:.3f}')
22
23    # Interpretation
24    print('\n' + '-'*60)
25    if result[1] <= 0.05:
26        print(f"✓ STATIONARY (p={result[1]:.4f} < 0.05)")
27        print(" → Reject null hypothesis")
28    else:
29        print(f"✗ NON-STATIONARY (p={result[1]:.4f} > 0.05)")
30        print(" → Fail to reject null hypothesis")
31    print('-'*60)
32

```

▼ Testing Our Series

```

1 # Test original series
2 is_stationary = check_stationarity(ts, 'Original Time Series')

```

```

=====
Stationarity Test: Original Time Series
=====
ADF Statistic:      -0.835398
p-value:            0.808502
# Lags Used:        12
# Observations:     1448

Critical Values:
 1%: -3.435
 5%: -2.864
 10%: -2.568

-----
✗ NON-STATIONARY (p=0.8085 > 0.05)
→ Fail to reject null hypothesis
-----
```

⌄ Making a Series Stationary

Method 1: Differencing

Differencing removes trends and can stabilize the mean:

First Difference: $Y'(t) = Y(t) - Y(t-1)$

```
1 # Apply first differencing
2 ts_diff = ts.diff().dropna()
3
4 # Test differenced series
5 is_stationary_diff = check_stationarity(ts_diff, 'First Differenced Series')
```

```
=====
Stationarity Test: First Differenced Series
=====
ADF Statistic:      -17.823692
p-value:            0.000000
# Lags Used:       11
# Observations:    1448

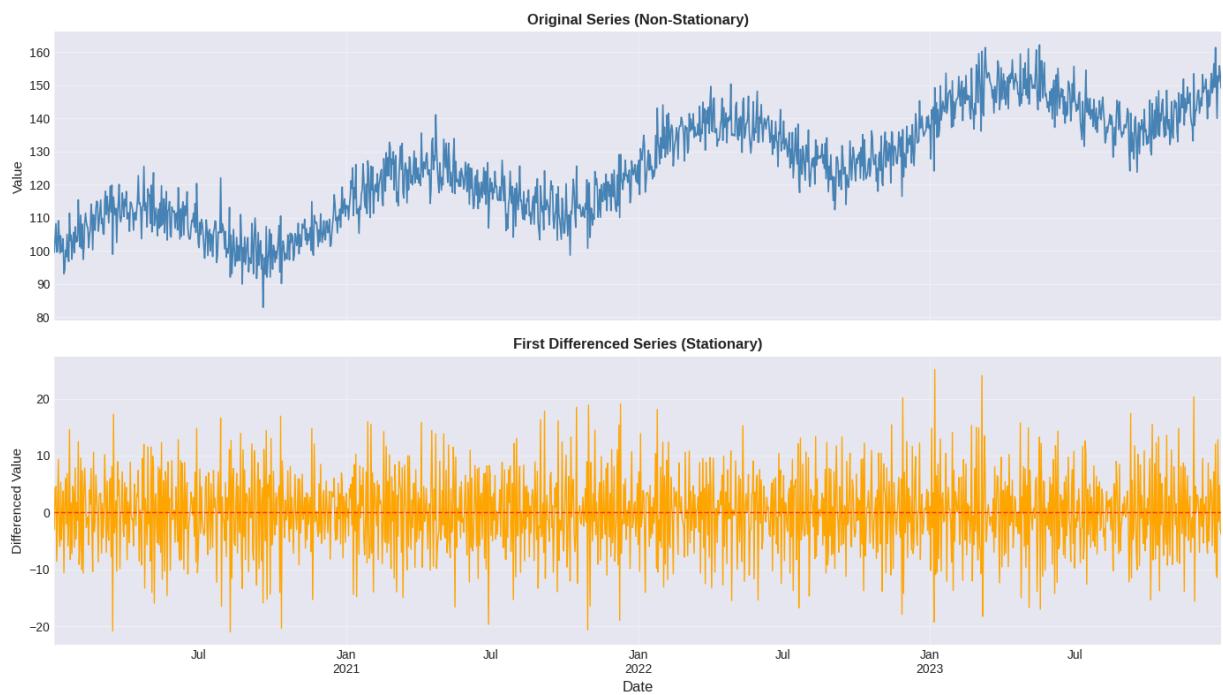
Critical Values:
1%: -3.435
5%: -2.864
10%: -2.568

-----
✓ STATIONARY (p=0.0000 < 0.05)
→ Reject null hypothesis
-----
```

⌄ Visualizing the Transformation

```
1 fig, axes = plt.subplots(2, 1, figsize=(14, 8), sharex=True)
2
3 # Original
4 ts.plot(ax=axes[0], linewidth=1.2, color='steelblue')
5 axes[0].set_title('Original Series (Non-Stationary)', fontsize=12, fontweight='bold')
6 axes[0].set_ylabel('Value', fontsize=11)
7 axes[0].grid(True, alpha=0.3)
8
9 # Differenced
10 ts_diff.plot(ax=axes[1], linewidth=1, color='orange')
11 axes[1].set_title('First Differenced Series (Stationary)', fontsize=12, fontweight='bold')
12 axes[1].set_ylabel('Differenced Value', fontsize=11)
13 axes[1].axhline(y=0, color='red', linestyle='--', linewidth=1, alpha=0.7)
14 axes[1].set_xlabel('Date', fontsize=12)
```

```
15 axes[1].grid(True, alpha=0.3)
16
17 plt.tight_layout()
18 plt.show()
```



4. Autocorrelation Analysis

▼ ACF: Autocorrelation Function

The ACF measures the correlation between observations at different time lags:

$$\text{ACF}(k) = \text{Corr}(Y_t, Y_{t-k})$$

- Lag 1: correlation between consecutive observations
- Lag 2: correlation between observations 2 time periods apart
- And so on...

PACF: Partial Autocorrelation Function

The PACF measures the correlation between Y_t and Y_{t-k} after removing the effect of intermediate lags.

Why Are ACF and PACF Important?

They help identify the order of AR and MA components:

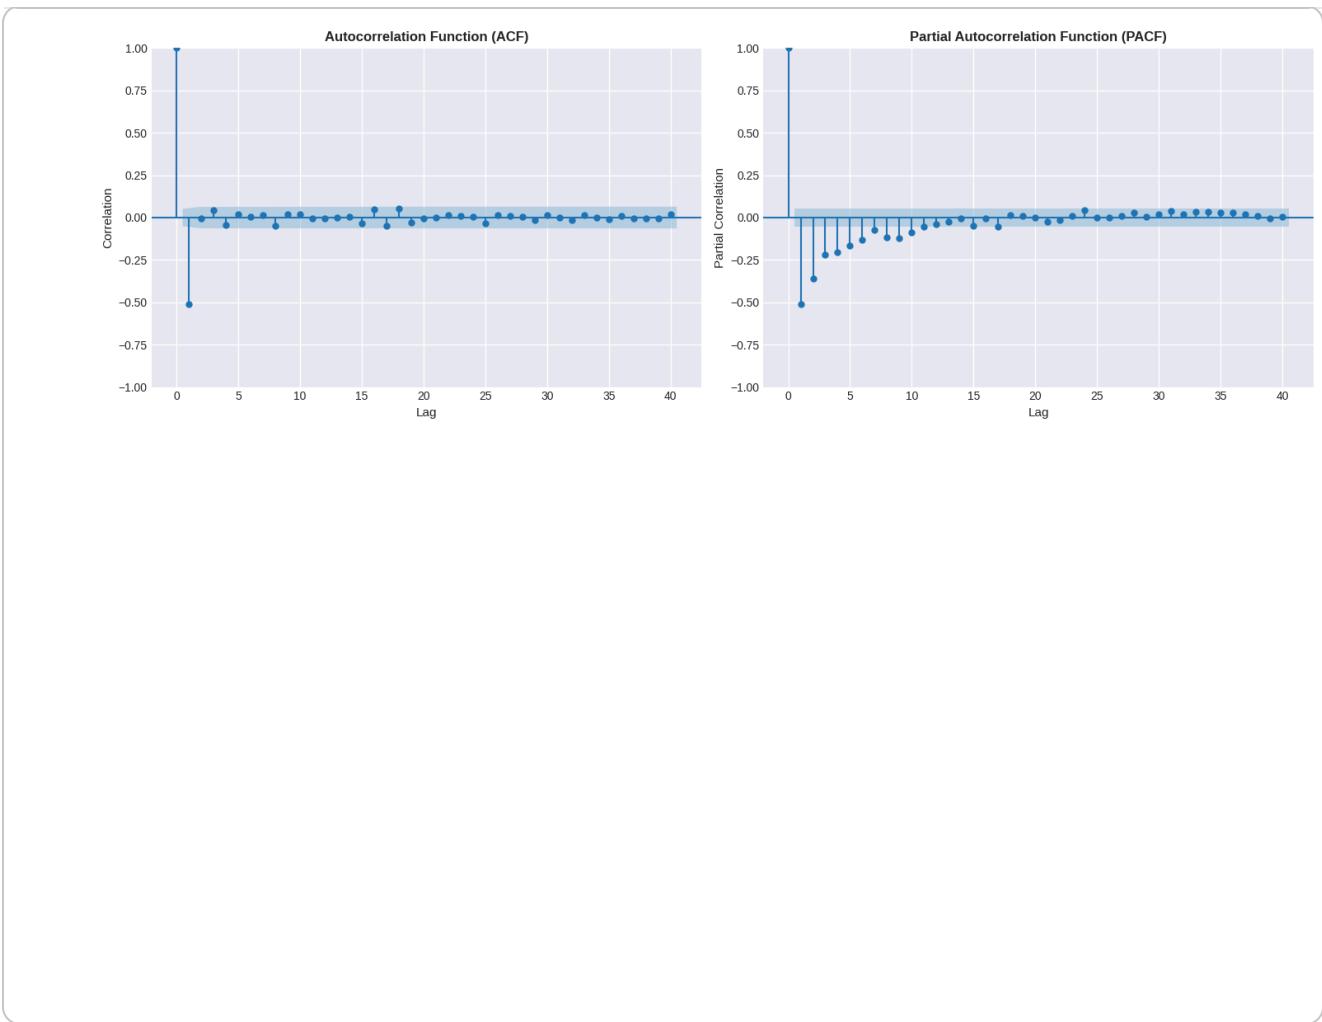
Pattern	Model Suggestion
PACF cuts off after lag p, ACF decays	AR(p)
ACF cuts off after lag q, PACF decays	MA(q)
Both decay gradually	ARMA(p,q)

Plotting ACF and PACF

```

1 # Use differenced series (should be stationary)
2 fig, axes = plt.subplots(1, 2, figsize=(15, 5))
3
4 # ACF
5 plot_acf(ts_diff, lags=40, ax=axes[0])
6 axes[0].set_title('Autocorrelation Function (ACF)', fontsize=12, fontweight='bold')
7 axes[0].set_xlabel('Lag', fontsize=11)
8 axes[0].set_ylabel('Correlation', fontsize=11)
9
10 # PACF
11 plot_pacf(ts_diff, lags=40, ax=axes[1])
12 axes[1].set_title('Partial Autocorrelation Function (PACF)', fontsize=12, fontweight='bold')
13 axes[1].set_xlabel('Lag', fontsize=11)
14 axes[1].set_ylabel('Partial Correlation', fontsize=11)
15
16 plt.tight_layout()
17 plt.show()

```



Interpretation:

- Blue shaded area represents the confidence interval
- Bars outside this area are statistically significant
- Look for where bars drop inside the confidence interval (cutoff point)

5. Classical Time Series Models

Model	Full Name	Equation	When to Use
AR(p)	Autoregressive	$Y_t = c + \phi_1 Y_{t-1} + \dots + \phi_p Y_{t-p} + \varepsilon_t$	PACF cuts off
MA(q)	Moving Average	$Y_t = \mu + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \dots + \theta_q \varepsilon_{t-q}$	ACF cuts off
ARMA(p,q)	AR + MA	Combines both	Both ACF & PACF decay
ARIMA(p,d,q)	Integrated ARMA	ARMA on d-times differenced data	Non-stationary series
SARIMA(p,d,q)(P,D,Q,s)	Seasonal ARIMA	Adds seasonal components	Seasonal patterns

Preparing Data for Modeling

Let's create a simpler dataset to demonstrate the models:

```

1 # Generate AR(1) process
2 np.random.seed(123)
3 n = 300
4 dates = pd.date_range(start='2022-01-01', periods=n, freq='D')
5
6 # AR(1): Y(t) = 0.7 * Y(t-1) + noise
7 ar_coef = 0.7
8 ar_data = [0]
9 for i in range(1, n):
10     ar_data.append(ar_coef * ar_data[i-1] + np.random.normal(0, 1))
11
12 ts_simple = pd.Series(ar_data, index=dates, name='Value')
13
14 # Train-test split (80-20)
15 train_size = int(len(ts_simple) * 0.8)
16 train = ts_simple[:train_size]
17 test = ts_simple[train_size:]
18
19 print(f"Train: {len(train)} observations")
20 print(f"Test: {len(test)} observations")

```

Train: 240 observations
 Test: 60 observations

Model 1: AR (Autoregressive)

Theory

An AR(p) model predicts the current value using p past values:

$$Y_t = c + \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \dots + \phi_p Y_{t-p} + \varepsilon_t$$

Where:

- c is a constant
- $\phi_1, \phi_2, \dots, \phi_p$ are coefficients
- ε_t is white noise error

Implementation

```

1 # Fit AR(1) model
2 # ARIMA(p,d,q) with p=1, d=0, q=0
3 ar_model = ARIMA(train, order=(1, 0, 0))
4 ar_fit = ar_model.fit()
5
6 print("AR(1) Model Summary:")
7 print(ar_fit.summary())

```

```
& print(f"\nEstimated coefficient: {ar[1]}")
```

AR(1) Model Summary:

SARIMAX Results

Dep. Variable:	Value	No. Observations:	240
Model:	ARIMA(1, 0, 0)	Log Likelihood	-345.782
Date:	Sat, 27 Dec 2025	AIC	697.563
Time:	07:05:21	BIC	708.005
Sample:	01-01-2022 - 08-28-2022	HQIC	701.771
Covariance Type:	opg		

	coef	std err	z	P> z	[0.025	0.975]
const	-0.0308	0.229	-0.135	0.893	-0.479	0.418
ar.L1	0.7119	0.047	15.074	0.000	0.619	0.805
sigma2	1.0415	0.094	11.072	0.000	0.857	1.226

Ljung-Box (L1) (Q):	0.15	Jarque-Bera (JB):	0
Prob(Q):	0.70	Prob(JB):	0
Heteroskedasticity (H):	0.60	Skew:	-0
Prob(H) (two-sided):	0.03	Kurtosis:	3

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-s

Estimated coefficient: 0.7119

True coefficient: 0.7

Interpreting AR(1) Model Output

What this really is (in one sentence)

AR(1) is just linear regression where the predictor is yesterday's value.

Connect it to regular regression (intuition first)

In ordinary regression you write:

$$y = \beta_0 + \beta_1 x_1 + \text{error}$$

In AR(1), you write:

$$y_t = \text{constant} + \beta_1 \times y_{t-1} + \text{error}$$

Same idea. The "feature" is just the lagged value of y itself.

Translate the output into regression language

1. The coefficients

```
const      = -0.0308
ar.L1      =  0.7119
sigma2     =  1.0415
```

Think of this as a regression table:

Feature	Coefficient	What it means
Intercept	-0.031	Baseline (when $y_{t-1} = 0$)
y_{t-1}	0.7119	Today = 71% of yesterday + noise
Error variance	1.042	Residual spread

The prediction equation is:

$$y_t = -0.031 + 0.7119 \times y_{t-1} + \varepsilon_t$$

This is identical to:

$$\text{sales_today} = -0.031 + 0.7119 \times \text{sales_yesterday} + \text{noise}$$

2. Standard errors and p-values

Exactly the same as regression:

	coef	std err	p-value
const	-0.031	0.229	0.893
ar.L1	0.712	0.047	0.000

const: $p = 0.893$ (high)

Not significant. Could drop it.

ar.L1: $p = 0.000$ (very low)

Highly significant. Yesterday strongly predicts today.

Same interpretation as:

feature_1: $p = 0.893 \rightarrow$ not useful feature_2: $p = 0.000 \rightarrow$ very useful

The model found **strong autocorrelation** in the data.

3. Model quality metrics

AIC	= 697.6
BIC	= 708.0
Log Likelihood	= -345.8

These are **goodness-of-fit measures**, not accuracy metrics.

AIC and BIC: Lower is better

Use these to compare models:

- AR(1): AIC = 697.6
- AR(2): AIC = 695.3 → better
- MA(1): AIC = 701.2 → worse

Think of AIC like:

penalized training error

It balances fit with model complexity.

Log Likelihood: Higher is better

Similar to minimizing loss in ML. Maximum likelihood estimation finds coefficients that make the observed data most probable.

4. Diagnostic tests (are residuals well-behaved?)

Ljung-Box (Q): 0.15 p = 0.70

Jarque-Bera (JB): 0.10 p = 0.95

Heteroskedasticity (H): 0.60 p = 0.03

Think of these as **residual plots in test form**.

Ljung-Box test

Question: "Do residuals have patterns left?"

- $p > 0.05 \rightarrow$ No patterns (good)
- $p < 0.05 \rightarrow$ Patterns remain (bad)

Here: $p = 0.70 \rightarrow$ residuals look random

This is like checking:

plot(residuals) shows no trend

Jarque-Bera test

Question: "Are residuals normally distributed?"

- $p > 0.05 \rightarrow$ Yes (good)
- $p < 0.05 \rightarrow$ No (may need transformation)

Here: $p = 0.95 \rightarrow$ very normal

This is like checking:

histogram(residuals) looks bell-shaped

Heteroskedasticity test

Question: "Is variance constant over time?"

- $p > 0.05 \rightarrow$ Yes (good)
- $p < 0.05 \rightarrow$ No (variance changes)

Here: $p = 0.03 \rightarrow$ some heteroskedasticity detected

This is like seeing:

residual spread increases over time

5. Model validation

Estimated coefficient: 0.7119

True coefficient: 0.7000

The model recovered the true data-generating process.

In ML terms:

The model learned the correct function

This would be like:

- You generate data: $y = 2x + \text{noise}$
- Your model learns: $y = 1.98x$
- Close match \rightarrow model works

Why it feels alien

1. No feature matrix visible

The "feature" (y_{t-1}) is created internally from the time series.

2. Maximum likelihood instead of MSE

Same goal (fit the data), different math. Likelihood is more general than squared error.

3. Heavy focus on diagnostics

Econometrics cares deeply about **why** the model works, not just **that** it works.

Statistical inference > predictive accuracy

4. Different vocabulary

- ML says: " $R^2 = 0.85$ "

- Econometrics says: "Log Likelihood = -345, AIC = 697"

Same information, different packaging.

How to read this output (decision rules)

Check coefficients:

- $p < 0.05 \rightarrow$ significant \rightarrow keep
- $p > 0.05 \rightarrow$ not significant \rightarrow consider dropping

Check diagnostics:

- Ljung-Box $p > 0.05 \rightarrow$ good (no autocorrelation left)
- Ljung-Box $p < 0.05 \rightarrow$ bad (model incomplete)

Compare models:

- Lower AIC \rightarrow better model
- Lower BIC \rightarrow better model (penalizes complexity more)

Validate:

- Residuals should look random
- Residuals should be roughly normal
- Variance should be constant (ideally)

One grounding sentence you can remember

AR(1) is linear regression where X is yesterday's y, and the output tells you both fit quality and residual behavior in statistical language instead of ML metrics.

Quick interpretation of this specific output

Model found: $y_t = 0.71 \times y_{t-1} + \text{noise}$

Constant term: Not significant ($p = 0.89$) \rightarrow probably zero in reality

Autocorrelation: Strong (coefficient = 0.71, $p < 0.001$) \rightarrow yesterday matters a lot

Residuals: Clean (Ljung-Box $p = 0.70$, JB $p = 0.95$) \rightarrow no patterns left

Minor issue: Slight heteroskedasticity ($p = 0.03$) \rightarrow variance not perfectly constant

Conclusion: This is a good model. The AR(1) structure correctly captures the data-generating process.

▼ Model 2: MA (Moving Average)

Theory

An MA(q) model predicts the current value using past forecast errors:

$$Y_t = \mu + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \dots + \theta_q \varepsilon_{t-q}$$

Where:

- μ is the mean
- $\theta_1, \theta_2, \dots, \theta_q$ are coefficients
- $\varepsilon_t, \varepsilon_{t-1}, \dots$ are error terms

Implementation

```
1 # Fit MA(1) model
2 ma_model = ARIMA(train, order=(0, 0, 1))
3 ma_fit = ma_model.fit()
4
5 print("MA(1) Model Summary:")
6 print(ma_fit.summary())
```

MA(1) Model Summary:

SARIMAX Results

```
=====
Dep. Variable:                      Value    No. Observations:                 240
Model:                            ARIMA(0, 0, 1)    Log Likelihood:            -370.431
Date:                            Sat, 27 Dec 2025    AIC:                         746.863
Time:                            07:05:22        BIC:                         757.305
Sample:                           01-01-2022    HQIC:                        751.070
                                  - 08-28-2022
Covariance Type:                  opg
=====
              coef    std err        z     P>|z|      [0.025      0.975]
-----
const    -0.0349     0.120   -0.290      0.771     -0.270     0.200
ma.L1      0.6287     0.051   12.246      0.000      0.528     0.729
sigma2     1.2801     0.119   10.766      0.000      1.047     1.513
=====
Ljung-Box (L1) (Q):                10.99    Jarque-Bera (JB):           0
Prob(Q):                           0.00    Prob(JB):                   0
Heteroskedasticity (H):             0.60    Skew:                       -0
Prob(H) (two-sided):                0.03    Kurtosis:                   2
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-s

Model 3: ARMA

Theory

ARMA(p,q) combines AR(p) and MA(q):

$$Y_t = c + \phi_1 Y_{t-1} + \dots + \phi_p Y_{t-p} + \theta_1 \epsilon_{t-1} + \dots + \theta_q \epsilon_{t-q} + \epsilon_t$$

Important: ARMA requires stationary data.

Implementation

```
1 # Fit ARMA(1,1) model
2 arma_model = ARIMA(train, order=(1, 0, 1))
3 arma_fit = arma_model.fit()
4
5 print("ARMA(1,1) Model Summary:")
6 print(arma_fit.summary())
```

ARMA(1,1) Model Summary:

SARIMAX Results

```
=====
Dep. Variable:                      Value    No. Observations:                 240
Model:                            ARIMA(1, 0, 1)    Log Likelihood:            -345.671
Date:                Sat, 27 Dec 2025      AIC:                         699.341
Time:                    07:05:23        BIC:                         713.264
Sample:               01-01-2022      HQIC:                        704.951
                           - 08-28-2022
Covariance Type:                  opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	-0.0308	0.223	-0.138	0.890	-0.468	0.407
ar.L1	0.6912	0.072	9.667	0.000	0.551	0.831
ma.L1	0.0424	0.096	0.441	0.659	-0.146	0.231
sigma2	1.0405	0.094	11.083	0.000	0.857	1.225

=====

	0.00	Jarque-Bera (JB):	0
Ljung-Box (L1) (Q):	0.96	Prob(JB):	0
Prob(Q):	0.60	Skew:	-0
Heteroskedasticity (H):	0.02	Kurtosis:	3
Prob(H) (two-sided):			

=====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-s

Model 4: ARIMA

Theory

ARIMA(p,d,q) is ARMA applied to d-times differenced data:

- **p:** Order of autoregressive part
- **d:** Degree of differencing

- q: Order of moving average part

Steps:

1. Difference the series d times to make it stationary
2. Apply ARMA(p,q) to the differenced series

Implementation

```

1 # Fit ARIMA(1,1,1)
2 arima_model = ARIMA(train, order=(1, 1, 1))
3 arima_fit = arima_model.fit()
4
5 print("ARIMA(1,1,1) Model Summary:")
6 print(arima_fit.summary())
7 print(f"\nAIC: {arima_fit.aic:.2f}")
8 print(f"BIC: {arima_fit.bic:.2f}")
9
10 # Forecast
11 forecast_steps = len(test)
12 arima_forecast = arima_fit.forecast(steps=forecast_steps)

```

ARIMA(1,1,1) Model Summary:

SARIMAX Results

```
=====
Dep. Variable:                      Value     No. Observations:                  240
Model:                            ARIMA(1, 1, 1)   Log Likelihood:           -346.335
Date:                            Sat, 27 Dec 2025   AIC:                   698.670
Time:                             07:05:25      BIC:                   709.100
Sample:                           01-01-2022   HQIC:                  702.873
                                  - 08-28-2022
Covariance Type:                    opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.7190	0.052	13.919	0.000	0.618	0.820
ma.L1	-1.0000	6.486	-0.154	0.877	-13.712	11.712
sigma2	1.0459	6.772	0.154	0.877	-12.227	14.319

```
=====
Ljung-Box (L1) (Q):                  0.08   Jarque-Bera (JB):            0
Prob(Q):                           0.78   Prob(JB):                  0
Heteroskedasticity (H):              0.59   Skew:                     -0
Prob(H) (two-sided):                0.02   Kurtosis:                  3
=====
```

Warnings:

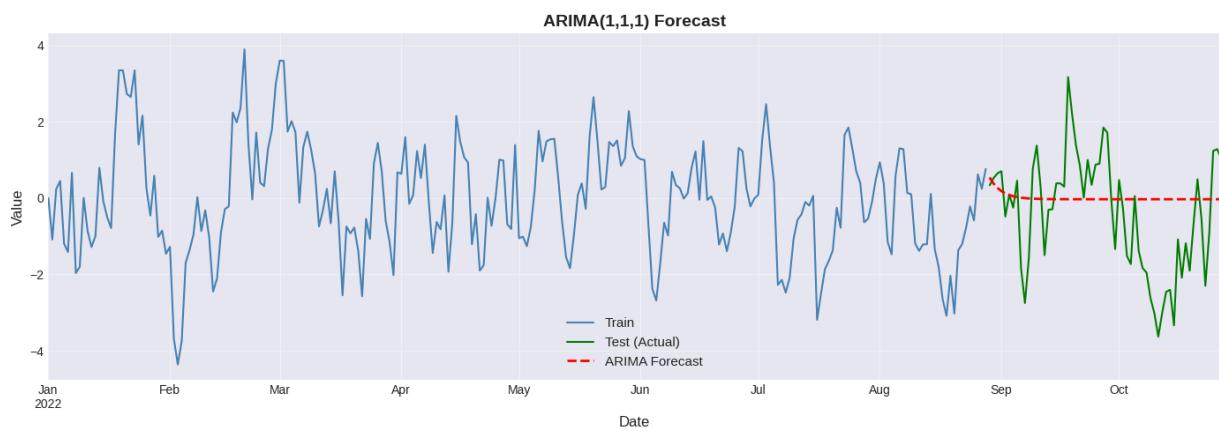
[1] Covariance matrix calculated using the outer product of gradients (complex-s

AIC: 698.67

BIC: 709.10

Visualizing Forecast

```
1 fig, ax = plt.subplots(figsize=(14, 5))
2
3 train.plot(ax=ax, label='Train', linewidth=1.5, color='steelblue')
4 test.plot(ax=ax, label='Test (Actual)', linewidth=1.5, color='green')
5 arima_forecast.plot(ax=ax, label='ARIMA Forecast', linewidth=2,
6                      linestyle='--', color='red')
7
8 ax.set_title('ARIMA(1,1,1) Forecast', fontsize=14, fontweight='bold')
9 ax.set_xlabel('Date', fontsize=12)
10 ax.set_ylabel('Value', fontsize=12)
11 ax.legend(fontsize=11)
12 ax.grid(True, alpha=0.3)
13 plt.tight_layout()
14 plt.show()
```



Model 5: SARIMA (Seasonal ARIMA)

Theory

SARIMA(p,d,q)(P,D,Q,s) extends ARIMA to handle seasonality:

Non-seasonal part: (p,d,q) **Seasonal part:** (P,D,Q,s)

- P: Seasonal AR order
- D: Seasonal differencing order
- Q: Seasonal MA order
- s: Seasonal period (e.g., 12 for monthly data with yearly seasonality)

Creating Seasonal Data

```

1 # Generate monthly data with seasonality
2 n_months = 240 # 20 years
3 dates_monthly = pd.date_range(start='2004-01-01', periods=n_months, freq='M')
4
5 # Components
6 trend = np.linspace(50, 100, n_months)
7 seasonal = 15 * np.sin(2 * np.pi * np.arange(n_months) / 12)
8 noise = np.random.normal(0, 3, n_months)
9
10 ts_seasonal = pd.Series(trend + seasonal + noise,
11                           index=dates_monthly,
12                           name='Sales')
13
14 # Split
15 train_seas = ts_seasonal[:int(len(ts_seasonal) * 0.8)]
16 test_seas = ts_seasonal[int(len(ts_seasonal) * 0.8):]
17
18 print(f"Seasonal series shape: {ts_seasonal.shape}")
19 print(f"Train: {len(train_seas)}, Test: {len(test_seas)}")

```

Seasonal series shape: (240,)
 Train: 192, Test: 48

Fit SARIMA

```

1 # Fit SARIMA(1,1,1)(1,1,1,12)
2 # Seasonal period = 12 months
3 sarima_model = SARIMAX(train_seas,
4                         order=(1, 1, 1),
5                         seasonal_order=(1, 1, 1, 12))
6 sarima_fit = sarima_model.fit(disp=False)
7
8 print("SARIMA(1,1,1)(1,1,1,12) Model Summary:")
9 print(sarima_fit.summary())
10
11 # Forecast
12 sarima_forecast = sarima_fit.forecast(steps=len(test_seas))

```

SARIMA(1,1,1)(1,1,1,12) Model Summary:

SARIMAX Results

```
=====
Dep. Variable:                      Sales      No. Observations:      930
Model:             SARIMAX(1, 1, 1)x(1, 1, 1, 12)   Log Likelihood   -112.5
Date:                Sat, 27 Dec 2025      AIC                  225.0
Time:                      07:05:31        BIC                  231.0
Sample:                 01-01-2004      HQIC                  225.0
                           - 12-01-2019
Covariance Type:            opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.0127	0.084	-0.151	0.880	-0.178	0.153
ma.L1	-0.9801	0.040	-24.629	0.000	-1.058	-0.902
ar.S.L12	-0.0239	0.104	-0.231	0.817	-0.227	0.179
ma.S.L12	-0.9156	0.140	-6.531	0.000	-1.190	-0.641
sigma2	8.9150	1.336	6.671	0.000	6.296	11.534

=====

Ljung-Box (L1) (Q):	0.03	Jarque-Bera (JB):	1
Prob(Q):	0.87	Prob(JB):	0
Heteroskedasticity (H):	1.30	Skew:	-0
Prob(H) (two-sided):	0.31	Kurtosis:	2

=====

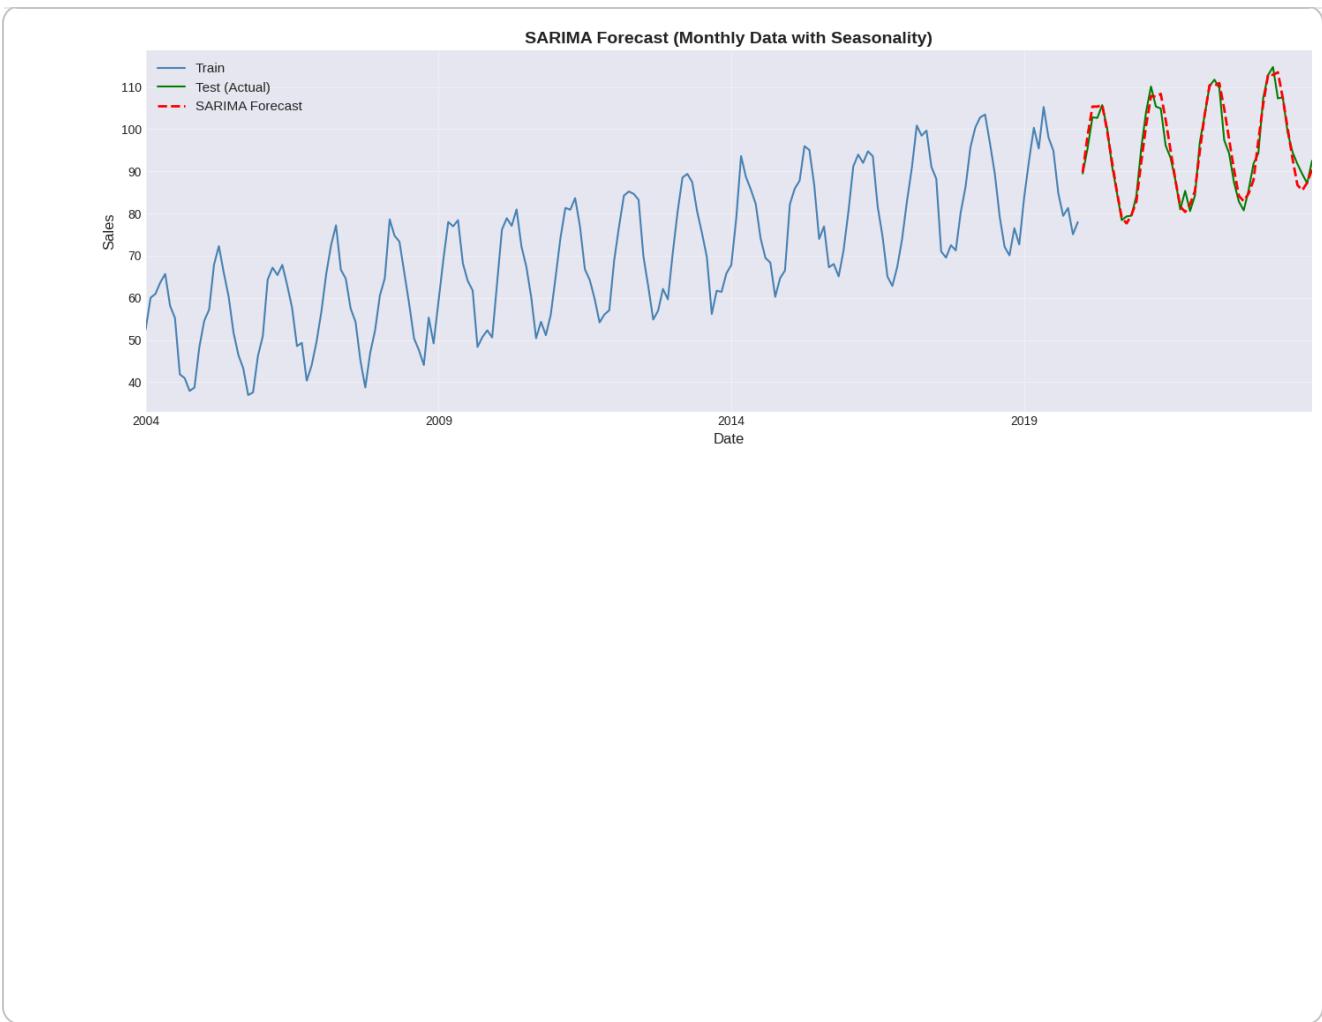
Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-s

Visualizing Seasonal Forecast

```

1 fig, ax = plt.subplots(figsize=(14, 5))
2
3 train_seas.plot(ax=ax, label='Train', linewidth=1.5, color='steelblue')
4 test_seas.plot(ax=ax, label='Test (Actual)', linewidth=1.5, color='green')
5 sarima_forecast.plot(ax=ax, label='SARIMA Forecast', linewidth=2,
6                         linestyle='--', color='red')
7
8 ax.set_title('SARIMA Forecast (Monthly Data with Seasonality)',
9               fontsize=14, fontweight='bold')
10 ax.set_xlabel('Date', fontsize=12)
11 ax.set_ylabel('Sales', fontsize=12)
12 ax.legend(fontsize=11)
13 ax.grid(True, alpha=0.3)
14 plt.tight_layout()
15 plt.show()
```



6. Model Selection

▼ Information Criteria

AIC (Akaike Information Criterion)

$$\text{AIC} = 2k - 2\ln(L)$$

Where:

- k = number of parameters
- L = maximum likelihood

BIC (Bayesian Information Criterion)

$$\text{BIC} = k \cdot \ln(n) - 2\ln(L)$$

Where:

- n = number of observations

Rule: Lower AIC/BIC = better model (balances fit and complexity)

Grid Search for Best Model

```

1 def evaluate_arima_models(data, p_range, d_range, q_range):
2     """
3     Evaluate different ARIMA configurations
4     """
5     results = []
6
7     for p in p_range:
8         for d in d_range:
9             for q in q_range:
10                try:
11                    model = ARIMA(data, order=(p, d, q))
12                    fitted = model.fit()
13
14                    results.append({
15                        'order': (p, d, q),
16                        'AIC': fitted.aic,
17                        'BIC': fitted.bic
18                    })
19                except:
20                    continue
21
22    return pd.DataFrame(results)
23
24 # Search for best model
25 print("Searching for best ARIMA model...")
26 results_df = evaluate_arima_models(
27     train,
28     p_range=range(0, 3),
29     d_range=range(0, 2),
30     q_range=range(0, 3)
31 )
32
33 # Sort and display
34 results_df = results_df.sort_values('AIC')
35 print("\nTop 5 Models by AIC:")
36 print(results_df.head())
37
38 best_order = results_df.iloc[0]['order']
39 print(f"\nBest Model: ARIMA{best_order}")

```

Searching for best ARIMA model...

Top 5 Models by AIC:

	order	AIC	BIC
6	(1, 0, 0)	697.563191	708.005108
13	(2, 0, 1)	698.028252	715.431446

```

10 (1, 1, 1) 698.670134 709.099525
12 (2, 0, 0) 699.335230 713.257786
7 (1, 0, 1) 699.341125 713.263681

```

Best Model: ARIMA(1, 0, 0)

7. Model Diagnostics

Why Check Diagnostics?

After fitting a model, we need to verify that:

1. **Residuals are white noise** (random, no pattern)
2. **Residuals are normally distributed**
3. **No autocorrelation in residuals**
4. **Model assumptions are satisfied**

Theory: Good Residuals

If the model is appropriate, residuals should have:

- Mean ≈ 0
- Constant variance (homoscedastic)
- No autocorrelation
- Normal distribution

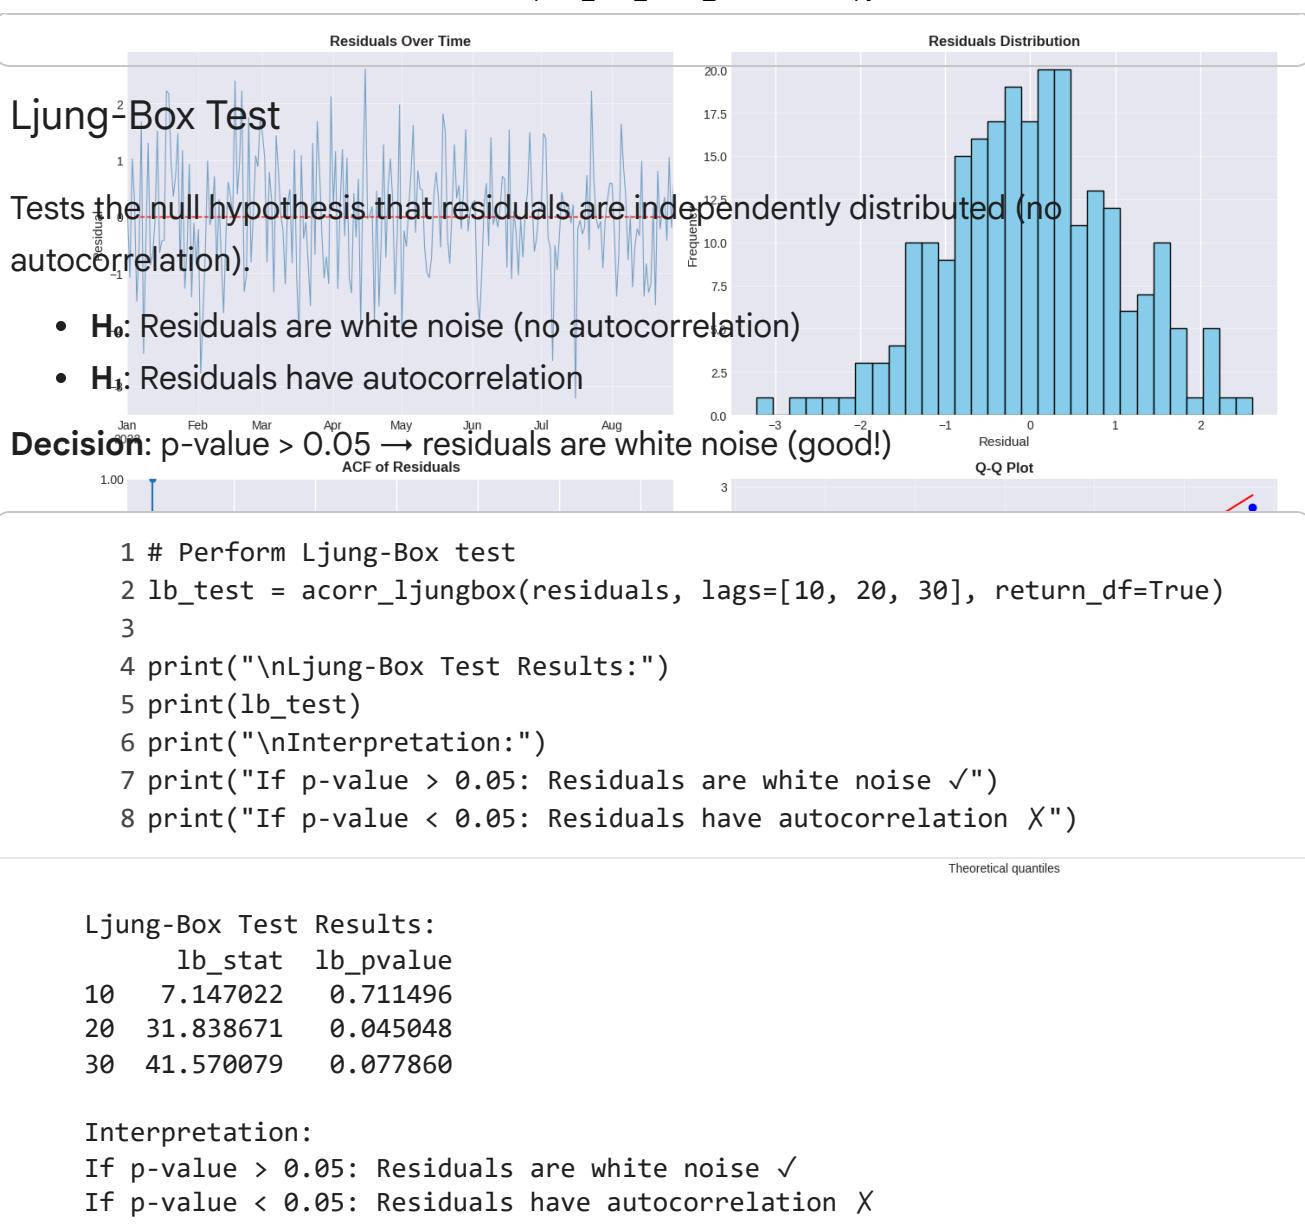
Diagnostic Plots

```

1 # Get residuals from best model
2 best_model = ARIMA(train, order=best_order).fit()
3 residuals = best_model.resid
4
5 # Create diagnostic plots
6 fig, axes = plt.subplots(2, 2, figsize=(14, 10))
7
8 # 1. Residuals over time
9 residuals.plot(ax=axes[0, 0], linewidth=0.8, color='steelblue', alpha=0.7)
10 axes[0, 0].axhline(y=0, color='red', linestyle='--', linewidth=1)
11 axes[0, 0].set_title('Residuals Over Time', fontsize=12, fontweight='bold')
12 axes[0, 0].set_ylabel('Residual', fontsize=11)
13 axes[0, 0].grid(True, alpha=0.3)
14
15 # 2. Histogram
16 residuals.hist(ax=axes[0, 1], bins=30, edgecolor='black', color='skyblue')
17 axes[0, 1].set_title('Residuals Distribution', fontsize=12, fontweight='bold')

```

```
18 axes[0, 1].set_xlabel('Residual', fontsize=11)
19 axes[0, 1].set_ylabel('Frequency', fontsize=11)
20 axes[0, 1].grid(True, alpha=0.3, axis='y')
21
22 # 3. ACF of residuals
23 plot_acf(residuals, lags=30, ax=axes[1, 0])
24 axes[1, 0].set_title('ACF of Residuals', fontsize=12, fontweight='bold')
25
26 # 4. Q-Q plot
27 stats.probplot(residuals, dist="norm", plot=axes[1, 1])
28 axes[1, 1].set_title('Q-Q Plot', fontsize=12, fontweight='bold')
29 axes[1, 1].grid(True, alpha=0.3)
30
31 plt.tight_layout()
32 ~~~~
```

Residual Statistics

```

1 print("\nResidual Summary Statistics:")
2 print("=*50)
3 print(f"Mean: {residuals.mean():>10.6f} (should be ≈ 0)")
4 print(f"Std Dev: {residuals.std():>10.6f}")
5 print(f"Min: {residuals.min():>10.6f}")
6 print(f"Max: {residuals.max():>10.6f}")
7 print(f"Skewness: {residuals.skew():>10.6f} (should be ≈ 0)")
8 print(f"Kurtosis: {residuals.kurtosis():>10.6f} (should be ≈ 0)")

```

Residual Summary Statistics:
=====

Mean:	-0.000090 (should be ≈ 0)
Std Dev:	1.022687
Min:	-3.222888

Max:	2.614799
Skewness:	-0.045184 (should be ≈ 0)
Kurtosis:	0.071627 (should be ≈ 0)

8. Model Evaluation

▼ Common Forecasting Metrics

1. MAE (Mean Absolute Error)

$$\text{MAE} = (1/n) \sum |y_i - \hat{y}_i|$$

- Average absolute difference
- Same units as original data
- Easy to interpret

2. RMSE (Root Mean Squared Error)

$$\text{RMSE} = \sqrt{(1/n) \sum (y_i - \hat{y}_i)^2}$$

- Penalizes large errors more
- Same units as original data
- More sensitive to outliers than MAE

3. MAPE (Mean Absolute Percentage Error)

$$\text{MAPE} = (100/n) \sum |(y_i - \hat{y}_i)/y_i|$$

- Expressed as percentage
- Scale-independent
- Undefined when $y_i = 0$

Computing Metrics

```

1 def calculate_metrics(actual, predicted):
2     """Calculate forecasting metrics"""
3     mae = mean_absolute_error(actual, predicted)
4     rmse = np.sqrt(mean_squared_error(actual, predicted))
5     mape = np.mean(np.abs((actual - predicted) / actual)) * 100
6
7     return {'MAE': mae, 'RMSE': rmse, 'MAPE': mape}
8
9 # Evaluate forecast
10 forecast = best_model.forecast(steps=len(test))
11 metrics = calculate_metrics(test, forecast)

```

```
12
13 print("\nForecast Evaluation Metrics:")
14 print("=="*50)
15 print(f"MAE: {metrics['MAE']:.4f}")
16 print(f"RMSE: {metrics['RMSE']:.4f}")
17 print(f"MAPE: {metrics['MAPE']:.2f}%)")
```

```
Forecast Evaluation Metrics:
=====
MAE: 1.2387
RMSE: 1.5522
MAPE: 102.35%
```

Summary and Key Takeaways