



Open in Colab

# Part 1: Fundamentals & Classical Statistical Methods

## Time Series Analysis in Python

### Setup and Imports

Before we begin, install the required packages:

```
pip install pandas numpy matplotlib statsmodels scipy scikit-learn
```

In [27]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.stattools import adfuller, kpss
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.stats.diagnostic import acorr_ljungbox
from scipy import stats
from sklearn.metrics import mean_absolute_error, mean_squared_error
import warnings
warnings.filterwarnings('ignore')

plt.style.use('seaborn-v0_8-darkgrid')
```

### 1. What Is a Time Series?

A **time series** is a sequence of observations collected **over time**, usually at regular intervals.

Examples:

- Daily stock prices
- Hourly electricity consumption
- Monthly sales revenue
- Sensor measurements every second

Why Time Series Are Different

- Observations are **dependent**

- The order of data points **cannot be shuffled**
- Past values influence future values

This dependency structure is the core challenge of time series analysis.

## Key Components

Time series typically contain four components:

1. **Trend (T)**: Long-term increase or decrease in the data
2. **Seasonality (S)**: Regular, predictable patterns that repeat over fixed periods (e.g., yearly, monthly)
3. **Cyclicality (C)**: Patterns that repeat but not at fixed intervals (e.g., economic cycles)
4. **Noise/Irregular (I)**: Random variation that cannot be attributed to trend, seasonality, or cyclicality

## Mathematical Representation

- **Additive Model**:  $Y(t) = T(t) + S(t) + C(t) + I(t)$ 
  - Use when seasonal variation is roughly constant over time
- **Multiplicative Model**:  $Y(t) = T(t) \times S(t) \times C(t) \times I(t)$ 
  - Use when seasonal variation increases with the level of the series

## Creating a Sample Time Series

The code below makes fake daily data that looks like a real-world time series.

More specifically:

- It creates **daily dates from 2020 to 2023**.
- It makes the values **slowly increase over time**.
- It adds a **repeating yearly up-and-down pattern**.
- It adds a bit of **random randomness** so it is not perfectly smooth.
- It combines all of that into one dataset and shows a small sample.

In short:

**It simulates realistic daily data with a trend, seasonality, and noise.**

```
In [28]: # Set random seed for reproducibility
np.random.seed(42)

# Generate date range
```

```

date_range = pd.date_range(start='2020-01-01', end='2023-12-31', freq='D')
n = len(date_range)

# Components
trend = np.linspace(100, 150, n) # Linear trend from 100 to 150
seasonality = 10 * np.sin(2 * np.pi * np.arange(n) / 365.25) # Yearly seasonality
noise = np.random.normal(0, 5, n) # Random noise

# Combine components (additive model)
ts_data = trend + seasonality + noise

# Create pandas Series with datetime index
ts = pd.Series(ts_data, index=date_range, name='Value')

print("Sample Time Series:")
print(ts.head(10))
print(f"\nShape: {ts.shape}")
print(f"Period: {ts.index.min()} to {ts.index.max()}")

```

Sample Time Series:

2020-01-01	102.483571
2020-01-02	99.514941
2020-01-03	103.650916
2020-01-04	108.233733
2020-01-05	99.653774
2020-01-06	99.859609
2020-01-07	109.131857
2020-01-08	105.278161
2020-01-09	99.298455
2020-01-10	104.563060

Freq: D, Name: Value, dtype: float64

Shape: (1461,)

Period: 2020-01-01 00:00:00 to 2023-12-31 00:00:00

## Visualizing the Time Series

The code below draws a line chart of the time-series data.


More specifically:

- It creates a **plotting area** of a fixed size.
- It **plots the time series as a line** on that area.
- It adds a **title** and **labels** for the x- and y-axes.
- It turns on a **light grid** to make the chart easier to read.
- It adjusts spacing so labels are not cut off.
- It **displays the chart**.

In short:

**It visualizes the time series so you can see the trend and seasonal pattern over time.**

```
In [29]: #subplot function returns a tuple (figure, axes), figure is the entire figure, axes
# Axes object is what we plot on
# figsize is in inches
fig, ax = plt.subplots(figsize=(14, 5))
ts.plot(ax=ax, linewidth=1, color='steelblue')
ax.set_title('Sample Time Series with Trend and Seasonality', fontsize=14, fontweight='bold')
ax.set_xlabel('Date', fontsize=12)
ax.set_ylabel('Value', fontsize=12)
ax.grid(True, alpha=0.3)
plt.tight_layout() # Adjust Layout to prevent clipping of labels
plt.show()
```



## 2. Time Series Decomposition

### Theory

Decomposition is the process of separating a time series into its constituent components. This helps us:

- Understand the underlying patterns
- Remove seasonality for better modeling
- Identify anomalies
- Choose appropriate forecasting methods

### Classical Decomposition Methods

The most common method is **moving average decomposition**:

1. Estimate trend using moving average
2. Detrend the series
3. Estimate seasonal component by averaging detrended values for each season
4. Calculate residuals: Residual = Original - Trend - Seasonal

## Performing Decomposition

The code below breaks the time series into its main parts.

More specifically:

- It tells Python the data has a **yearly pattern** (365 days).
- It **separates the data** into:
  - a **trend** (long-term direction),
  - a **seasonal pattern** (repeating yearly ups and downs),
  - and **random noise** (what's left over).
- It stores each part separately.
- It prints the size of each component to confirm they match the original data.

In short:

**It decomposes the time series into trend, seasonality, and noise.**

```
In [30]: # Perform additive decomposition
# Period = 365 because we have daily data with yearly seasonality
decomposition = seasonal_decompose(ts, model='additive', period=365)

# Extract components
trend_component = decomposition.trend
seasonal_component = decomposition.seasonal
residual_component = decomposition.resid

print("Decomposition Components:")
print(f"Trend: {trend_component.shape}")
print(f"Seasonal: {seasonal_component.shape}")
print(f"Residual: {residual_component.shape}")
```

Decomposition Components:

```
Trend: (1461,)
Seasonal: (1461,)
Residual: (1461,)
```

## Visualizing Decomposition

The code below draws four stacked charts to show the decomposition results.

More specifically:

- It creates **four plots arranged vertically**.
- The first plot shows the **original time series**.

- The second shows the **trend** (long-term movement).
- The third shows the **seasonal pattern** (repeating yearly cycle).
- The fourth shows the **residuals** (random noise left over).
- It labels each plot so you can clearly see what each line represents.
- It draws a zero line on the residual plot to make deviations easy to spot.
- Finally, it displays the figure.

In short:

**It visually explains how the original data is split into trend, seasonality, and noise.**

```
In [31]: fig, axes = plt.subplots(4, 1, figsize=(14, 10))


# Original
ts.plot(ax=axes[0], linewidth=1, color='black')
axes[0].set_ylabel('Original', fontsize=11, fontweight='bold')
axes[0].set_title('Time Series Decomposition (Additive Model)', fontsize=14, fontweight='bold')

# Trend
trend_component.plot(ax=axes[1], linewidth=1.5, color='orange')
axes[1].set_ylabel('Trend', fontsize=11, fontweight='bold')

# Seasonal
seasonal_component.plot(ax=axes[2], linewidth=1, color='green')
axes[2].set_ylabel('Seasonal', fontsize=11, fontweight='bold')

# Residual
residual_component.plot(ax=axes[3], linewidth=0.8, color='red', alpha=0.7)
axes[3].set_ylabel('Residual', fontsize=11, fontweight='bold')
axes[3].set_xlabel('Date', fontsize=12)
axes[3].axhline(y=0, color='black', linestyle='--', linewidth=0.8, alpha=0.5)

plt.tight_layout()
plt.show()
```



### 3. Stationarity

#### What is Stationarity?

A time series is **stationary** if its statistical properties (mean, variance, autocorrelation) do not change over time.

#### Why Does Stationarity Matter?

Most classical time series models (ARIMA, SARIMA) assume stationarity because:

- Statistical properties are easier to model when constant
- Predictions are more reliable
- Mathematical theory is simpler

#### Types of Stationarity

1. **Strict Stationarity:** Joint distribution is time-invariant (very restrictive)
2. **Weak/Covariance Stationarity:** Only mean, variance, and autocorrelation are constant (commonly used)

#### Common Non-Stationary Patterns

- **Trend:** Mean changes over time
- **Seasonality:** Pattern repeats at regular intervals
- **Heteroscedasticity:** Variance changes over time

## Augmented Dickey-Fuller (ADF) Test

The ADF test checks the null hypothesis that the series has a unit root (non-stationary).

- **$H_0$ :** Series has a unit root (non-stationary)
- **$H_1$ :** Series is stationary

**Decision Rule:** If p-value < 0.05, reject  $H_0$  (series is stationary)

```
In [32]: def check_stationarity(timeseries, name='Series'):
    """
    Perform Augmented Dickey-Fuller test
    """
    print(f"\n{'*'*60}")
    print(f"Stationarity Test: {name}")
    print('='*60)

    # Remove NaN values
    ts_clean = timeseries.dropna()

    # Perform ADF test
    result = adfuller(ts_clean, autolag='AIC')

    print(f'ADF Statistic:      {result[0]:.6f}')
    print(f'p-value:            {result[1]:.6f}')
    print(f'# Lags Used:       {result[2]}')
    print(f'# Observations:     {result[3]}')
    print('\nCritical Values:')
    for key, value in result[4].items():
        print(f'  {key}: {value:.3f}')

    # Interpretation
    print('\n' + '-'*60)
    if result[1] <= 0.05:
        print(f'✓ STATIONARY (p={result[1]:.4f} < 0.05)')
        print("  → Reject null hypothesis")
    else:
        print(f'✗ NON-STATIONARY (p={result[1]:.4f} > 0.05)')
        print("  → Fail to reject null hypothesis")
    print('-'*60)

    return result[1] <= 0.05
```

## Testing Our Series

```
In [33]: # Test original series
is_stationary = check_stationarity(ts, 'Original Time Series')

=====
Stationarity Test: Original Time Series
=====
ADF Statistic: -0.835398
p-value: 0.808502
# Lags Used: 12
# Observations: 1448

Critical Values:
1%: -3.435
5%: -2.864
10%: -2.568

-----
X NON-STATIONARY (p=0.8085 > 0.05)
→ Fail to reject null hypothesis
-----
```

## Making a Series Stationary using the Differencing Method

Differencing removes trends and can stabilize the mean:

**First Difference:**  $Y'(t) = Y(t) - Y(t-1)$

```
In [34]: # Apply first differencing
ts_diff = ts.diff().dropna()

# Test differenced series
is_stationary_diff = check_stationarity(ts_diff, 'First Differenced Series')

=====
Stationarity Test: First Differenced Series
=====
ADF Statistic: -17.823692
p-value: 0.000000
# Lags Used: 11
# Observations: 1448

Critical Values:
1%: -3.435
5%: -2.864
10%: -2.568

-----
✓ STATIONARY (p=0.0000 < 0.05)
→ Reject null hypothesis
-----
```

## Visualizing the Transformation

The code below compares the original data with a transformed version that removes trend.

More specifically:

- The top plot shows the **original time series**, which changes over time (non-stationary).
- The bottom plot shows the **first-differenced series**, where each value is replaced by the change from the previous one.
- The differenced series fluctuates around zero, making it **more stable (stationary)**.
- The horizontal zero line helps you see this stability.
- Both plots share the same time axis for easy comparison.

In short:

It shows how differencing turns a trending series into a stationary one.

```
In [35]: fig, axes = plt.subplots(2, 1, figsize=(14, 8), sharex=True)

# Original
ts.plot(ax=axes[0], linewidth=1.2, color='steelblue')
axes[0].set_title('Original Series (Non-Stationary)', fontsize=12, fontweight='bold')
axes[0].set_ylabel('Value', fontsize=11)
axes[0].grid(True, alpha=0.3)

# Differenced
ts_diff.plot(ax=axes[1], linewidth=1, color='orange')
axes[1].set_title('First Differenced Series (Stationary)', fontsize=12, fontweight='bold')
axes[1].set_ylabel('Differenced Value', fontsize=11)
axes[1].axhline(y=0, color='red', linestyle='--', linewidth=1, alpha=0.7)
axes[1].set_xlabel('Date', fontsize=12)
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



## 4. Autocorrelation Analysis

### ACF: Autocorrelation Function

The ACF measures the correlation between observations at different time lags:

$$\text{ACF}(k) = \text{Corr}(Y_t, Y_{t-k})$$

- Lag 1: correlation between consecutive observations
- Lag 2: correlation between observations 2 time periods apart
- And so on...

### PACF: Partial Autocorrelation Function

The PACF measures the correlation between  $Y_t$  and  $Y_{t-k}$  after removing the effect of intermediate lags.

### Why Are ACF and PACF Important?

They help identify the order of AR and MA components:

Pattern	Model Suggestion
PACF cuts off after lag p, ACF decays	AR(p)

Pattern	Model Suggestion
ACF cuts off after lag q, PACF decays	MA(q)
Both decay gradually	ARMA(p,q)

## Plotting ACF and PACF

The code below looks at how the data is related to its past values.

More specifically:

- It uses the **stationary (differenced) series**.
- The **ACF plot** shows how today's value is correlated with previous days.
- The **PACF plot** shows the *direct* effect of past days, removing indirect effects.
- The plots help decide how many past values a time-series model (like ARIMA) should use.

In short:


**It helps you choose the right ARIMA model by showing which past values matter.**

```
In [36]: # Use differenced series (should be stationary)
fig, axes = plt.subplots(1, 2, figsize=(15, 5))

# ACF
plot_acf(ts_diff, lags=40, ax=axes[0])
axes[0].set_title('Autocorrelation Function (ACF)', fontsize=12, fontweight='bold')
axes[0].set_xlabel('Lag', fontsize=11)
axes[0].set_ylabel('Correlation', fontsize=11)

# PACF
plot_pacf(ts_diff, lags=40, ax=axes[1])
axes[1].set_title('Partial Autocorrelation Function (PACF)', fontsize=12, fontweight='bold')
axes[1].set_xlabel('Lag', fontsize=11)
axes[1].set_ylabel('Partial Correlation', fontsize=11)

plt.tight_layout()
plt.show()
```



### Interpretation:

- Blue shaded area represents the confidence interval
- Bars outside this area are statistically significant
- Look for where bars drop inside the confidence interval (cutoff point)

The ACF plot shows a single strong and statistically significant spike at lag 1, followed by correlations that lie within the confidence bounds at higher lags. This indicates a clear cutoff in the autocorrelation function after the first lag.

The PACF plot does not exhibit a sharp cutoff. Instead, the partial autocorrelations decay gradually toward zero over several lags, indicating a tailing-off pattern.

This combination—ACF cutting off at lag 1 and PACF tailing off—is characteristic of a moving-average process of order 1. When applied to a first-differenced series, this pattern is consistent with an ARIMA(0,1,1) model.

## 5. Classical Time Series Models

Model	Full Name	Equation	When to Use
AR(p)	Autoregressive	$Y_t = c + \varphi_1 Y_{t-1} + \dots + \varphi_p Y_{t-p} + \varepsilon_t$	PACF cuts off
MA(q)	Moving Average	$Y_t = \mu + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \dots + \theta_q \varepsilon_{t-q}$	ACF cuts off
ARMA(p,q)	AR + MA	Combines both	Both ACF & PACF decay
ARIMA(p,d,q)	Integrated ARMA	ARMA on d-times differenced data	Non-stationary series
SARIMA(p,d,q) (P,D,Q,s)	Seasonal ARIMA	Adds seasonal components	Seasonal patterns

## Preparing Data for Modeling

Let's create a simpler dataset to demonstrate the models:

The code below **creates a simple autoregressive time series and splits it into training and test data.**

More specifically:

- It generates **300 days of fake daily data**.
- Each day's value depends on **70% of the previous day's value**, plus some random noise. This is an **AR(1) process**.
- The result is a series where values are **correlated with their immediate past**.
- The data is then split into:
  - **80% for training** (used to fit a model),
  - **20% for testing** (used to evaluate predictions).
- Finally, it prints how many observations are in each part.

In short:

The code simulates a simple AR(1) time series and prepares it for model training and evaluation.

```
In [37]: # Generate AR(1) process
np.random.seed(123)
n = 300
dates = pd.date_range(start='2022-01-01', periods=n, freq='D')

# AR(1): Y(t) = 0.7 * Y(t-1) + noise
ar_coef = 0.7
ar_data = [0]
for i in range(1, n):
    ar_data.append(ar_coef * ar_data[i-1] + np.random.normal(0, 1))

ts_simple = pd.Series(ar_data, index=dates, name='Value')

# Train-test split (80-20)
train_size = int(len(ts_simple) * 0.8)
train = ts_simple[:train_size]
test = ts_simple[train_size:]

print(f"Train: {len(train)} observations")
print(f"Test: {len(test)} observations")
```


Train: 240 observations  
Test: 60 observations

Plot the dataset

```
In [38]: import matplotlib.pyplot as plt

# Plot the AR(1) time series
```

```
plt.figure(figsize=(14, 5))
ts_simple.plot(linewidth=1.2)
plt.title('Simulated AR(1) Time Series', fontsize=14, fontweight='bold')
plt.xlabel('Date', fontsize=12)
plt.ylabel('Value', fontsize=12)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



## Model 1: AR (Autoregressive)

### Theory

An AR( $p$ ) model predicts the current value using  $p$  past values:

$$Y_t = c + \varphi_1 Y_{t-1} + \varphi_2 Y_{t-2} + \dots + \varphi_p Y_{t-p} + \varepsilon_t$$

Where:

- $c$  is a constant
- $\varphi_1, \varphi_2, \dots, \varphi_p$  are coefficients
- $\varepsilon_t$  is white noise error

### Implementation

The code below **fits an AR(1) model to the training data and compares what the model learns to the true value used to generate the data.**

More specifically:

- It fits an **ARIMA(1,0,0)** model, which is an **AR(1)** model (one autoregressive term, no differencing, no moving average).
- The model is trained using the **training portion** of the simulated series.
- It prints a **summary** showing estimated parameters and diagnostics.
- It extracts the **estimated AR(1) coefficient** from the fitted model.

- It compares that estimate to the **true coefficient (0.7)** used to generate the data.

In short:

The code checks whether an AR(1) model can correctly recover the underlying dependency in the simulated time series.

```
In [39]: # Fit AR(1) model
# ARIMA(p,d,q) with p=1, d=0, q=0
ar_model = ARIMA(train, order=(1, 0, 0))
ar_fit = ar_model.fit()

print("AR(1) Model Summary:")
print(ar_fit.summary())
print(f"\nEstimated coefficient: {ar_fit.params[1]:.4f}")
print(f"True coefficient: {ar_coef}")
```

AR(1) Model Summary:

#### SARIMAX Results

```
=====
Dep. Variable:                      Value    No. Observations:                  240
Model:                            ARIMA(1, 0, 0)    Log Likelihood:                -345.782
Date:                            Tue, 30 Dec 2025    AIC:                         697.563
Time:                            08:17:58        BIC:                         708.005
Sample:                           01-01-2022    HQIC:                        701.771
                                  - 08-28-2022
Covariance Type:                    opg
=====
            coef      std err      z      P>|z|      [0.025      0.975]
-----
const    -0.0308      0.229    -0.135      0.893     -0.479      0.418
ar.L1     0.7119      0.047    15.074      0.000      0.619      0.805
sigma2    1.0415      0.094    11.072      0.000      0.857      1.226
=====
Ljung-Box (L1) (Q):                 0.15    Jarque-Bera (JB):                 0.10
Prob(Q):                           0.70    Prob(JB):                     0.95
Heteroskedasticity (H):              0.60    Skew:                          -0.04
Prob(H) (two-sided):                  0.03    Kurtosis:                     3.05
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

Estimated coefficient: 0.7119

True coefficient: 0.7

## Interpreting AR(1) Model Output - OPTIONAL

**What this really is (in one sentence)**

**AR(1) is just linear regression where the predictor is yesterday's value.**

## Connect it to regular regression (intuition first)

In ordinary regression you write:

$$y = \beta_0 + \beta_1 x_1 + \text{error}$$

In AR(1), you write:

$$y_t = \text{constant} + \beta_1 \times y_{t-1} + \text{error}$$

Same idea. The "feature" is just the lagged value of  $y$  itself.

## Translate the output into regression language

### 1. The coefficients

const	=	-0.0308
ar.L1	=	0.7119
sigma2	=	1.0415

Think of this as a regression table:

Feature	Coefficient	What it means
Intercept	-0.031	Baseline (when $y_{t-1} = 0$ )
$y_{t-1}$	0.7119	Today = 71% of yesterday + noise
Error variance	1.042	Residual spread

### The prediction equation is:

$$y_t = -0.031 + 0.7119 \times y_{t-1} + \varepsilon_t$$

This is identical to:

$$\text{sales\_today} = -0.031 + 0.7119 \times \text{sales\_yesterday} + \text{noise}$$

### 2. Standard errors and p-values

Exactly the same as regression:

	coef	std err	p-value
const	-0.031	0.229	0.893
ar.L1	0.712	0.047	0.000

**const:**  $p = 0.893$  (high)

Not significant. Could drop it.

**ar.L1:**  $p = 0.000$  (very low)

Highly significant. Yesterday strongly predicts today.

Same interpretation as:

feature\_1:  $p = 0.893 \rightarrow$  not useful feature\_2:  $p = 0.000 \rightarrow$  very useful

The model found **strong autocorrelation** in the data.

### 3. Model quality metrics

AIC = 697.6

BIC = 708.0

Log Likelihood = -345.8

These are **goodness-of-fit measures**, not accuracy metrics.

**AIC and BIC**: Lower is better

Use these to compare models:

- AR(1): AIC = 697.6
- AR(2): AIC = 695.3  $\rightarrow$  better
- MA(1): AIC = 701.2  $\rightarrow$  worse

Think of AIC like:

penalized training error

It balances fit with model complexity.

**Log Likelihood**: Higher is better

Similar to minimizing loss in ML. Maximum likelihood estimation finds coefficients that make the observed data most probable.

### 4. Diagnostic tests (are residuals well-behaved?)

Ljung-Box (Q):	0.15	$p = 0.70$
Jarque-Bera (JB):	0.10	$p = 0.95$
Heteroskedasticity (H):	0.60	$p = 0.03$

Think of these as **residual plots in test form**.

#### Ljung-Box test

Question: "Do residuals have patterns left?"

- $p > 0.05 \rightarrow$  No patterns (good)
- $p < 0.05 \rightarrow$  Patterns remain (bad)

Here:  $p = 0.70 \rightarrow$  residuals look random

This is like checking:

plot(residuals) shows no trend

### **Jarque-Bera test**

Question: "Are residuals normally distributed?"

- $p > 0.05 \rightarrow$  Yes (good)
- $p < 0.05 \rightarrow$  No (may need transformation)

Here:  $p = 0.95 \rightarrow$  very normal

This is like checking:

histogram(residuals) looks bell-shaped

### **Heteroskedasticity test**

Question: "Is variance constant over time?"

- $p > 0.05 \rightarrow$  Yes (good)
- $p < 0.05 \rightarrow$  No (variance changes)

Here:  $p = 0.03 \rightarrow$  some heteroskedasticity detected

This is like seeing:

residual spread increases over time

## **5. Model validation**

Estimated coefficient: 0.7119  
 True coefficient: 0.7000

The model recovered the true data-generating process.

In ML terms:

The model learned the correct function

This would be like:

- You generate data:  $y = 2x + \text{noise}$
- Your model learns:  $y = 1.98x$
- Close match  $\rightarrow$  model works

## **Why it feels alien**

## 1. No feature matrix visible

The "feature" ( $y_{t-1}$ ) is created internally from the time series.

## 2. Maximum likelihood instead of MSE

Same goal (fit the data), different math. Likelihood is more general than squared error.

## 3. Heavy focus on diagnostics

Econometrics cares deeply about **why** the model works, not just **that** it works.

Statistical inference > predictive accuracy

## 4. Different vocabulary

- ML says: " $R^2 = 0.85$ "
- Econometrics says: "Log Likelihood = -345, AIC = 697"

Same information, different packaging.

## How to read this output (decision rules)

### Check coefficients:

- $p < 0.05 \rightarrow$  significant  $\rightarrow$  keep
- $p > 0.05 \rightarrow$  not significant  $\rightarrow$  consider dropping

### Check diagnostics:

- Ljung-Box  $p > 0.05 \rightarrow$  good (no autocorrelation left)
- Ljung-Box  $p < 0.05 \rightarrow$  bad (model incomplete)

### Compare models:

- Lower AIC  $\rightarrow$  better model
- Lower BIC  $\rightarrow$  better model (penalizes complexity more)

### Validate:

- Residuals should look random
- Residuals should be roughly normal
- Variance should be constant (ideally)

## One grounding sentence you can remember

AR(1) is linear regression where X is yesterday's y, and the output tells you both fit quality and residual behavior in statistical language instead of ML metrics.

## Quick interpretation of this specific output

**Model found:**  $y_t = 0.71 \times y_{t-1} + \text{noise}$

**Constant term:** Not significant ( $p = 0.89$ ) → probably zero in reality

**Autocorrelation:** Strong (coefficient = 0.71,  $p < 0.001$ ) → yesterday matters a lot

**Residuals:** Clean (Ljung-Box  $p = 0.70$ , JB  $p = 0.95$ ) → no patterns left

**Minor issue:** Slight heteroskedasticity ( $p = 0.03$ ) → variance not perfectly constant

**Conclusion:** This is a good model. The AR(1) structure correctly captures the data-generating process.

## Model 2: MA (Moving Average)

### Theory

An MA( $q$ ) model predicts the current value using past forecast errors:

$$Y_t = \mu + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \dots + \theta_q \varepsilon_{t-q}$$

Where:

- $\mu$  is the mean
- $\theta_1, \theta_2, \dots, \theta_q$  are coefficients
- $\varepsilon_t, \varepsilon_{t-1}, \dots$  are error terms

### Implementation

The code below **fits a moving-average model of order 1 to the training data**.

More specifically:

- It fits an **ARIMA(0,0,1)** model, which is an **MA(1)** model.
- The model assumes the series depends on **one past shock (error term)** rather than past values.
- It estimates the MA coefficient using the **training data**.
- It prints a **model summary** with parameter estimates and diagnostics.

In short:

The code fits an MA(1) model so it can be compared with the AR(1) model on the same data.

```
In [40]: # Fit MA(1) model
ma_model = ARIMA(train, order=(0, 0, 1))
ma_fit = ma_model.fit()
```

```
print("MA(1) Model Summary:")
print(ma_fit.summary())
```

MA(1) Model Summary:

### SARIMAX Results

```
=====
Dep. Variable:                      Value     No. Observations:                  240
Model:                          ARIMA(0, 0, 1)   Log Likelihood:           -370.431
Date:                Tue, 30 Dec 2025   AIC:                            746.863
Time:                      08:17:58      BIC:                            757.305
Sample:                01-01-2022   HQIC:                           751.070
                           - 08-28-2022
Covariance Type:                    opg
=====
            coef    std err        z     P>|z|      [0.025    0.975]
-----
const    -0.0349    0.120    -0.290     0.771    -0.270     0.200
ma.L1      0.6287    0.051   12.246     0.000     0.528     0.729
sigma2     1.2801    0.119   10.766     0.000     1.047     1.513
=====
Ljung-Box (L1) (Q):                 10.99   Jarque-Bera (JB):             0.40
Prob(Q):                           0.00   Prob(JB):                   0.82
Heteroskedasticity (H):              0.60   Skew:                     -0.10
Prob(H) (two-sided):                0.03   Kurtosis:                  2.94
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step p).

## Model 3: ARMA

### Theory

ARMA(p,q) combines AR(p) and MA(q):

$$Y_t = c + \varphi_1 Y_{t-1} + \dots + \varphi_p Y_{t-p} + \theta_1 \varepsilon_{t-1} + \dots + \theta_q \varepsilon_{t-q} + \varepsilon_t$$

**Important:** ARMA requires stationary data.

### Implementation

The code below **fits a combined autoregressive and moving-average model of order (1,1)** to the training data.

More specifically:

- It fits an **ARIMA(1,0,1)** model, also known as **ARMA(1,1)**.
- The model allows the series to depend on:

- **one past value** (AR part), and
- **one past shock or error** (MA part).
- The model is trained on the **training dataset**.
- It prints a **summary** showing parameter estimates and diagnostics.

In short:

The code fits an ARMA(1,1) model so its performance and parameters can be compared against the simpler AR(1) and MA(1) models.

```
In [41]: # Fit ARMA(1,1) model
arma_model = ARIMA(train, order=(1, 0, 1))
arma_fit = arma_model.fit()

print("ARMA(1,1) Model Summary:")
print(arma_fit.summary())
```

ARMA(1,1) Model Summary:

#### SARIMAX Results

Dep. Variable:	Value	No. Observations:	240			
Model:	ARIMA(1, 0, 1)	Log Likelihood	-345.671			
Date:	Tue, 30 Dec 2025	AIC	699.341			
Time:	08:17:58	BIC	713.264			
Sample:	01-01-2022 - 08-28-2022	HQIC	704.951			
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
const	-0.0308	0.223	-0.138	0.890	-0.468	0.407
ar.L1	0.6912	0.072	9.667	0.000	0.551	0.831
ma.L1	0.0424	0.096	0.441	0.659	-0.146	0.231
sigma2	1.0405	0.094	11.083	0.000	0.857	1.225
Ljung-Box (L1) (Q):	0.00	Jarque-Bera (JB):	0.13			
Prob(Q):	0.96	Prob(JB):	0.94			
Heteroskedasticity (H):	0.60	Skew:	-0.05			
Prob(H) (two-sided):	0.02	Kurtosis:	3.05			

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

## Model 4: ARIMA

## Theory

ARIMA(p,d,q) is ARMA applied to d-times differenced data:

- **p**: Order of autoregressive part
- **d**: Degree of differencing
- **q**: Order of moving average part

### Steps:

1. Difference the series d times to make it stationary
2. Apply ARMA(p,q) to the differenced series

## Implementation

The code below **fits a full ARIMA model, evaluates it, and uses it to make forecasts.**

More specifically:

- It fits an **ARIMA(1,1,1)** model:
  - one autoregressive term,
  - one differencing step to remove trend,
  - one moving-average term.
- The model is trained on the **training data**.
- It prints a **model summary** with estimated parameters.
- It reports **AIC and BIC**, which measure model quality while penalizing complexity (lower is better).
- It then **predicts future values** for the same number of time steps as the test set.

In short:

The code builds a trend-aware ARIMA model and uses it to forecast the unseen portion of the time series.

```
In [42]: # Fit ARIMA(1,1,1)
arima_model = ARIMA(train, order=(1, 1, 1))
arima_fit = arima_model.fit()

print("ARIMA(1,1,1) Model Summary:")
print(arima_fit.summary())
print(f"\nAIC: {arima_fit.aic:.2f}")
print(f"BIC: {arima_fit.bic:.2f}")

# Forecast
forecast_steps = len(test)
arima_forecast = arima_fit.forecast(steps=forecast_steps)
```

## ARIMA(1,1,1) Model Summary:

## SARIMAX Results

```
=====
Dep. Variable:                      Value     No. Observations:                  240
Model:                 ARIMA(1, 1, 1)   Log Likelihood:           -346.335
Date:            Tue, 30 Dec 2025   AIC:                         698.670
Time:                08:17:58      BIC:                         709.100
Sample:             01-01-2022   HQIC:                        702.873
                           - 08-28-2022
Covariance Type:                    opg
=====
              coef    std err        z      P>|z|      [0.025      0.975]
-----
ar.L1      0.7190    0.052    13.919      0.000      0.618      0.820
ma.L1     -0.9999    2.479    -0.403      0.687     -5.860      3.860
sigma2     1.0460    2.583     0.405      0.685     -4.016      6.108
=====
Ljung-Box (L1) (Q):                  0.08  Jarque-Bera (JB):          0.08
Prob(Q):                            0.78  Prob(JB):                0.96
Heteroskedasticity (H):               0.59  Skew:                   -0.03
Prob(H) (two-sided):                  0.02  Kurtosis:                 3.06
=====
```

## Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step p).

AIC: 698.67


BIC: 709.10

## Visualizing Forecast

```
In [43]: fig, ax = plt.subplots(figsize=(14, 5))

train.plot(ax=ax, label='Train', linewidth=1.5, color='steelblue')
test.plot(ax=ax, label='Test (Actual)', linewidth=1.5, color='green')
arima_forecast.plot(ax=ax, label='ARIMA Forecast', linewidth=2,
                     linestyle='--', color='red')

ax.set_title('ARIMA(1,1,1) Forecast', fontsize=14, fontweight='bold')
ax.set_xlabel('Date', fontsize=12)
ax.set_ylabel('Value', fontsize=12)
ax.legend(fontsize=11)
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



## Model 5: SARIMA (Seasonal ARIMA)

### Theory

SARIMA( $p,d,q$ )( $P,D,Q,s$ ) extends ARIMA to handle seasonality:

**Non-seasonal part:** ( $p,d,q$ ) **Seasonal part:** ( $P,D,Q,s$ )

- P: Seasonal AR order
- D: Seasonal differencing order
- Q: Seasonal MA order
- s: Seasonal period (e.g., 12 for monthly data with yearly seasonality)

### Creating Seasonal Data

The code below **creates fake monthly data that looks like long-term seasonal sales and splits it into training and test sets.**

More specifically:

- It creates **20 years of monthly dates**.
- It builds data with:
  - a **slow upward trend** over time,
  - a **repeating yearly pattern** (seasonality every 12 months),
  - and some **random noise**.
- These components are added together to look like realistic sales data.
- The data is then split into:
  - **80% for training**,
  - **20% for testing**.
- Finally, it prints the size of the full series and each split.

In short:

The code simulates a realistic monthly time series with trend and seasonality and prepares it for seasonal time-series modeling.

```
In [44]: # Generate monthly data with seasonality
n_months = 240 # 20 years
dates_monthly = pd.date_range(start='2004-01-01', periods=n_months, freq='MS')

# Components
trend = np.linspace(50, 100, n_months)
seasonal = 15 * np.sin(2 * np.pi * np.arange(n_months) / 12)
noise = np.random.normal(0, 3, n_months)

ts_seasonal = pd.Series(trend + seasonal + noise,
                       index=dates_monthly,
                       name='Sales')

# Split
train_seas = ts_seasonal[:int(len(ts_seasonal) * 0.8)]
test_seas = ts_seasonal[int(len(ts_seasonal) * 0.8):]

print(f"Seasonal series shape: {ts_seasonal.shape}")
print(f"Train: {len(train_seas)}, Test: {len(test_seas)}")
```

Seasonal series shape: (240,)

Train: 192, Test: 48

## Fitting SARIMA

The code below **fits a seasonal ARIMA model and uses it to make forecasts**.

More specifically:

- It fits a **SARIMA(1,1,1)(1,1,1,12)** model, which means:
  - one autoregressive term, one differencing, and one moving-average term for the **non-seasonal part**,
  - one autoregressive term, one differencing, and one moving-average term for the **seasonal part**,
  - with a **12-month seasonal cycle**.
- The model is trained on the **seasonal training data**.
- It prints a **summary** showing estimated parameters and diagnostics.
- It then **forecasts future values** for the length of the test period.

In short:

The code models both long-term trends and repeating yearly patterns, then uses that model to predict future seasonal behavior.

```
In [45]: # Fit SARIMA(1,1,1)(1,1,1,12)
# Seasonal period = 12 months
sarima_model = SARIMAX(train_seas,
                       order=(1, 1, 1),
                       seasonal_order=(1, 1, 1, 12))
sarima_fit = sarima_model.fit(disp=False)

print("SARIMA(1,1,1)(1,1,1,12) Model Summary:")
print(sarima_fit.summary())

# Forecast
sarima_forecast = sarima_fit.forecast(steps=len(test_seas))
```

SARIMA(1,1,1)(1,1,1,12) Model Summary:

		SARIMAX Results				
Dep. Variable:	Sales	No. Observations:				
192			-4			
Model:	SARIMAX(1, 1, 1)x(1, 1, 1, 12)	Log Likelihood				
63.423			9			
Date:	Tue, 30 Dec 2025	AIC				
36.845			9			
Time:	08:17:59	BIC				
52.782			9			
Sample:	01-01-2004	HQIC				
43.307			9			
	- 12-01-2019					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.0128	0.084	-0.151	0.880	-0.178	0.152
ma.L1	-0.9801	0.040	-24.634	0.000	-1.058	-0.902
ar.S.L12	-0.0238	0.104	-0.230	0.818	-0.227	0.179
ma.S.L12	-0.9156	0.140	-6.530	0.000	-1.190	-0.641
sigma2	8.9146	1.336	6.671	0.000	6.295	11.534
Ljung-Box (L1) (Q):		0.03	Jarque-Bera (JB):		1.26	
Prob(Q):		0.87	Prob(JB):		0.53	
Heteroskedasticity (H):		1.30	Skew:		-0.07	
Prob(H) (two-sided):		0.31	Kurtosis:		2.61	

Warnings:


[1] Covariance matrix calculated using the outer product of gradients (complex-step).

## Visualizing Seasonal Forecast

```
In [46]: fig, ax = plt.subplots(figsize=(14, 5))

train_seas.plot(ax=ax, label='Train', linewidth=1.5, color='steelblue')
test_seas.plot(ax=ax, label='Test (Actual)', linewidth=1.5, color='green')
sarima_forecast.plot(ax=ax, label='SARIMA Forecast', linewidth=2,
                      linestyle='--', color='red')

ax.set_title('SARIMA Forecast (Monthly Data with Seasonality)',
             fontsize=14, fontweight='bold')
ax.set_xlabel('Date', fontsize=12)
ax.set_ylabel('Sales', fontsize=12)
ax.legend(fontsize=11)
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



## 6. Model Selection

### Information Criteria

#### AIC (Akaike Information Criterion)

$$\text{AIC} = 2k - 2\ln(L)$$

Where:

- k = number of parameters
- L = maximum likelihood

#### BIC (Bayesian Information Criterion)

$$\text{BIC} = k \cdot \ln(n) - 2\ln(L)$$

Where:

- n = number of observations

**Rule:** Lower AIC/BIC = better model (balances fit and complexity)

## Grid Search for Best Model

```
In [47]: def evaluate_arima_models(data, p_range, d_range, q_range):
    """
    Evaluate different ARIMA configurations
    """
    results = []

    for p in p_range:
        for d in d_range:
            for q in q_range:
                try:
                    model = ARIMA(data, order=(p, d, q))
                    fitted = model.fit()

                    results.append({
                        'order': (p, d, q),
                        'AIC': fitted.aic,
                        'BIC': fitted.bic
                    })
                except:
                    continue

    return pd.DataFrame(results)

# Search for best model
print("Searching for best ARIMA model...")
results_df = evaluate_arima_models(
    train,
    p_range=range(0, 3),
    d_range=range(0, 2),
    q_range=range(0, 3)
)

# Sort and display
results_df = results_df.sort_values('AIC')
print("\nTop 5 Models by AIC:")
print(results_df.head())

best_order = results_df.iloc[0]['order']
print(f"\nBest Model: ARIMA{best_order}")
```

Searching for best ARIMA model...

Top 5 Models by AIC:

	order	AIC	BIC
6	(1, 0, 0)	697.563191	708.005108
13	(2, 0, 1)	698.028252	715.431446
10	(1, 1, 1)	698.670168	709.099558
12	(2, 0, 0)	699.335230	713.257786
7	(1, 0, 1)	699.341125	713.263681

Best Model: ARIMA(1, 0, 0)

## 7. Model Diagnostics

### Why Check Diagnostics?

After fitting a model, we need to verify that:

1. **Residuals are white noise** (random, no pattern)
2. **Residuals are normally distributed**
3. **No autocorrelation in residuals**
4. **Model assumptions are satisfied**

### Theory: Good Residuals

If the model is appropriate, residuals should have:

- Mean  $\approx 0$
- Constant variance (homoscedastic)
- No autocorrelation
- Normal distribution

### Diagnostic Plots

```
In [48]: # Get residuals from best model
best_model = ARIMA(train, order=best_order).fit()
residuals = best_model.resid

# Create diagnostic plots
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# 1. Residuals over time
residuals.plot(ax=axes[0, 0], linewidth=0.8, color='steelblue', alpha=0.7)
axes[0, 0].axhline(y=0, color='red', linestyle='--', linewidth=1)
axes[0, 0].set_title('Residuals Over Time', fontsize=12, fontweight='bold')
axes[0, 0].set_ylabel('Residual', fontsize=11)
axes[0, 0].grid(True, alpha=0.3)
```

```


# 2. Histogram
residuals.hist(ax=axes[0, 1], bins=30, edgecolor='black', color='skyblue')
axes[0, 1].set_title('Residuals Distribution', fontsize=12, fontweight='bold')
axes[0, 1].set_xlabel('Residual', fontsize=11)
axes[0, 1].set_ylabel('Frequency', fontsize=11)
axes[0, 1].grid(True, alpha=0.3, axis='y')

# 3. ACF of residuals
plot_acf(residuals, lags=30, ax=axes[1, 0])
axes[1, 0].set_title('ACF of Residuals', fontsize=12, fontweight='bold')

# 4. Q-Q plot
stats.probplot(residuals, dist="norm", plot=axes[1, 1])
axes[1, 1].set_title('Q-Q Plot', fontsize=12, fontweight='bold')
axes[1, 1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```



## Ljung-Box Test

Tests the null hypothesis that residuals are independently distributed (no autocorrelation).

- $H_0$ : Residuals are white noise (no autocorrelation)
- $H_1$ : Residuals have autocorrelation

**Decision:** p-value > 0.05 → residuals are white noise (good!)

```
In [49]: # Perform Ljung-Box test
lb_test = acorr_ljungbox(residuals, lags=[10, 20, 30], return_df=True)

print("\nLjung-Box Test Results:")
print(lb_test)
print("\nInterpretation:")
print("If p-value > 0.05: Residuals are white noise √")
print("If p-value < 0.05: Residuals have autocorrelation X")
```

Ljung-Box Test Results:  
 lb\_stat lb\_pvalue  
 10 7.147022 0.711496  
 20 31.838671 0.045048  
 30 41.570079 0.077860

Interpretation:  
 If p-value > 0.05: Residuals are white noise √  
 If p-value < 0.05: Residuals have autocorrelation X

## Residual Statistics

```
In [50]: print("\nResidual Summary Statistics:")
print("*50)
print(f"Mean: {residuals.mean():>10.6f} (should be ≈ 0)")
print(f"Std Dev: {residuals.std():>10.6f}")
print(f"Min: {residuals.min():>10.6f}")
print(f"Max: {residuals.max():>10.6f}")
print(f"Skewness: {residuals.skew():>10.6f} (should be ≈ 0)")
print(f"Kurtosis: {residuals.kurtosis():>10.6f} (should be ≈ 0)")
```

Residual Summary Statistics:  
=====
Mean: -0.000090 (should be ≈ 0)
Std Dev: 1.022687
Min: -3.222888
Max: 2.614799
Skewness: -0.045184 (should be ≈ 0)
Kurtosis: 0.071627 (should be ≈ 0)

## 8. Model Evaluation

### Common Forecasting Metrics

#### 1. MAE (Mean Absolute Error)

$$\text{MAE} = (1/n) \sum |y_i - \hat{y}_i|$$

- Average absolute difference
- Same units as original data

- Easy to interpret

## 2. RMSE (Root Mean Squared Error)

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum (y_i - \hat{y}_i)^2}$$

- Penalizes large errors more
- Same units as original data
- More sensitive to outliers than MAE

## 3. MAPE (Mean Absolute Percentage Error)

$$\text{MAPE} = \left( \frac{100}{n} \right) \sum |(y_i - \hat{y}_i)/y_i|$$

- Expressed as percentage
- Scale-independent
- Undefined when  $y_i = 0$

### Computing Metrics

```
In [51]: def calculate_metrics(actual, predicted):
    """Calculate forecasting metrics"""
    mae = mean_absolute_error(actual, predicted)
    rmse = np.sqrt(mean_squared_error(actual, predicted))
    mape = np.mean(np.abs((actual - predicted) / actual)) * 100

    return {'MAE': mae, 'RMSE': rmse, 'MAPE': mape}

# Evaluate forecast
forecast = best_model.forecast(steps=len(test))
metrics = calculate_metrics(test, forecast)

print("\nForecast Evaluation Metrics:")
print("=="*50)
print(f"MAE: {metrics['MAE']:.4f}")
print(f"RMSE: {metrics['RMSE']:.4f}")
print(f"MAPE: {metrics['MAPE']:.2f}%")
```

Forecast Evaluation Metrics:  
=====

MAE: 1.2387  
RMSE: 1.5522  
MAPE: 102.35%

### Summary and Key Takeaways

### Workflow for Time Series Modeling

## 1. Explore Data

- Plot the series
- Identify trend, seasonality, noise
- Perform decomposition

## 2. Check Stationarity

- Use ADF test
- Apply differencing if needed
- Verify stationarity after transformation

## 3. Identify Model Orders

- Examine ACF plot → suggests MA order (q)
- Examine PACF plot → suggests AR order (p)
- Consider seasonal patterns

## 4. Fit Candidate Models

- Start with simple models (AR, MA)
- Try ARIMA for non-stationary data
- Use SARIMA for seasonal data

## 5. Select Best Model

- Compare AIC/BIC values
- Lower is better
- Balance complexity and fit

## 6. Validate Model

- Check residual plots
- Perform Ljung-Box test
- Ensure residuals are white noise

## 7. Forecast and Evaluate

- Generate predictions
- Calculate MAE, RMSE, MAPE
- Compare with test data

## Model Quick Reference

If you see...	Consider...
Clear trend	Differencing ( $d > 0$ ) or ARIMA
Seasonality	SARIMA with appropriate period
PACF cuts off at lag p	AR(p)
ACF cuts off at lag q	MA(q)

If you see...	Consider...
Both decay slowly	ARMA or ARIMA
Non-constant variance	Log transformation or multiplicative model

## Python Libraries

- **pandas**: Data manipulation and time series structures
- **statsmodels**: ARIMA, SARIMA, statistical tests, diagnostics
- **matplotlib**: Visualization
- **scipy**: Statistical functions
- **sklearn**: Evaluation metrics

## Next Steps

In **Part 2**, we'll cover:

- Advanced pandas techniques for time series
- Resampling and frequency conversion
- Rolling windows and moving averages
- Time-based indexing and slicing
- Practical data preparation workflows

In **Part 3**, we'll explore:

- Feature engineering for ML models
- Machine learning approaches (XGBoost, Random Forest)
- Deep learning (LSTM, GRU)
- Prophet and modern forecasting tools

*End of Part 1*