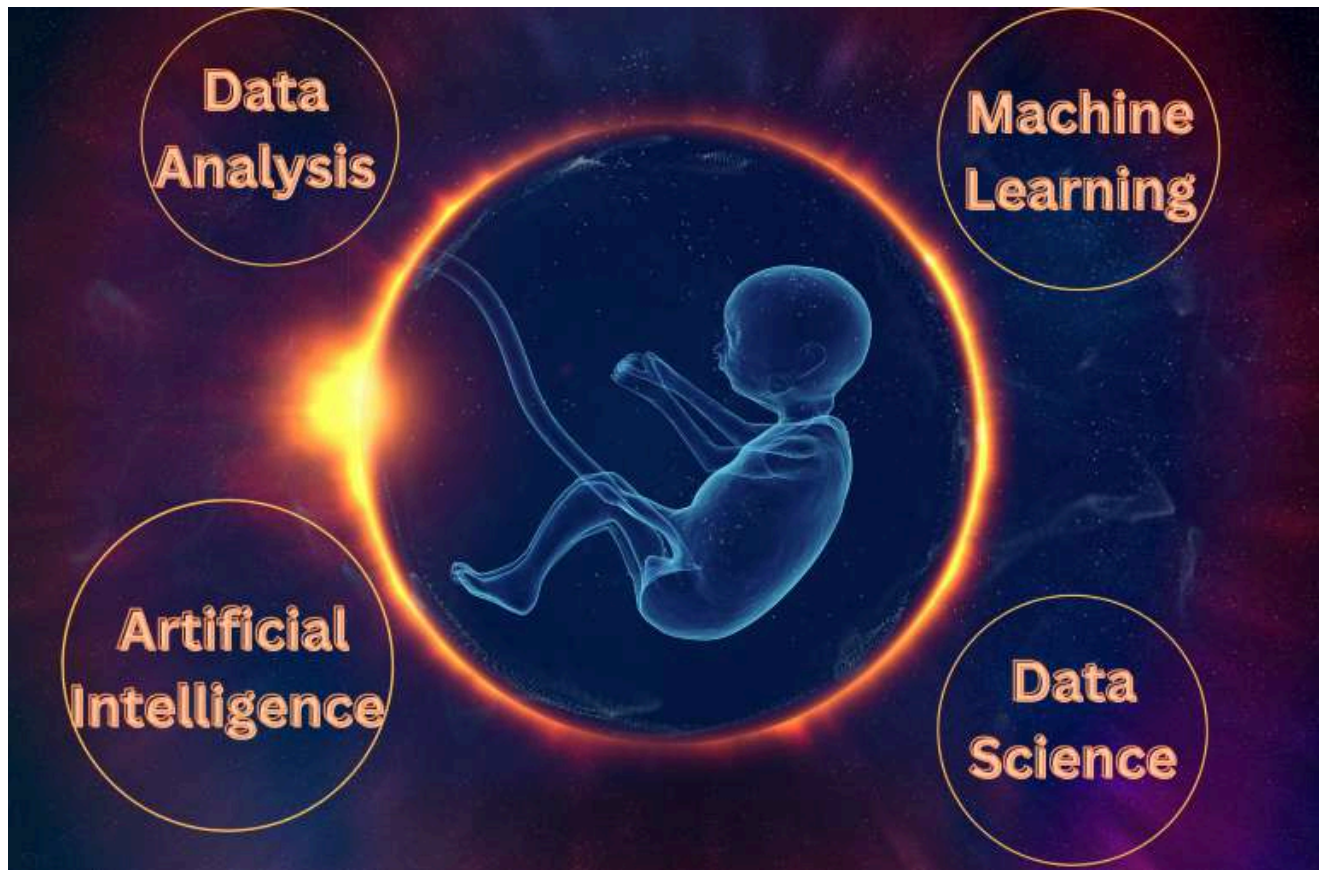


Fetal Health Classification



Fetal Health Classification

The **reduction of child and maternal mortality** is a critical focus for global health and human development. Several of the United Nations' **Sustainable Development Goals** reflect this priority, aiming to **end preventable deaths of newborns and children under 5 by 2030**, while reducing maternal mortality, which claimed over **295,000 lives in 2017**, the vast majority in low-resource settings.

Cardiotocograms (CTGs) provide a simple and cost-effective method to assess fetal health, allowing healthcare professionals to monitor **fetal heart rate, movements, and uterine contractions**. By leveraging CTG data, it is possible to detect potential issues early, preventing complications that could lead to **child or maternal mortality**.

This dataset contains **2126 records** derived from CTG exams, each classified by three expert obstetricians into three categories: **Normal**, **Suspect**, or **Pathological**. The goal of this project is to build a **machine learning model** to accurately predict the **health status of a fetus** based on the CTG data, thereby aiding in early diagnosis and timely medical intervention.



Join Faiz Siddiqui's Network

[LinkedIn](#)[GitHub](#)[Kaggle](#)

Faiz Siddiqui

A student of data science who loves learning and solving problems with machine learning. Dedicated to improving and helping others succeed in data science. Always excited to take on new challenges and share knowledge with the community.

Aim & Objectives

Aim

The aim of this notebook is to develop and evaluate machine learning models for classifying fetal health based on cardiotocogram (CTG) data. The goal is to accurately predict the health status of a fetus as either **Normal**, **Suspect**, or **Pathological**, facilitating early detection of potential issues that could lead to maternal or child mortality.

Objectives

- To perform an exploratory data analysis (EDA) on the fetal health dataset to understand key patterns and relationships within the data, specifically focusing on the **Normal**, **Suspect**, and **Pathological** health categories.
- To preprocess the CTG data, ensuring it is clean, standardized, and suitable for model training.
- To train and evaluate multiple machine learning models, including **XGBoost**, **Random Forest**, **SVM**, **KNN**, **Logistic Regression**, and ensemble techniques, for fetal health classification.
- To compare model performance based on **Accuracy** and **Matthews Correlation Coefficient (MCC)** and identify the best model for fetal health prediction.
- To provide insights and recommendations for healthcare professionals on the potential application of the selected model in clinical settings for early detection of fetal health issues.

```
In [2]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier, VotingClassifier, StackingClassifier
from sklearn.neighbors import KNeighborsClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, matthews_corrcoef
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize
from itertools import cycle
import plotly.graph_objs as go
import plotly.subplots as sp
plt.style.use('seaborn-v0_8-whitegrid')
```

```
In [3]: df = pd.read_csv('/kaggle/input/fetal-health-classification/fetal_health.csv')
df
```

Out[3]:

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolonged
0	120.0	0.000	0.000	0.000	0.000	0.000	0.0
1	132.0	0.006	0.000	0.006	0.003	0.003	0.0
2	133.0	0.003	0.000	0.008	0.003	0.003	0.0
3	134.0	0.003	0.000	0.008	0.003	0.003	0.0
4	132.0	0.007	0.000	0.008	0.000	0.000	0.0
...
2121	140.0	0.000	0.000	0.007	0.000	0.000	0.0
2122	140.0	0.001	0.000	0.007	0.000	0.000	0.0
2123	140.0	0.001	0.000	0.007	0.000	0.000	0.0
2124	140.0	0.001	0.000	0.006	0.000	0.000	0.0
2125	142.0	0.002	0.002	0.008	0.000	0.000	0.0

2126 rows × 22 columns

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2126 entries, 0 to 2125
```

```
Data columns (total 22 columns):
```

#	Column	Non-Null Count	Dtype
0	baseline value	2126 non-null	float64
1	accelerations	2126 non-null	float64
2	fetal_movement	2126 non-null	float64
3	uterine_contractions	2126 non-null	float64
4	light_decelerations	2126 non-null	float64
5	severe_decelerations	2126 non-null	float64
6	prolongued_decelerations	2126 non-null	float64
7	abnormal_short_term_variability	2126 non-null	float64
8	mean_value_of_short_term_variability	2126 non-null	float64
9	percentage_of_time_with_abnormal_long_term_variability	2126 non-null	float64
10	mean_value_of_long_term_variability	2126 non-null	float64
11	histogram_width	2126 non-null	float64
12	histogram_min	2126 non-null	float64
13	histogram_max	2126 non-null	float64
14	histogram_number_of_peaks	2126 non-null	float64
15	histogram_number_of_zeroes	2126 non-null	float64
16	histogram_mode	2126 non-null	float64
17	histogram_mean	2126 non-null	float64
18	histogram_median	2126 non-null	float64
19	histogram_variance	2126 non-null	float64
20	histogram_tendency	2126 non-null	float64
21	fetal_health	2126 non-null	float64

```
dtypes: float64(22)
```

```
memory usage: 365.5 KB
```

```
In [5]: df.isnull().sum()
```

```
Out[5]: baseline value      0
accelerations      0
fetal_movement     0
uterine_contractions  0
light_decelerations  0
severe_decelerations  0
prolongued_decelerations  0
abnormal_short_term_variability  0
mean_value_of_short_term_variability  0
percentage_of_time_with_abnormal_long_term_variability  0
mean_value_of_long_term_variability  0
histogram_width     0
histogram_min       0
histogram_max       0
histogram_number_of_peaks  0
histogram_number_of_zeroes  0
histogram_mode      0
histogram_mean      0
histogram_median    0
histogram_variance  0
histogram_tendency  0
fetal_health        0
dtype: int64
```

Exploratory Data Analysis (EDA)

• Univariate Analysis

```
In [6]: # Bar plot for fetal_health value counts
```

```
bar_fig = go.Bar(
    x=(df['fetal_health']-1).value_counts().index,
    y=(df['fetal_health']-1).value_counts().values,
    marker=dict(color='#66C2A5')
)
```

```

# Pie chart for fetal_health value counts
pie_fig = go.Pie(
    labels=(df['fetal_health']-1).value_counts().index,
    values=(df['fetal_health']-1).value_counts().values,
    hole=0.3 # Optional: for a donut-style pie chart
)

# Create subplots with Plotly
fig = sp.make_subplots(
    rows=1, cols=2,
    subplot_titles=("Bar Plot", "Pie Chart"),
    specs=[[{"type": "bar"}, {"type": "pie"}]]
)

# Add the bar plot to the first subplot
fig.add_trace(bar_fig, row=1, col=1)

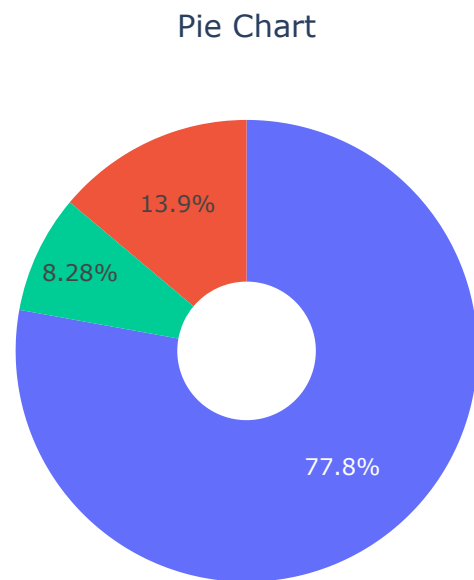
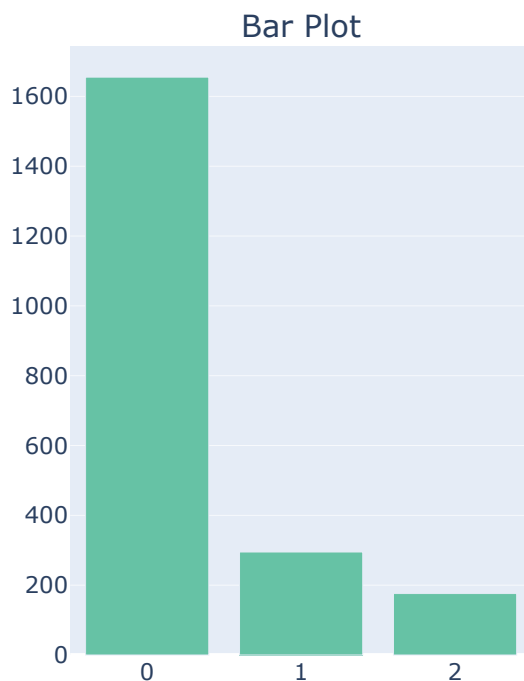
# Add the pie chart to the second subplot
fig.add_trace(pie_fig, row=1, col=2)

# Update layout for better visualization
fig.update_layout(
    title_text="Fetal Health Distribution",
    showlegend=False, # Hide legend to avoid redundancy in pie chart
    height=500,
    width=700
)

# Show the figure
fig.show()

```

Fetal Health Distribution



• Visualizing Data Distributions with `histplot`

In [7]: `import warnings`
`import seaborn as sns`

```

import matplotlib.pyplot as plt

# Suppress warnings
warnings.filterwarnings('ignore')

# Set seaborn style
plt.style.use('seaborn-v0_8-whitegrid')

# Drop the target variable from the features
X = df.drop('fetal_health', axis=1)

# Create the figure
plt.figure(figsize=(20, 50))
num_row = 1

# Loop through each column to create the subplots
for col in X.columns:
    plt.subplot(11, 2, num_row)

    # Set title with larger font size
    plt.title(f"Distribution of {col} Data", fontsize=22)

    # Plot histogram with KDE and hue
    sns.histplot(x=df[col], kde=True, hue=df['fetal_health'], palette='bright')

    # Set x and y axis labels with larger font size
    plt.xlabel(col, fontsize=20)
    plt.ylabel('Frequency', fontsize=20)

    # Set larger font size for ticks
    plt.xticks(fontsize=20)
    plt.yticks(fontsize=20)

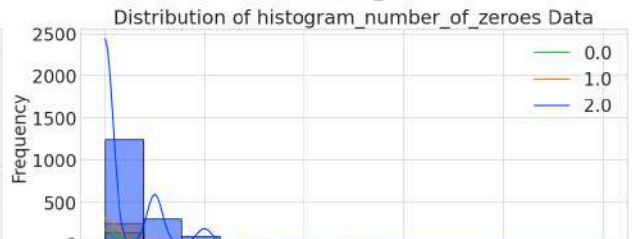
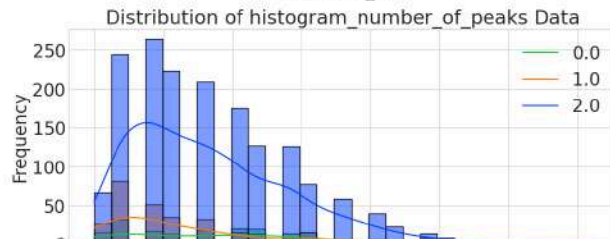
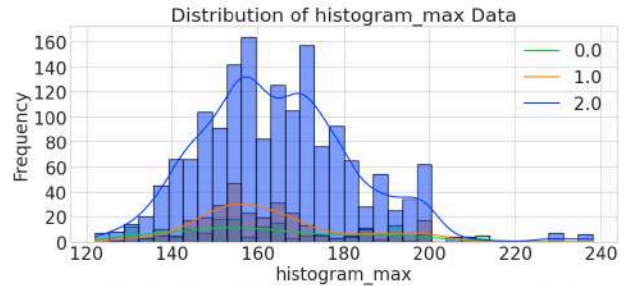
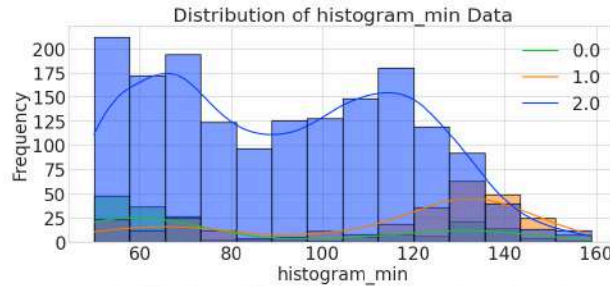
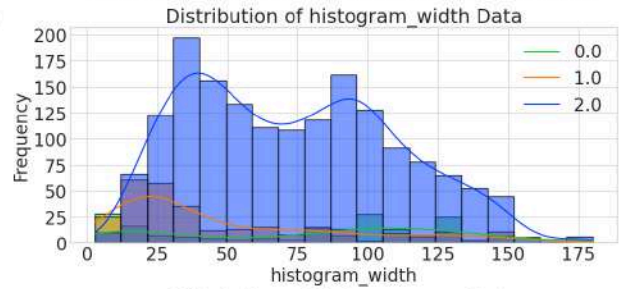
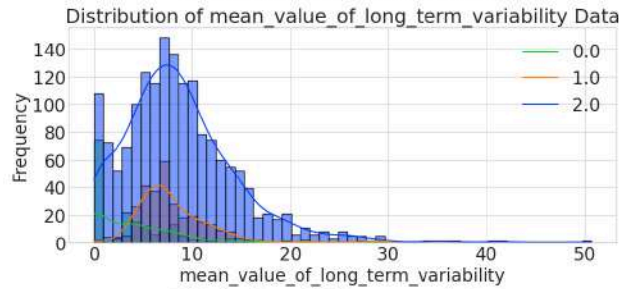
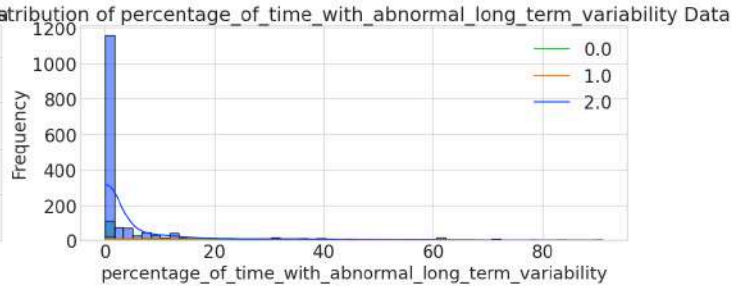
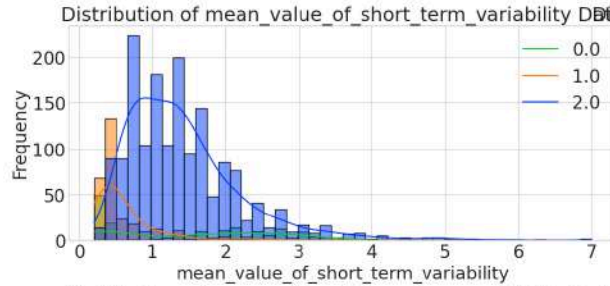
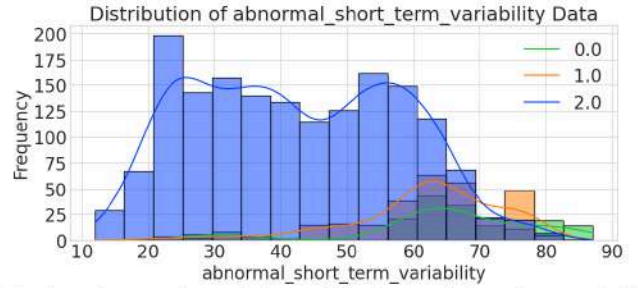
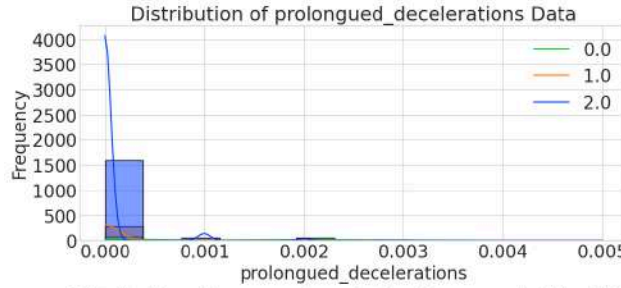
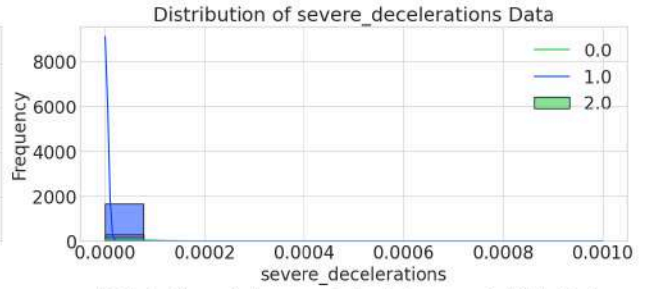
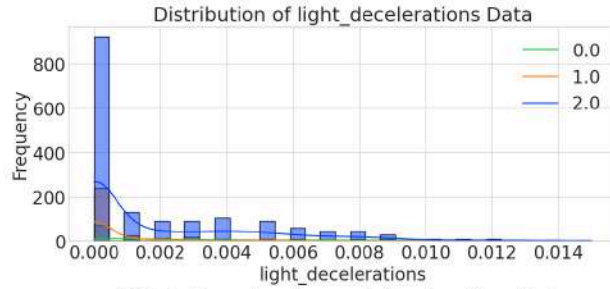
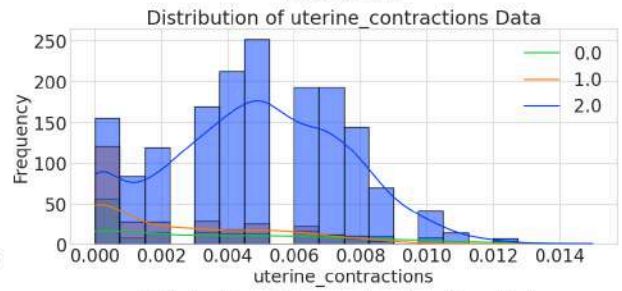
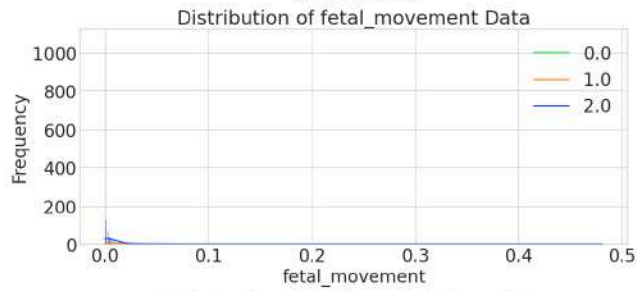
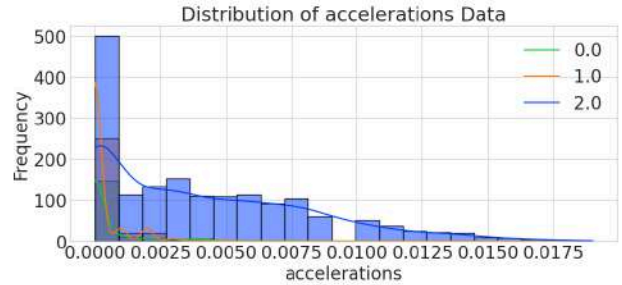
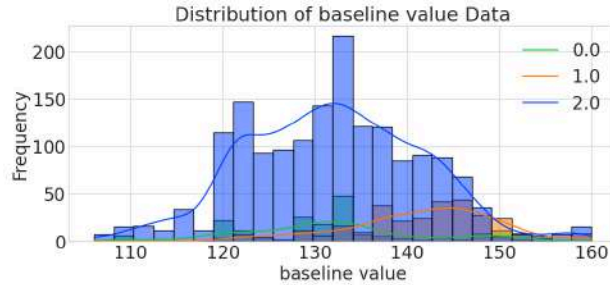
    # Set larger font size for the Legend
    plt.legend(df['fetal_health'].value_counts().index-1, title_fontsize='20', fontsize='20')
    #plt.legend(title='fetal_health', title_fontsize='15', fontsize='15')
    # Adjust layout
    plt.tight_layout()

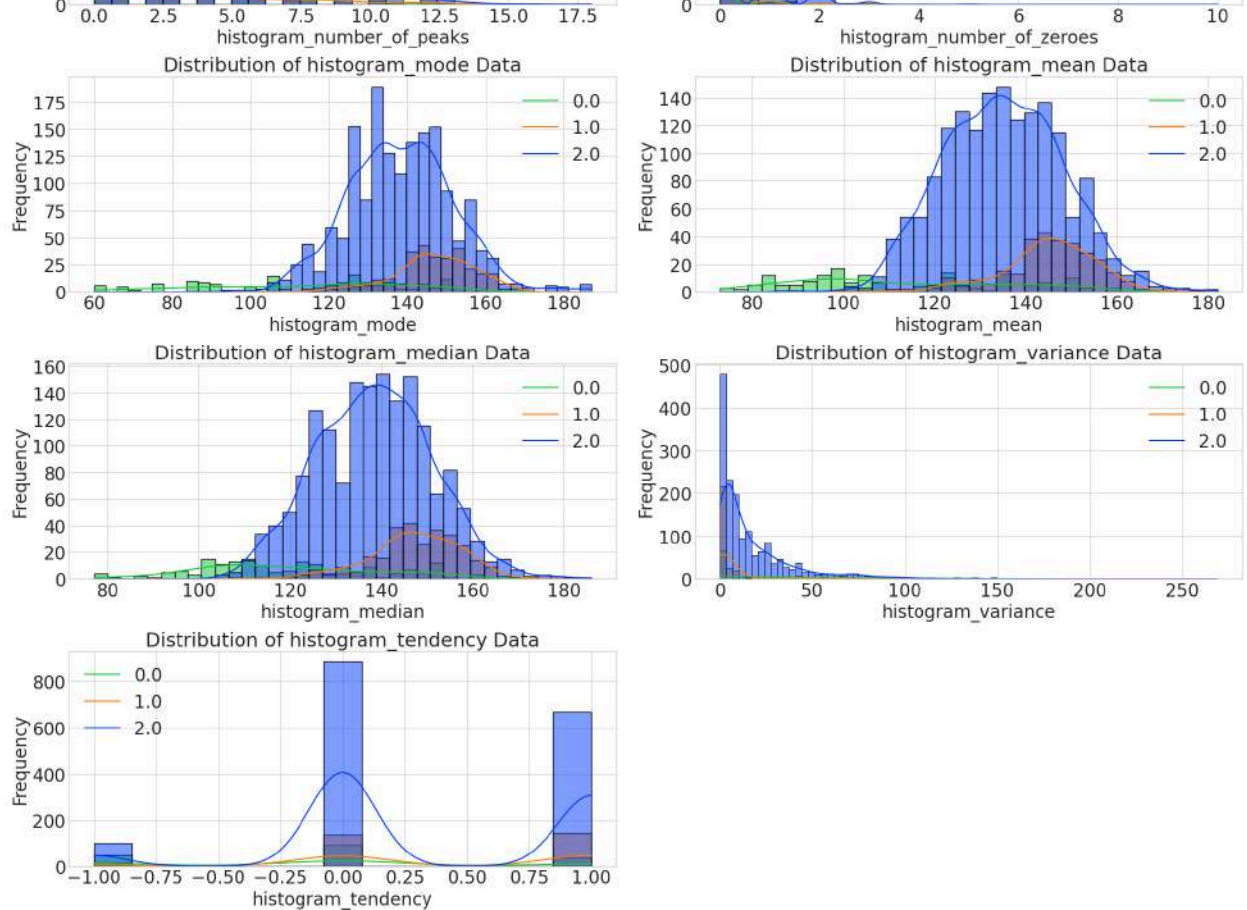
    num_row += 1

# Show the plots

plt.show()

```



• Visualizing Data Distributions with Box-Plot

```
In [8]: import warnings
import seaborn as sns
import matplotlib.pyplot as plt

# Suppress warnings
warnings.filterwarnings('ignore')

# Set seaborn style
plt.style.use('seaborn-v0_8-whitegrid')

# Drop the target variable from the features
X = df.drop('fetal_health', axis=1)

# Create the figure
plt.figure(figsize=(20, 50))
num_row = 1

# Loop through each column to create the subplots
for col in X.columns:
    plt.subplot(11, 2, num_row)

    # Set title with larger font size
    plt.title(f"Distribution of {col} Data", fontsize=22)

    # Plot histogram with KDE and hue
    sns.boxplot(y=df[col], x=df['fetal_health'], hue=df['fetal_health'])

    # Set x and y axis labels with larger font size
    plt.xlabel(col, fontsize=20)
    plt.ylabel('Frequency', fontsize=20)

    # Set larger font size for ticks
    plt.xticks(fontsize=20)
    plt.yticks(fontsize=20)

    # Set larger font size for the legend
```

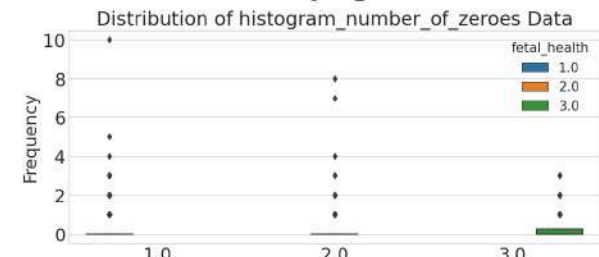
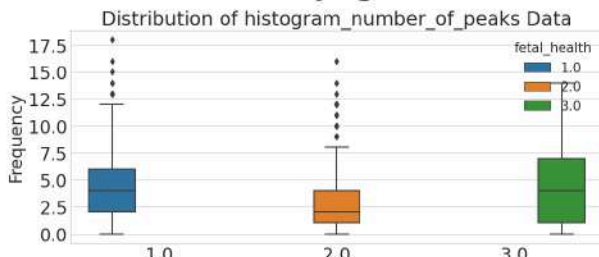
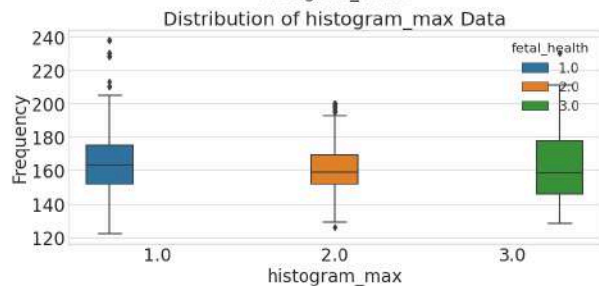
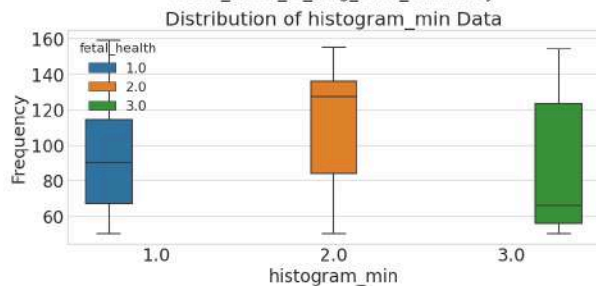
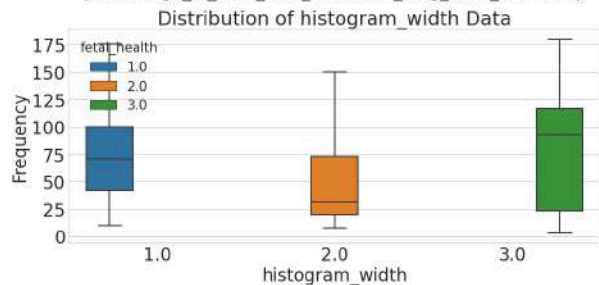
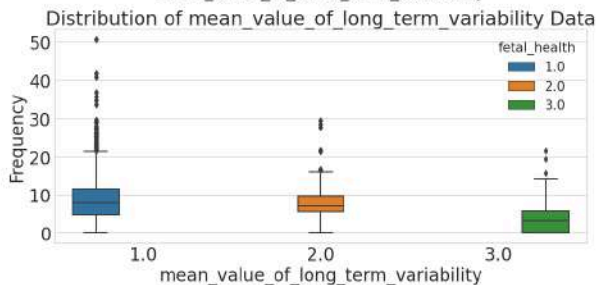
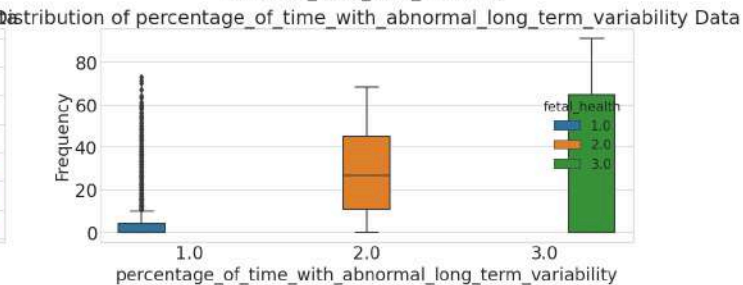
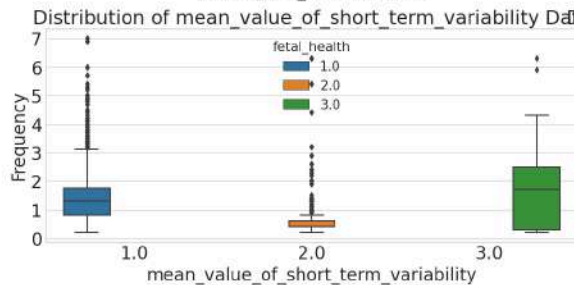
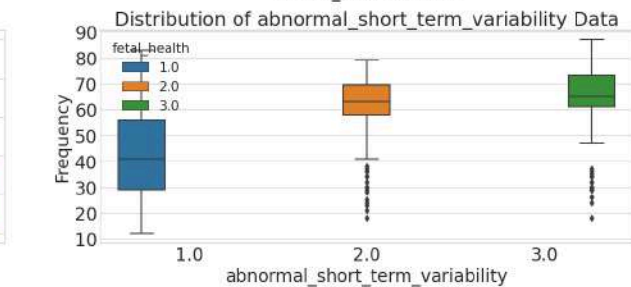
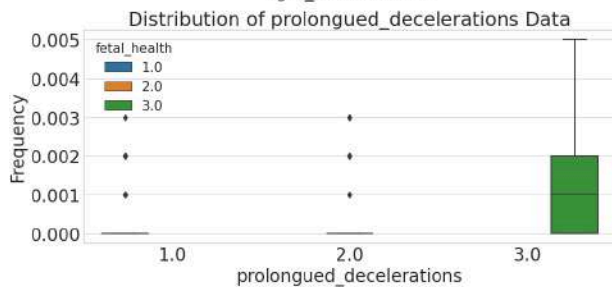
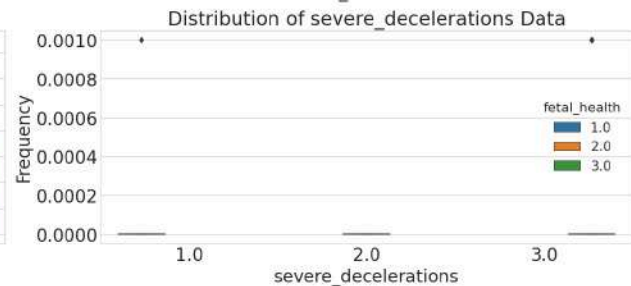
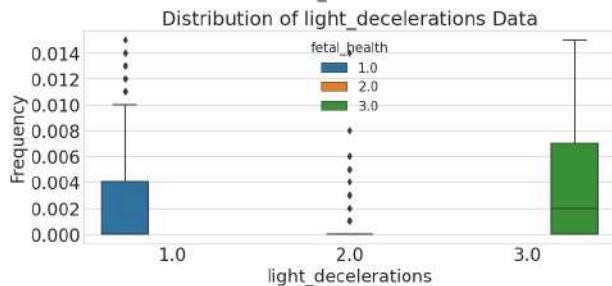
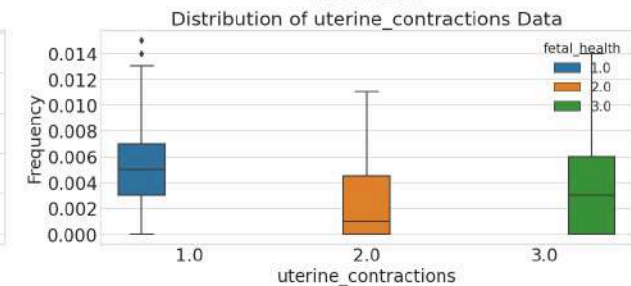
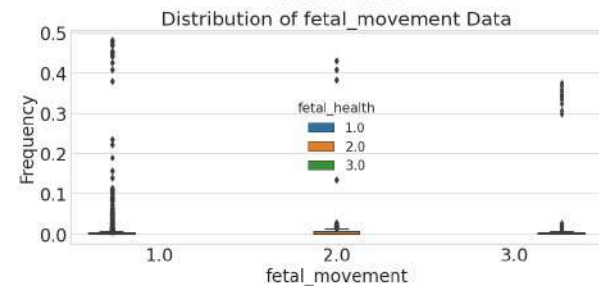
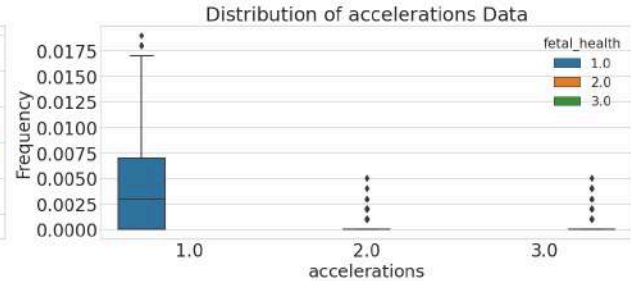
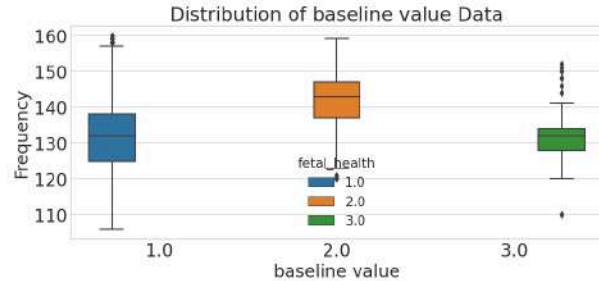


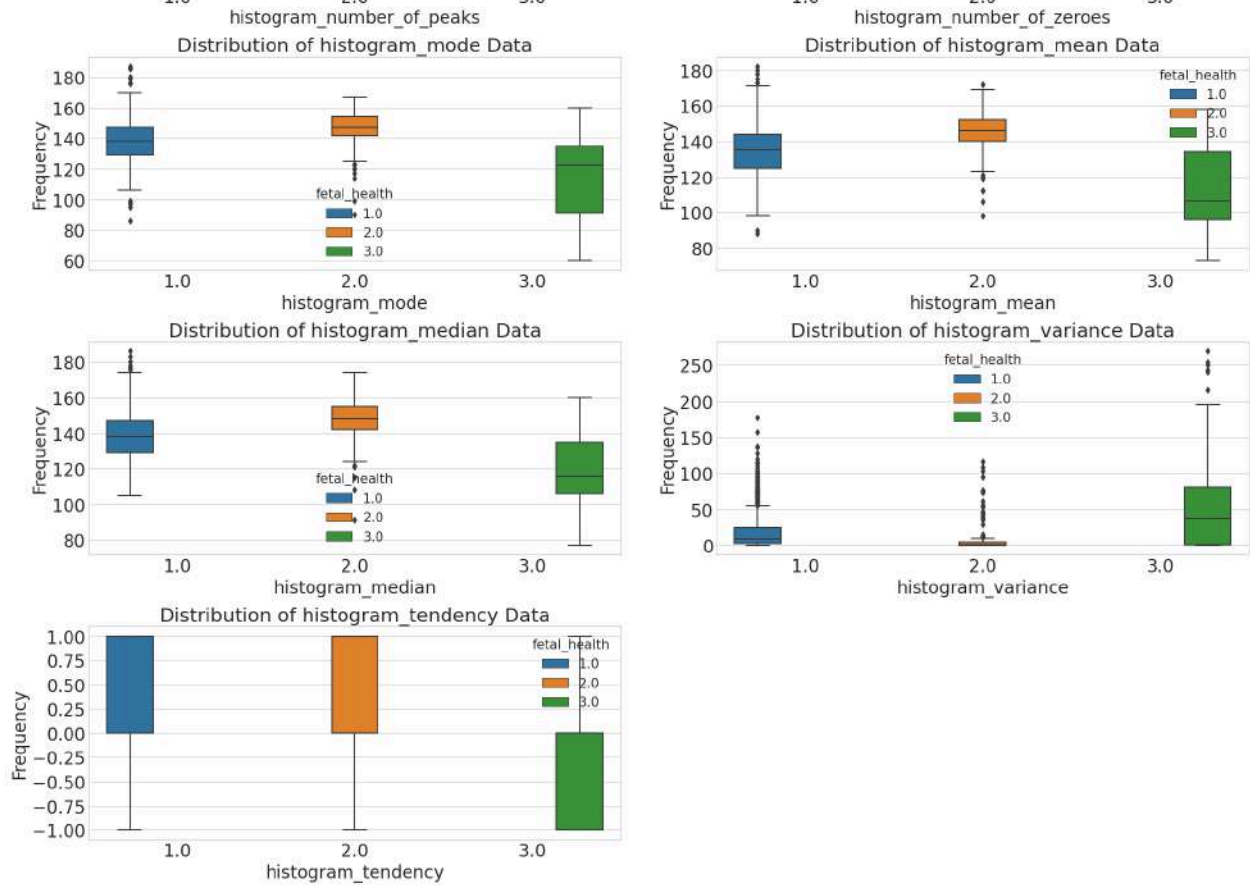
```
plt.legend(title='fetal_health', title_fontsize='15', fontsize='15')

# Adjust Layout
plt.tight_layout()

num_row += 1

# Show the plots
plt.show()
```

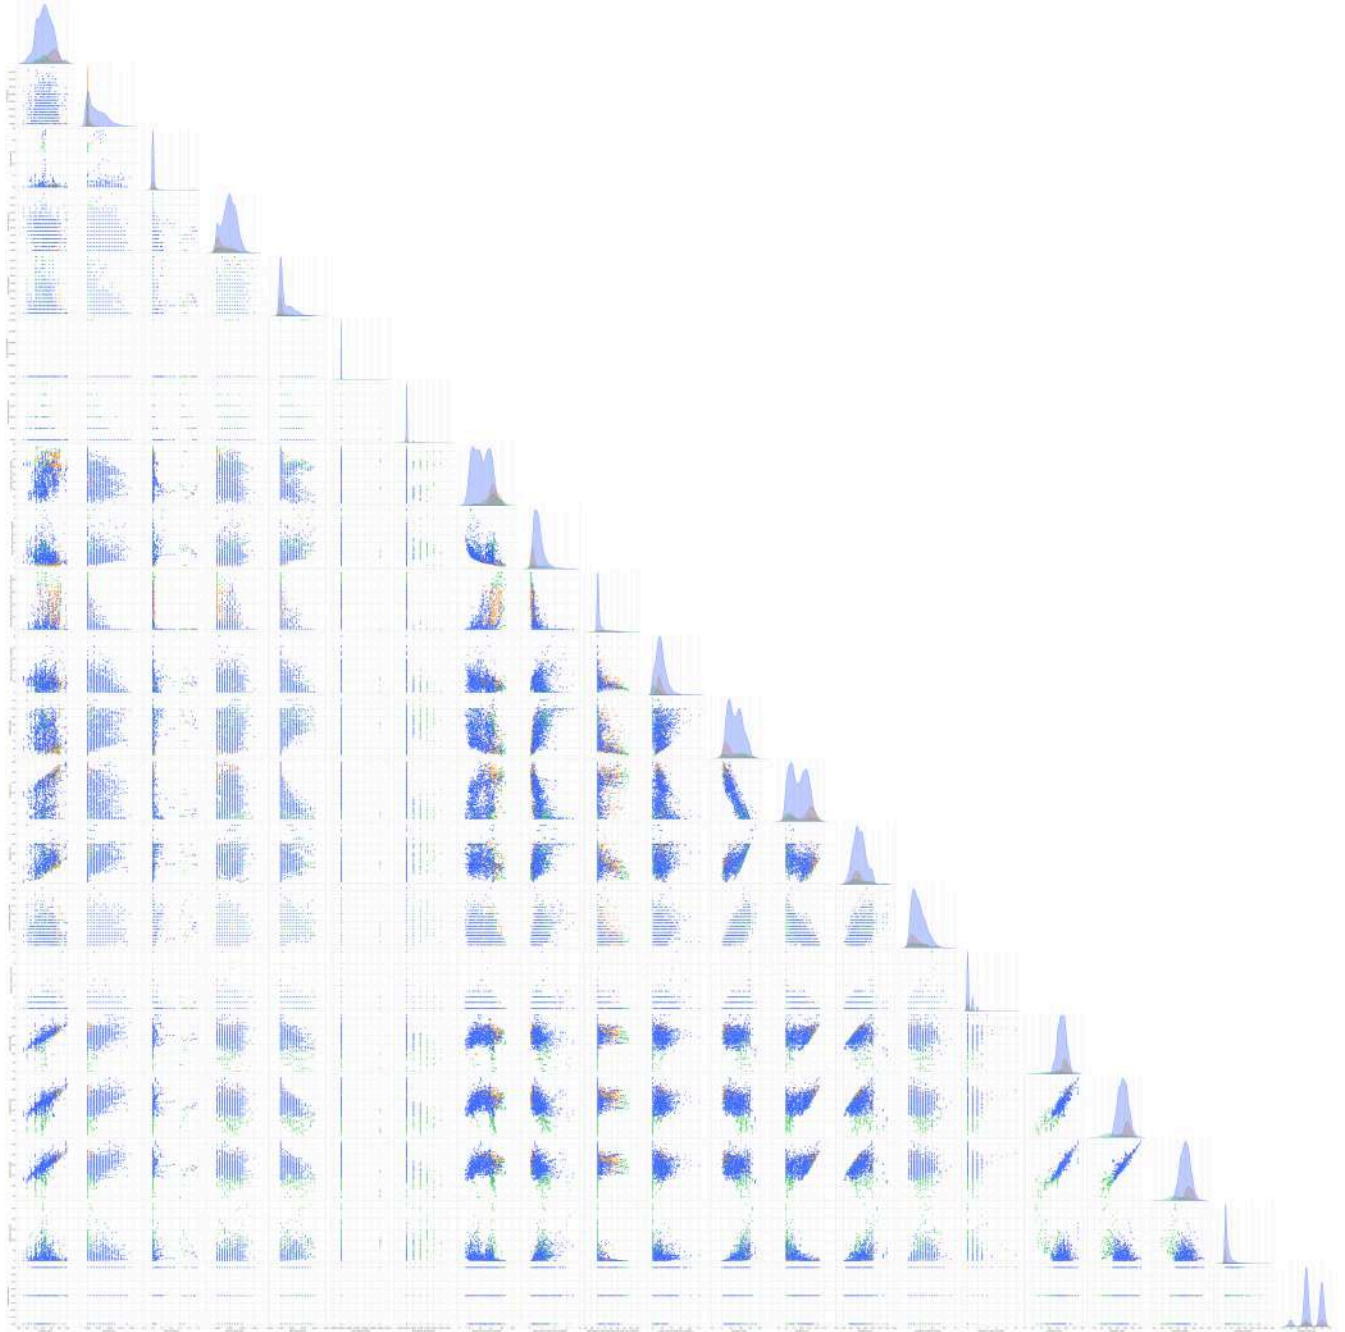




Bivariate Analysis with pairplot

```
In [9]: sns.pairplot(df, hue='fetal_health', corner=True, palette='bright', height=5)

plt.show()
```



Bivariate Analysis with Heatmap

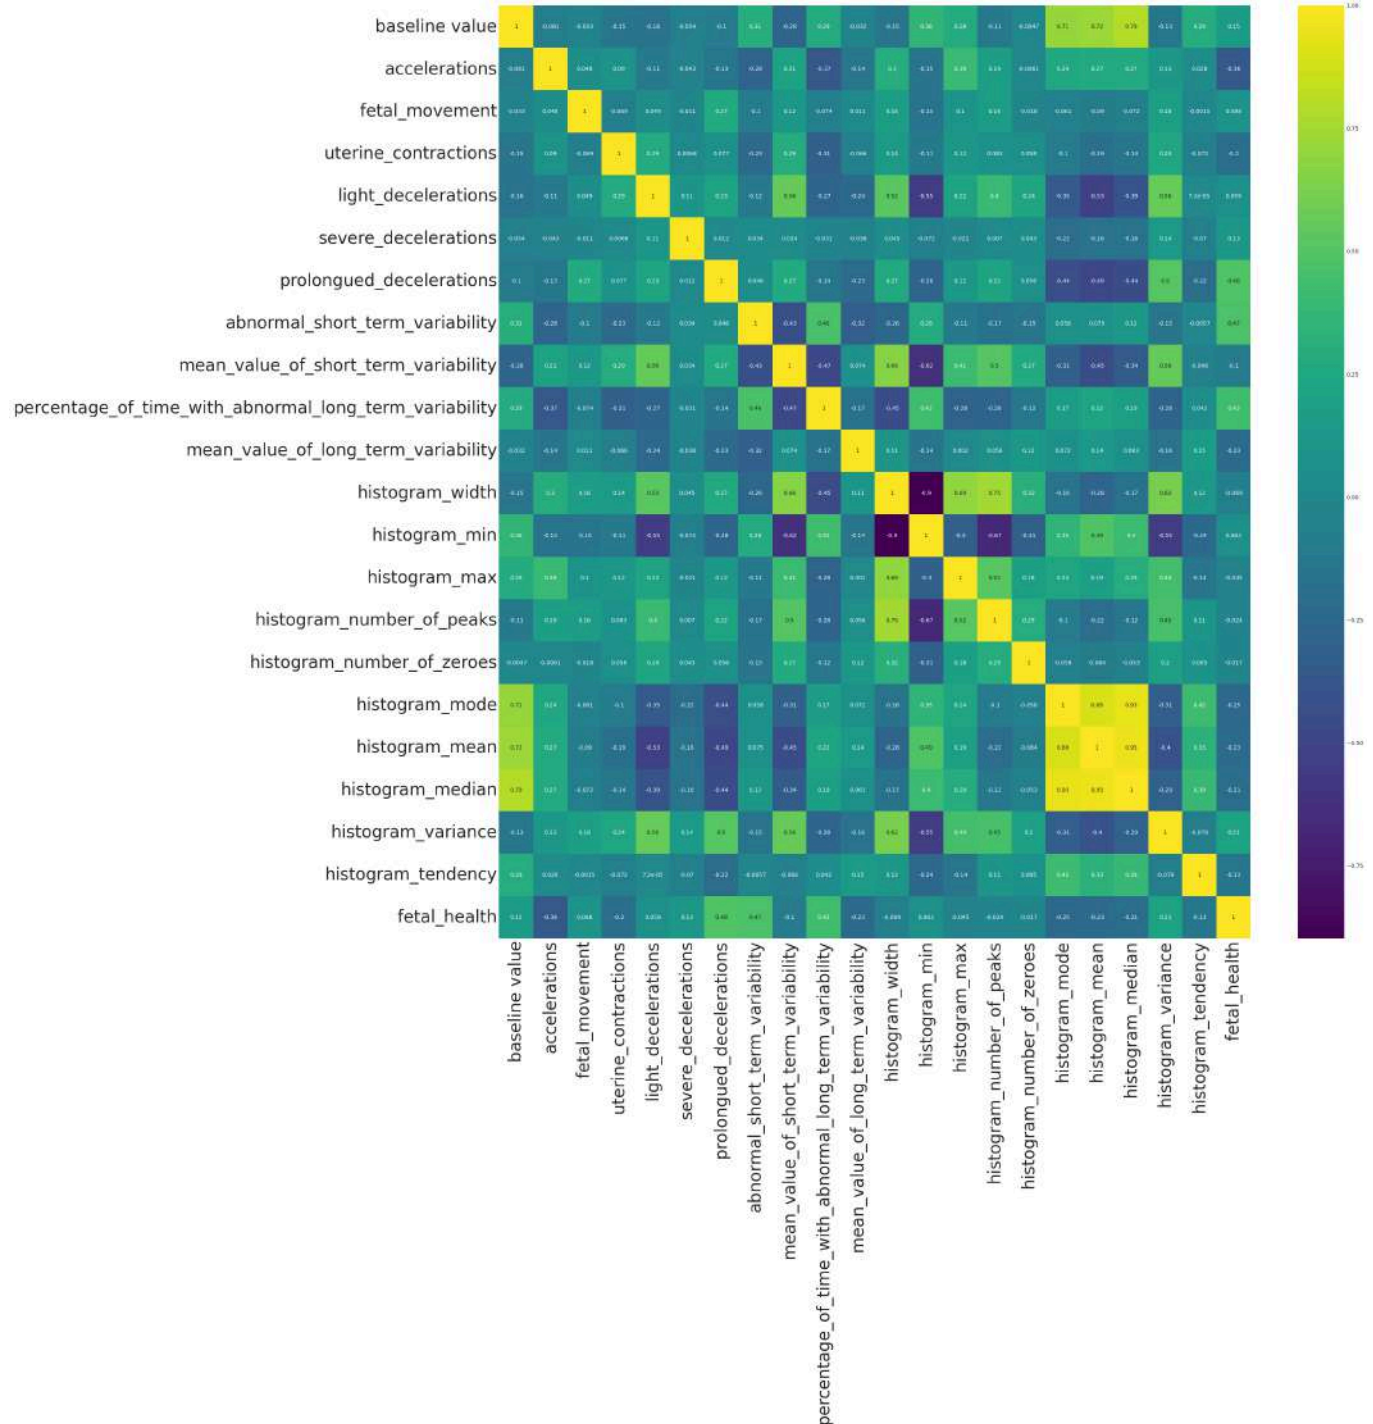
```
In [10]: import seaborn as sns
import matplotlib.pyplot as plt

# Create the figure
plt.figure(figsize=(30, 30))

# Plot the heatmap
sns.heatmap(df.corr(), annot=True, cmap='viridis')

plt.xticks(fontsize=30)
plt.yticks(fontsize=30)

# Show the plot
plt.show()
```



Model Training and Prediction

```
In [11]: # Splitting data into features and target
X = df.drop(columns=['fetal_health'])
y = df['fetal_health']
y=y-1
# Scaling the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

```
In [12]: def plot_multi_class_roc(model, X_test, y_test, num_classes=3):
# Binarize the labels for multi-class ROC calculation
y_test_bin = label_binarize(y_test, classes=[0, 1, 2])
y_pred_proba = model.predict_proba(X_test)

fpr = dict()
tpr = dict()
roc_auc = dict()
```



```

# Compute ROC curve and ROC area for each class
for i in range(num_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_pred_proba[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot ROC curve for each class
plt.figure(figsize=(8, 6))
colors = cycle(['blue', 'red', 'green'])

for i, color in zip(range(num_classes), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=2, label=f'Class {i+1} (AUC = {roc_auc[i]:.4f})')

# Plot the diagonal (random classifier)
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Multi-Class ROC Curve')
plt.legend(loc="lower right")
plt.show()

```

• Logistic Regression

```

In [13]: # Logistic Regression
log_model = LogisticRegression(random_state=42)
log_model.fit(X_train, y_train)
log_y_pred = log_model.predict(X_test)
log_accuracy = accuracy_score(y_test, log_y_pred)

# Logistic Regression - Classification Report & Confusion Matrix
print("Logistic Regression Classification Report")
print(classification_report(y_test, log_y_pred))
print("Logistic Regression Confusion Matrix")
print(confusion_matrix(y_test, log_y_pred))

# Calculate MCC for the model's predictions
mcc_log = matthews_corrcoef(y_test, log_y_pred)

# Print MCC
print(f"Matthews Correlation Coefficient (MCC): {mcc_log:.4f}")

```

```

Logistic Regression Classification Report
              precision    recall  f1-score   support

     0.0         0.94      0.93      0.94         333
     1.0         0.63      0.64      0.64          64
     2.0         0.73      0.76      0.75          29

 accuracy          0.88         0.88         0.88         426
 macro avg         0.77         0.78         0.77         426
weighted avg         0.88         0.88         0.88         426

```

```

Logistic Regression Confusion Matrix
[[311  19   3]
 [ 18  41   5]
 [   2   5  22]]
Matthews Correlation Coefficient (MCC): 0.6655

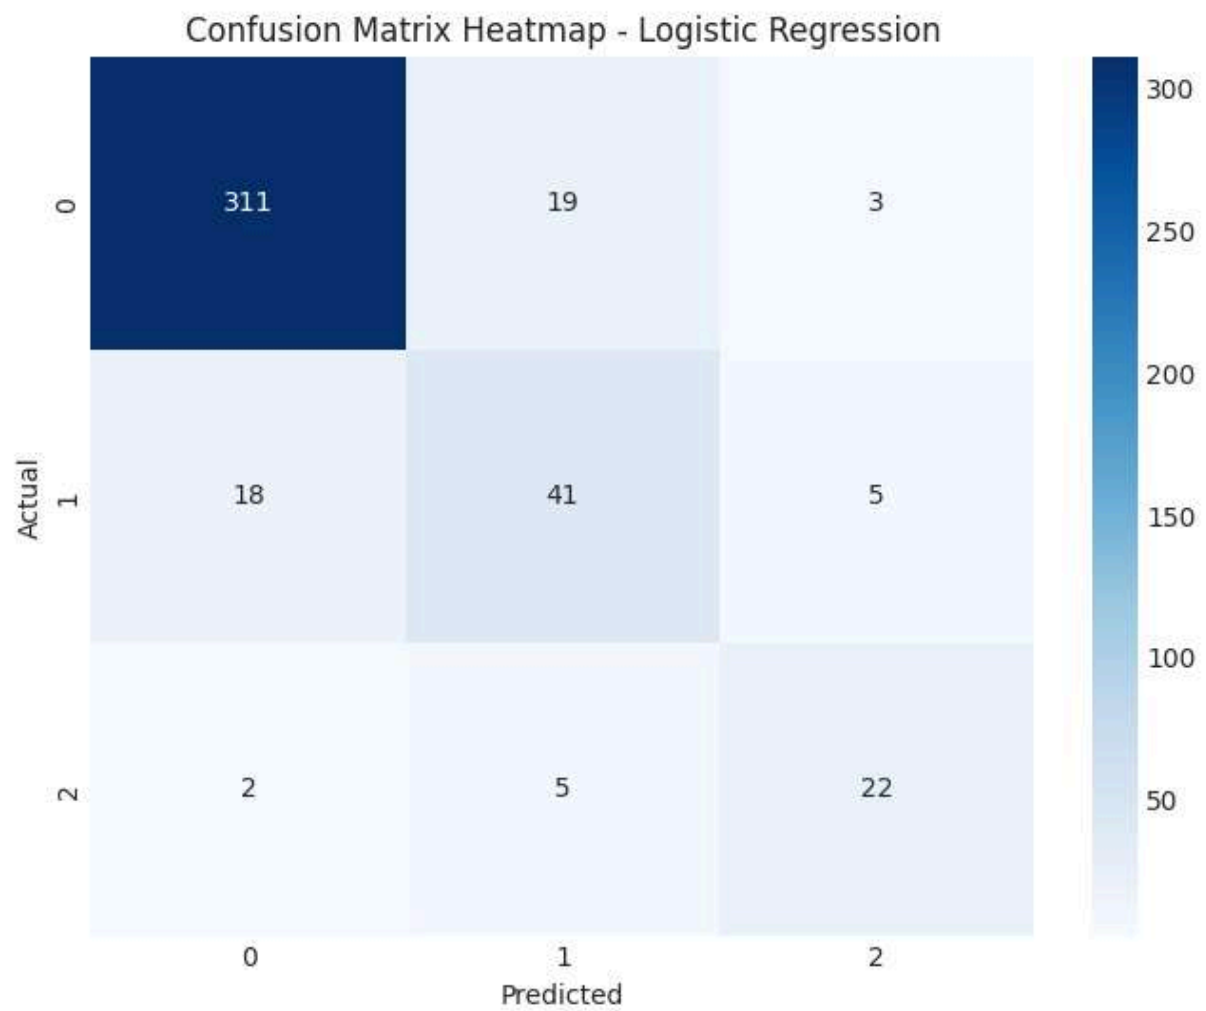
```

```

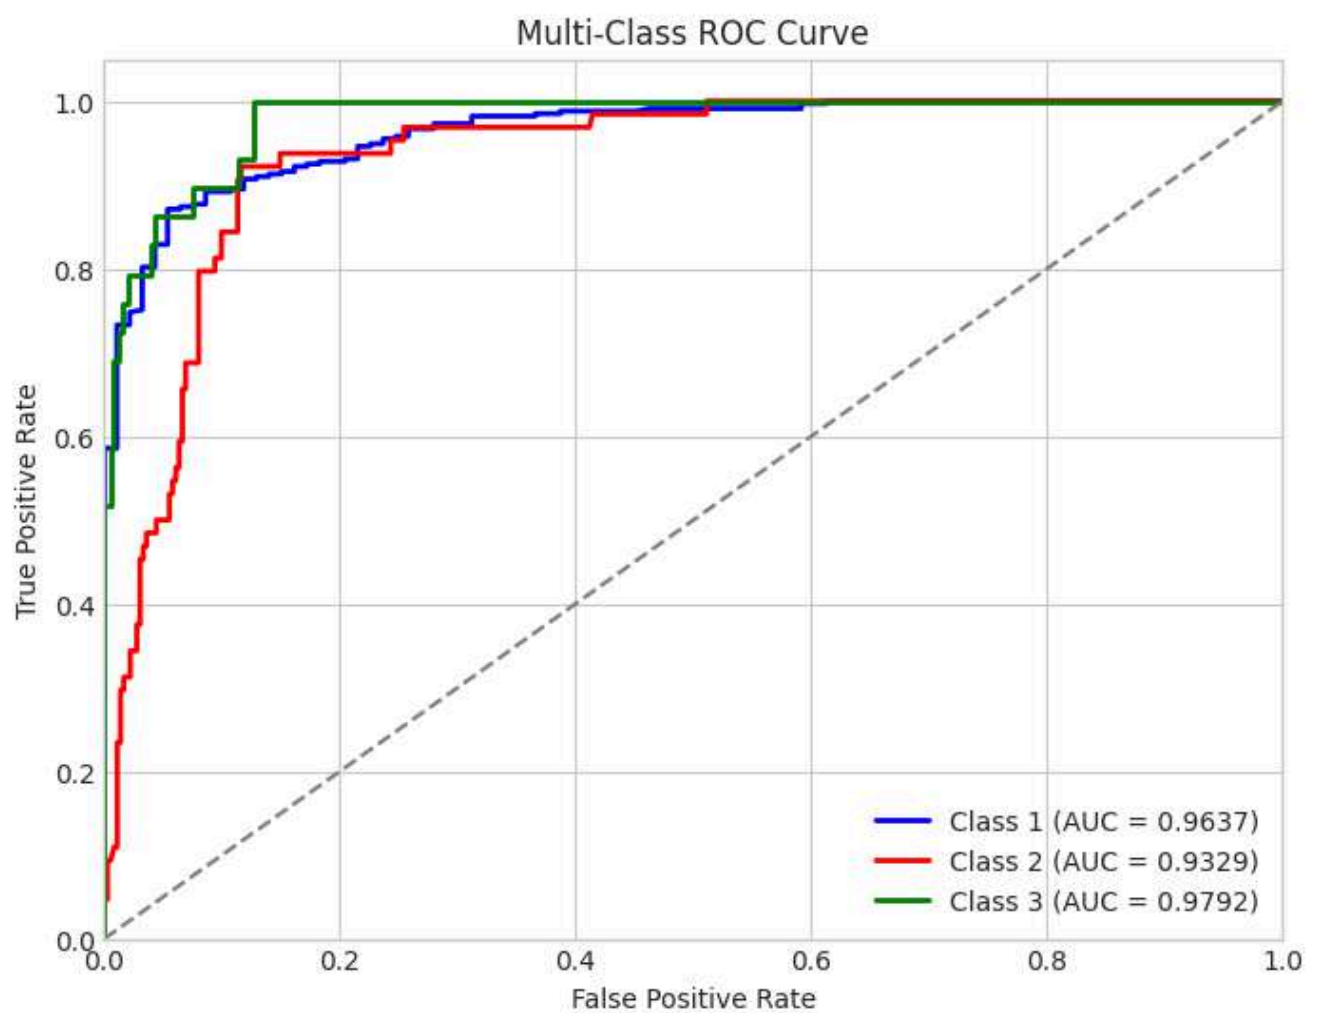
In [14]: def plot_confusion_matrix_heatmap(y_test, y_pred, model_name):
    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=[0, 1, 2], yticklabels=[0, 1, 2])
    plt.title(f'Confusion Matrix Heatmap - {model_name}')
    plt.xlabel('Predicted')
    plt.ylabel('Actual')

```

```
plt.show()
plot_confusion_matrix_heatmap(y_test, log_y_pred, 'Logistic Regression')
```



```
In [15]: plot_multi_class_roc(log_model, X_test, y_test, num_classes=3)
```



• Support Vector Machine (SVM)

```
In [16]: # Support Vector Machine (SVM)
svm_model = SVC(random_state=42, probability=True)
svm_model.fit(X_train, y_train)
svm_y_pred = svm_model.predict(X_test)
svm_accuracy = accuracy_score(y_test, svm_y_pred)

# Support Vector Machine (SVM) - Classification Report & Confusion Matrix
print(svm_accuracy)
print("SVM Classification Report")
print(classification_report(y_test, svm_y_pred))
print("SVM Confusion Matrix")
print(confusion_matrix(y_test, svm_y_pred))

# Calculate MCC for the model's predictions
mcc_svm = matthews_corrcoef(y_test, svm_y_pred)

# Print MCC
print(f"Matthews Correlation Coefficient (MCC): {mcc_svm:.4f}")
```

0.903755868544601

SVM Classification Report

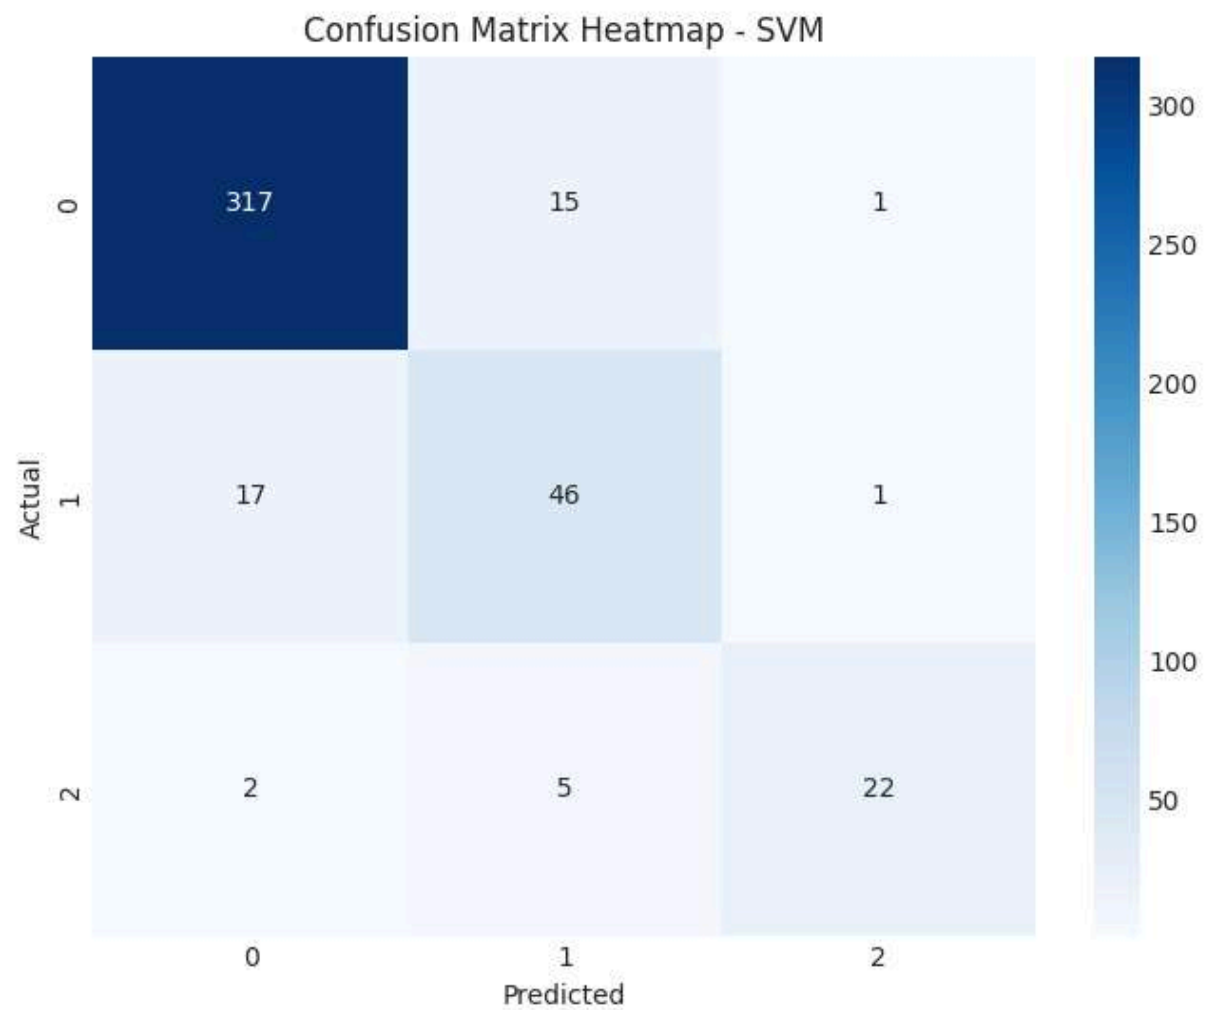
	precision	recall	f1-score	support
0.0	0.94	0.95	0.95	333
1.0	0.70	0.72	0.71	64
2.0	0.92	0.76	0.83	29
accuracy			0.90	426
macro avg	0.85	0.81	0.83	426
weighted avg	0.90	0.90	0.90	426

SVM Confusion Matrix

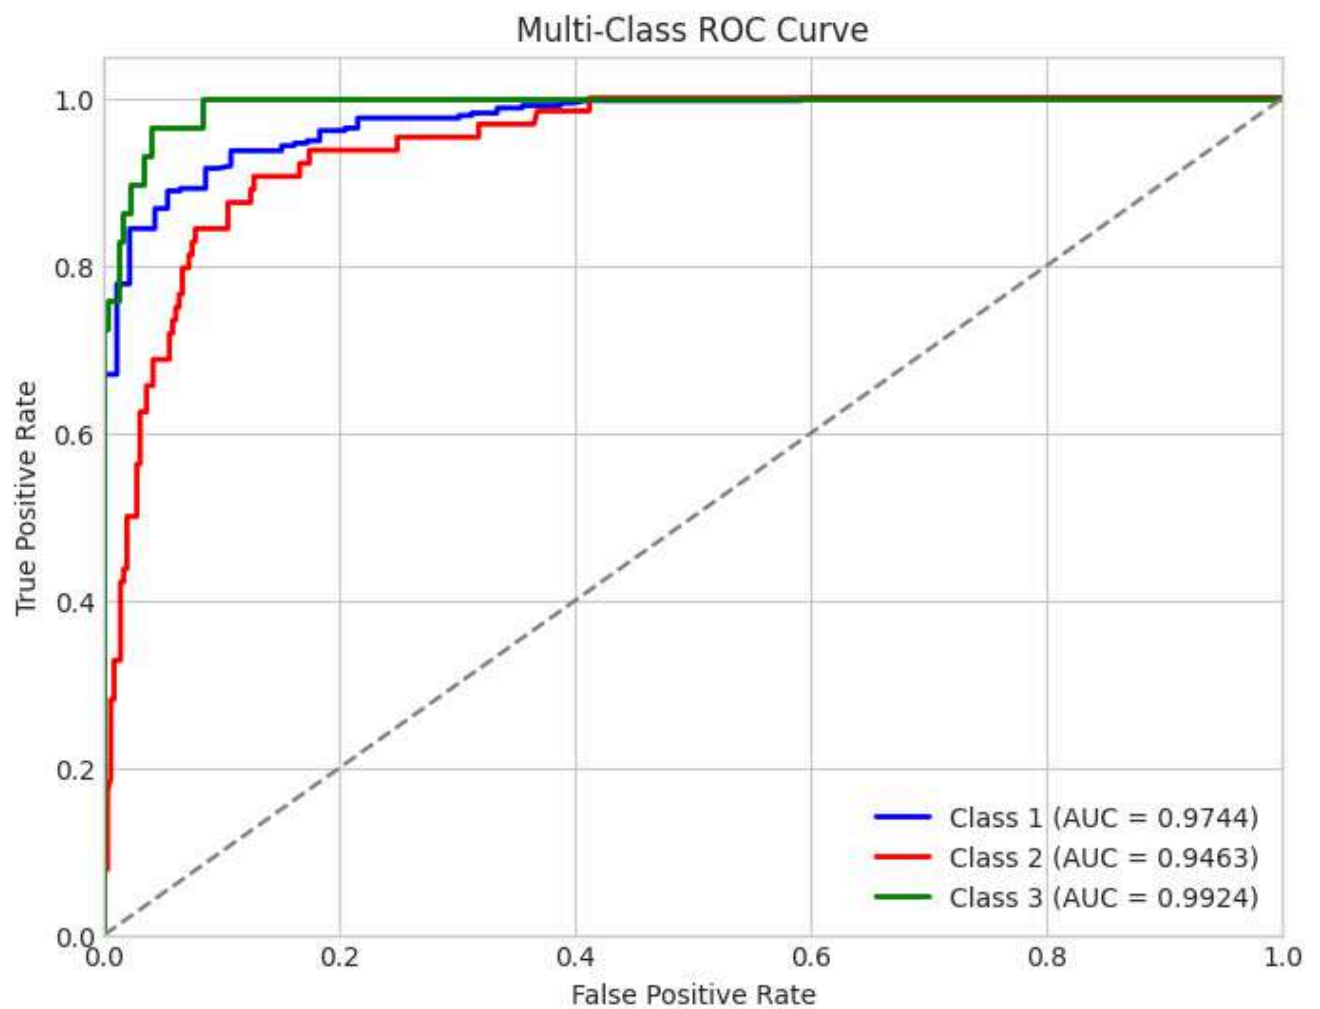
```
[[317  15   1]
 [ 17  46   1]
 [   2   5  22]]
```

Matthews Correlation Coefficient (MCC): 0.7302

```
In [17]: plot_confusion_matrix_heatmap(y_test, svm_y_pred, 'SVM')
```



```
In [18]: plot_multi_class_roc(svm_model ,X_test, y_test, num_classes=3)
```



• Random Forest

```
In [19]: # Random Forest
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)
rf_y_pred = rf_model.predict(X_test)
rf_accuracy = accuracy_score(y_test, rf_y_pred)

# Random Forest - Classification Report & Confusion Matrix
print("Random Forest Classification Report")
print(classification_report(y_test, rf_y_pred))
print("Random Forest Confusion Matrix")
print(confusion_matrix(y_test, rf_y_pred))

# Calculate MCC for the model's predictions
mcc_rf = matthews_corrcoef(y_test, rf_y_pred)

# Print MCC
print(f"Matthews Correlation Coefficient (MCC): {mcc_rf:.4f}")
```

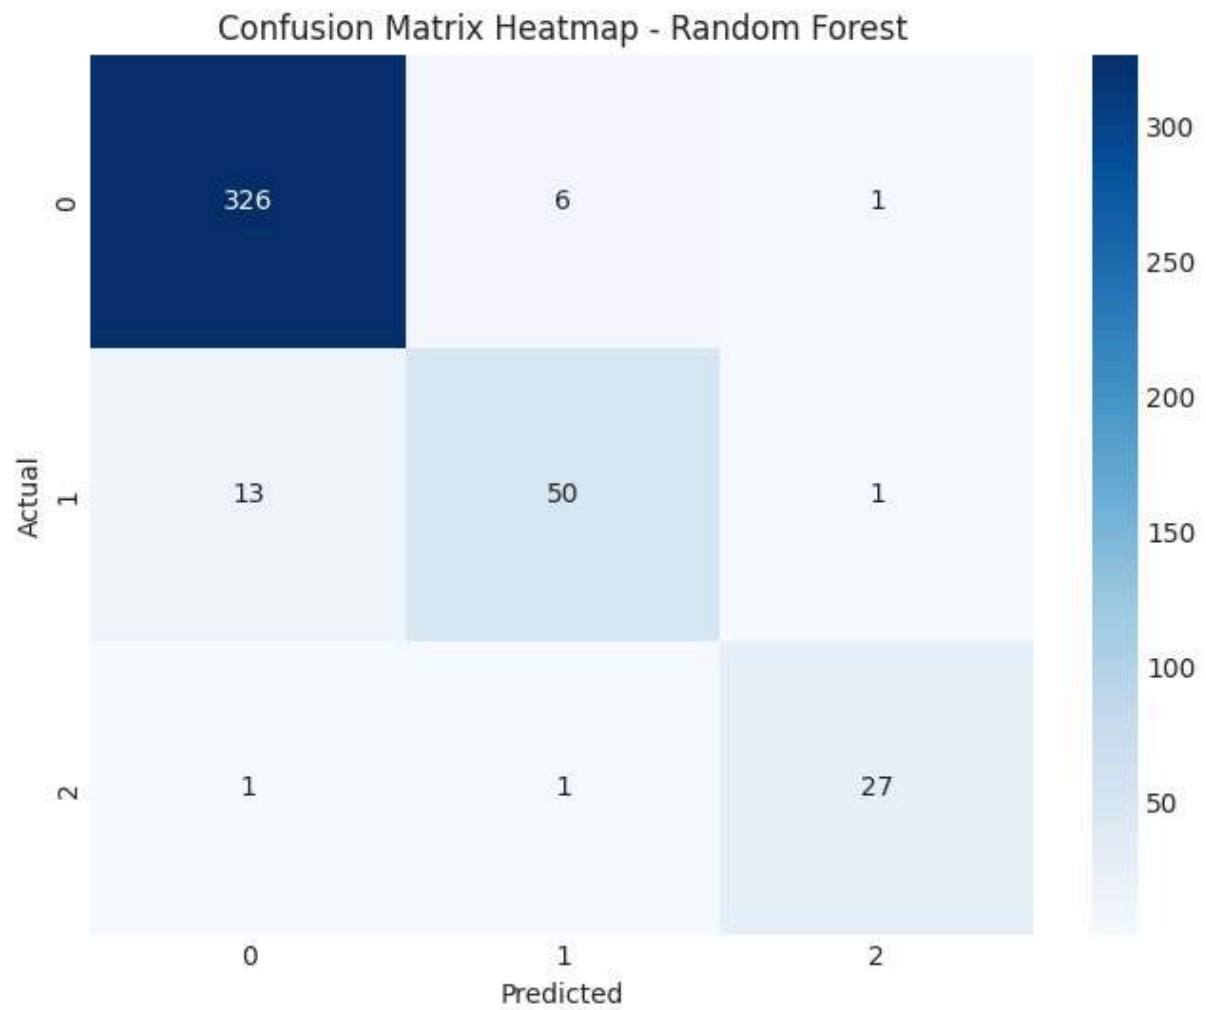

Random Forest Classification Report				
	precision	recall	f1-score	support
0.0	0.96	0.98	0.97	333
1.0	0.88	0.78	0.83	64
2.0	0.93	0.93	0.93	29
accuracy			0.95	426
macro avg			0.92	426
weighted avg			0.94	426

Random Forest Confusion Matrix

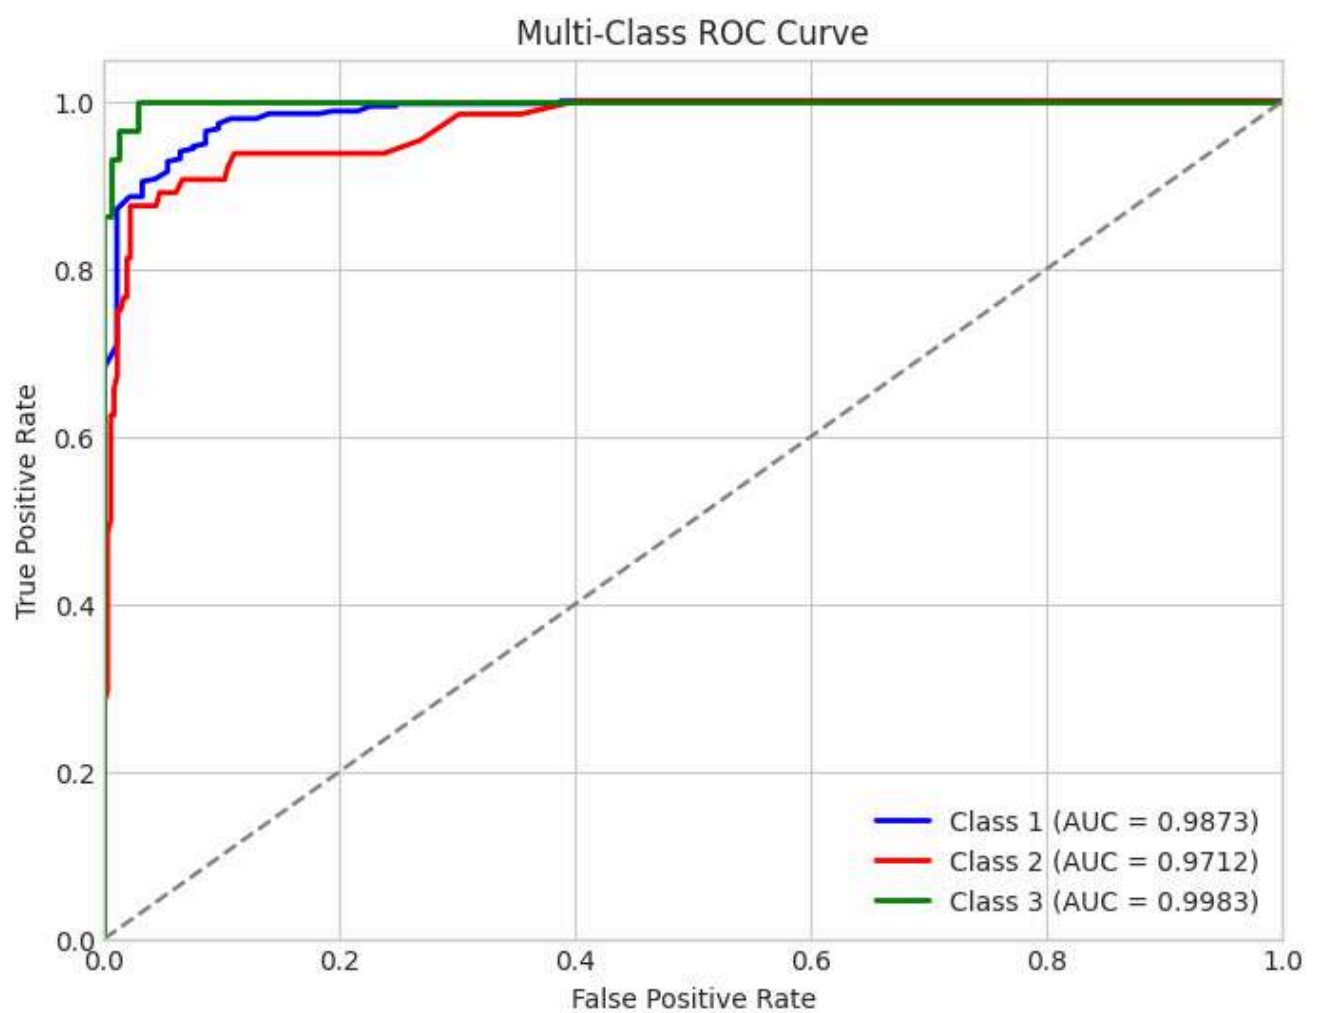
```
[[326  6  1]
 [ 13 50  1]
 [  1  1 27]]
```

Matthews Correlation Coefficient (MCC): 0.8474

```
In [20]: plot_confusion_matrix_heatmap(y_test, rf_y_pred, 'Random Forest')
```



```
In [21]: plot_multi_class_roc(rf_model ,X_test, y_test, num_classes=3)
```



• K-Nearest Neighbors (KNN)

```
In [22]: # K-Nearest Neighbors (KNN)
knn_model = KNeighborsClassifier()
knn_model.fit(X_train, y_train)
knn_y_pred = knn_model.predict(X_test)
knn_accuracy = accuracy_score(y_test, knn_y_pred)

# K-Nearest Neighbors (KNN) - Classification Report & Confusion Matrix
print("KNN Classification Report")
print(classification_report(y_test, knn_y_pred))
print("KNN Confusion Matrix")
print(confusion_matrix(y_test, knn_y_pred))

# Calculate MCC for the model's predictions
mcc_knn = matthews_corrcoef(y_test, knn_y_pred)

# Print MCC
print(f"Matthews Correlation Coefficient (MCC): {mcc_knn:.4f}")
```

```

KNN Classification Report
              precision    recall  f1-score   support

     0.0         0.94      0.97      0.96        333
     1.0         0.81      0.69      0.75         64
     2.0         0.79      0.76      0.77         29

 accuracy          0.92        426
 macro avg         0.85      0.81      0.82        426
 weighted avg      0.91      0.92      0.91        426

```

```

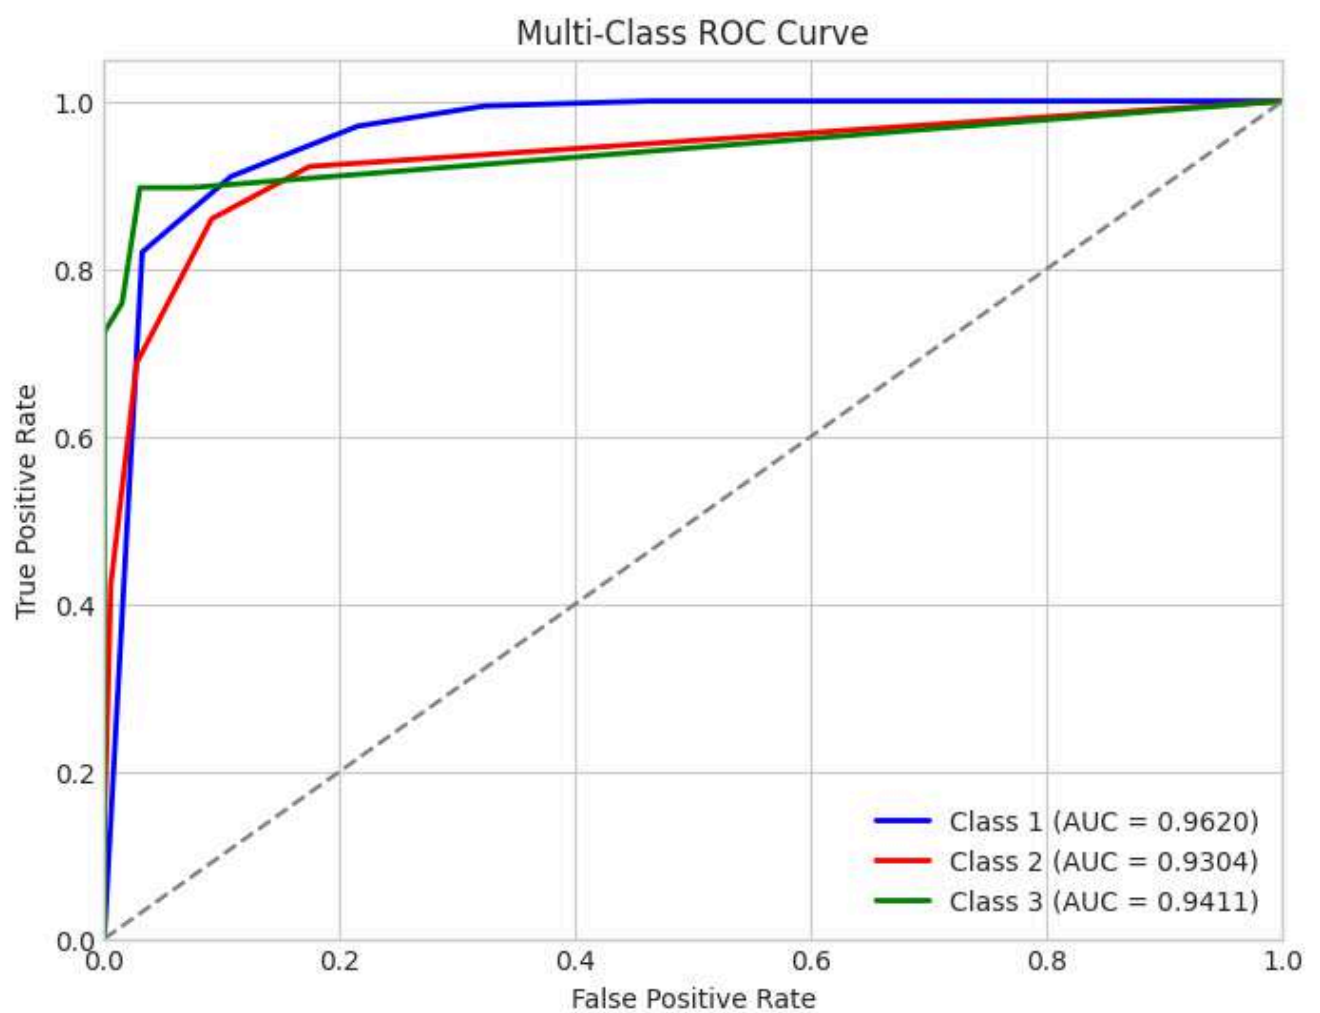
KNN Confusion Matrix
[[324  7  2]
 [ 16 44  4]
 [  4  3 22]]
Matthews Correlation Coefficient (MCC): 0.7575

```

```
In [23]: plot_confusion_matrix_heatmap(y_test, knn_y_pred, 'KNN')
```



```
In [24]: plot_multi_class_roc(knn_model ,X_test, y_test, num_classes=3)
```



• XGBoost

```
In [25]: # XGBoost
xgb_model = XGBClassifier(random_state=42, use_label_encoder=False, eval_metric='mlogloss')
xgb_model.fit(X_train, y_train)
xgb_y_pred = xgb_model.predict(X_test)
xgb_accuracy = accuracy_score(y_test, xgb_y_pred)

# XGBoost - Classification Report & Confusion Matrix

print("XGBoost Classification Report")
print(classification_report(y_test, xgb_y_pred))
print("XGBoost Confusion Matrix")
print(confusion_matrix(y_test, xgb_y_pred))

# Calculate MCC for the model's predictions
mcc_XGb = matthews_corrcoef(y_test, xgb_y_pred)

# Print MCC
print(f"Matthews Correlation Coefficient (MCC): {mcc_XGb:.4f}")
```

XGBoost Classification Report

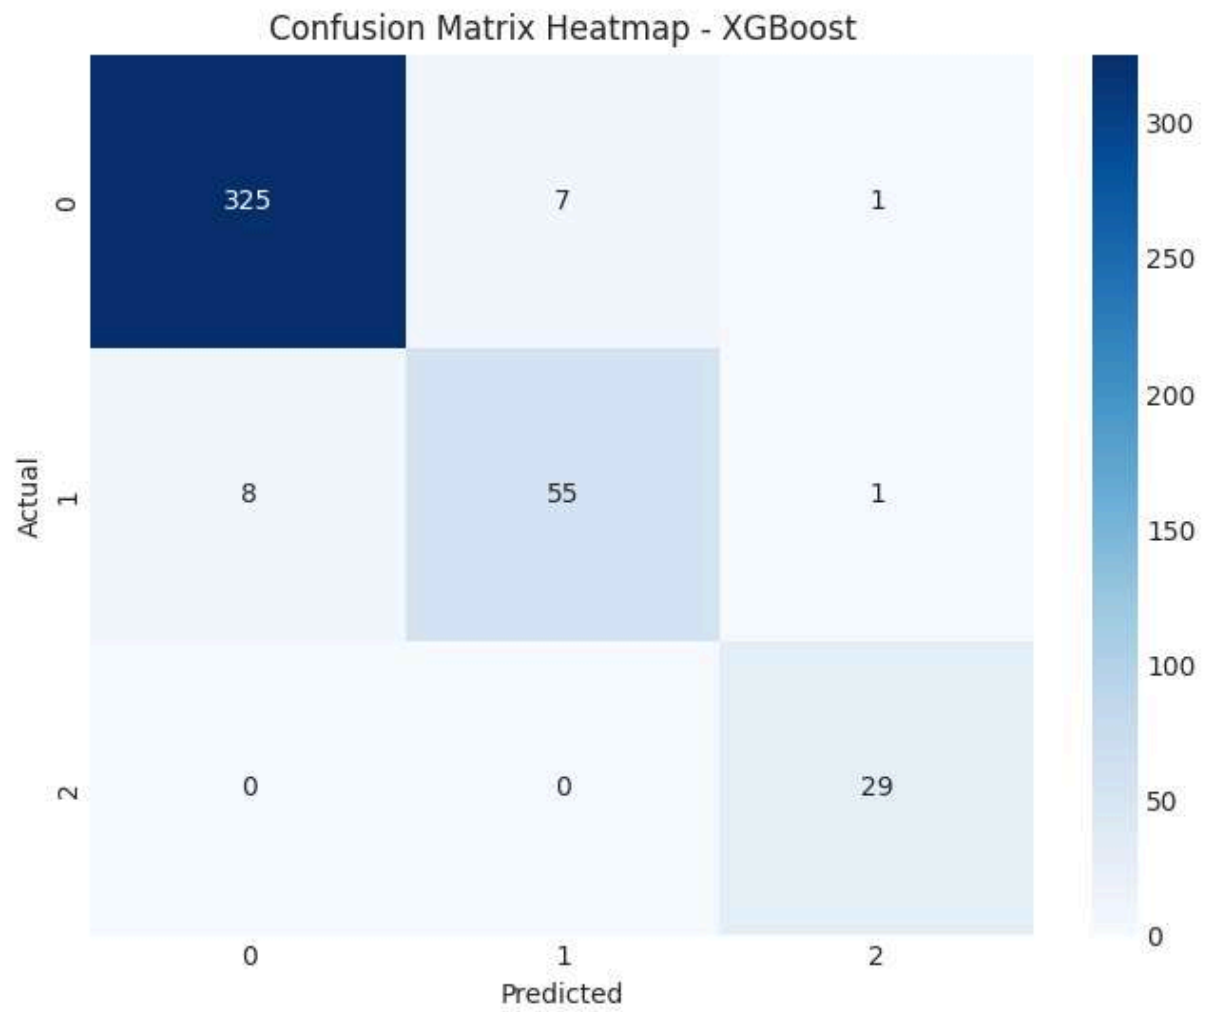
	precision	recall	f1-score	support
0.0	0.98	0.98	0.98	333
1.0	0.89	0.86	0.87	64
2.0	0.94	1.00	0.97	29
accuracy			0.96	426
macro avg	0.93	0.95	0.94	426
weighted avg	0.96	0.96	0.96	426

XGBoost Confusion Matrix

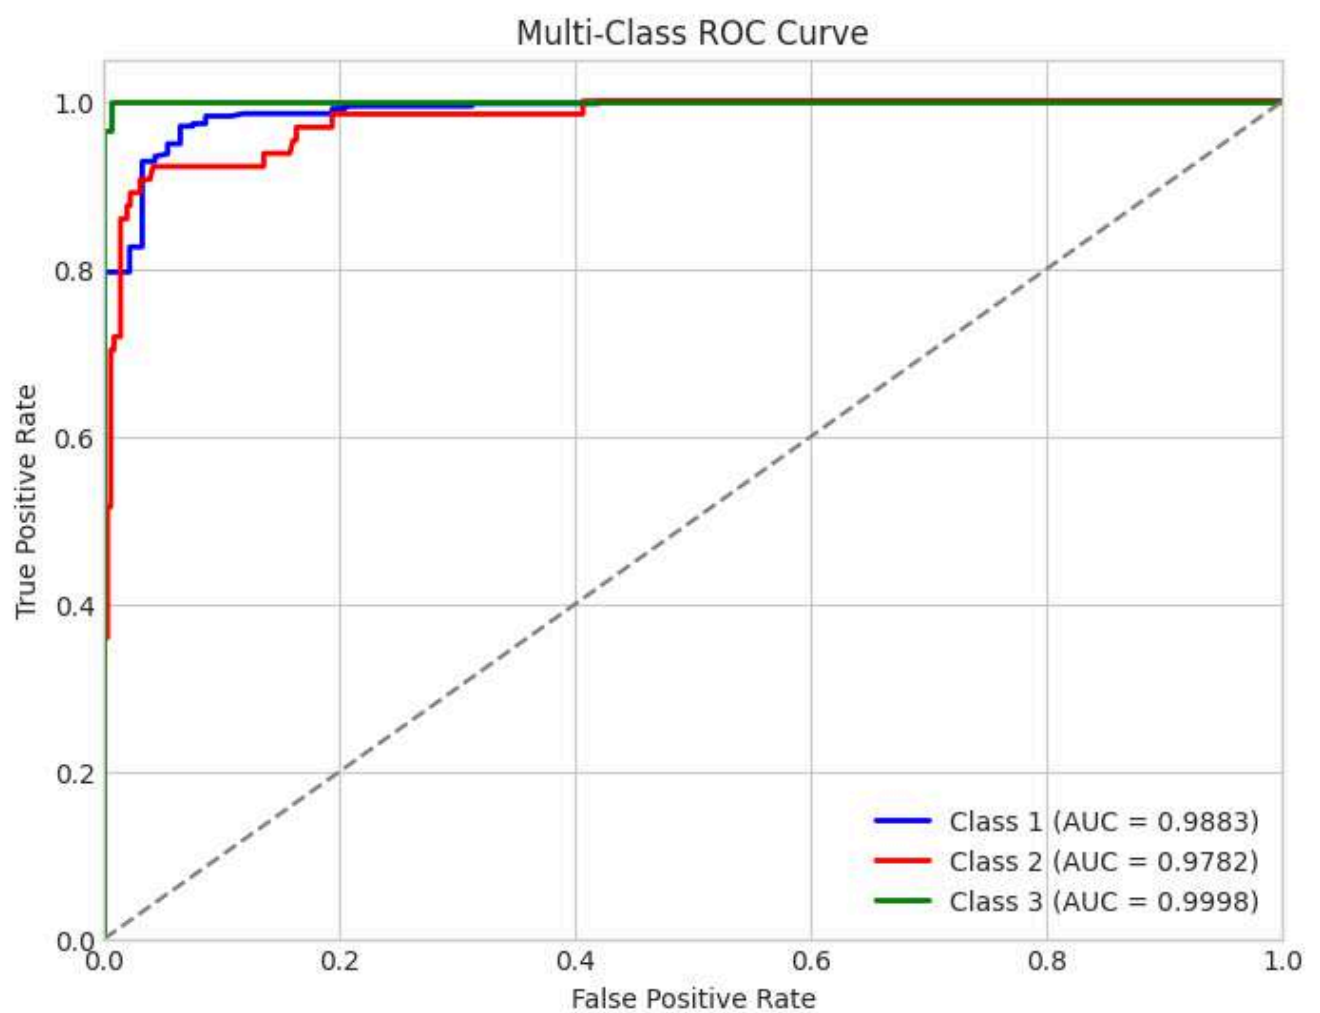
```
[[325  7  1]
 [ 8 55  1]
 [ 0  0 29]]
```

Matthews Correlation Coefficient (MCC): 0.8899

```
In [26]: plot_confusion_matrix_heatmap(y_test, xgb_y_pred, 'XGBoost')
```



```
In [27]: plot_multi_class_roc(xgb_model ,X_test, y_test, num_classes=3)
```

• Blending (Voting Classifier)

```
In [28]: # Blending (Voting Classifier)
voting_model = VotingClassifier(
    estimators=[('xgb', xgb_model), ('rf', rf_model)],
    voting='soft' # soft voting averages the predicted probabilities
)
voting_model.fit(X_train, y_train)

# Predictions
voting_y_pred = voting_model.predict(X_test)

# Accuracy
voting_accuracy = accuracy_score(y_test, voting_y_pred)

# Classification Report & Confusion Matrix for Blending
print("Blending (Voting) Classification Report")
print(classification_report(y_test, voting_y_pred))
print("Blending (Voting) Confusion Matrix")
print(confusion_matrix(y_test, voting_y_pred))

# Accuracy of Blending
print(f"Blending Accuracy: {voting_accuracy:.4f}")

# Calculate MCC for the model's predictions
mcc_voting = matthews_corrcoef(y_test, voting_y_pred)

# Print MCC
print(f"Matthews Correlation Coefficient (MCC): {mcc_voting:.4f}")
```

```

Blending (Voting) Classification Report
      precision    recall  f1-score   support

    0.0         0.97     0.97     0.97        333
    1.0         0.87     0.86     0.87         64
    2.0         0.93     0.97     0.95         29

 accuracy         0.96         426
 macro avg         0.93     0.93     0.93         426
weighted avg         0.96     0.96     0.96         426

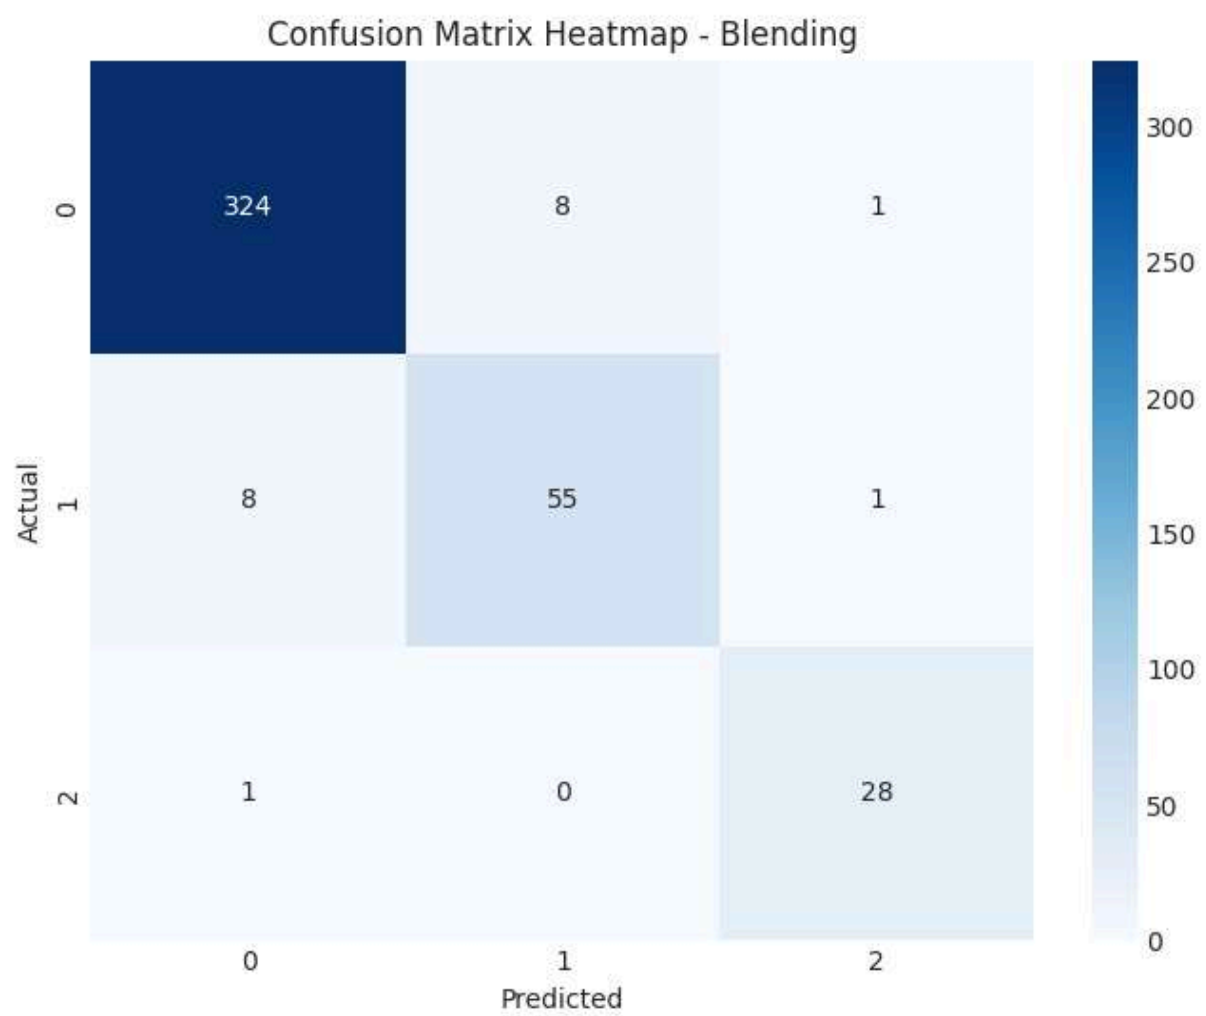
```

```

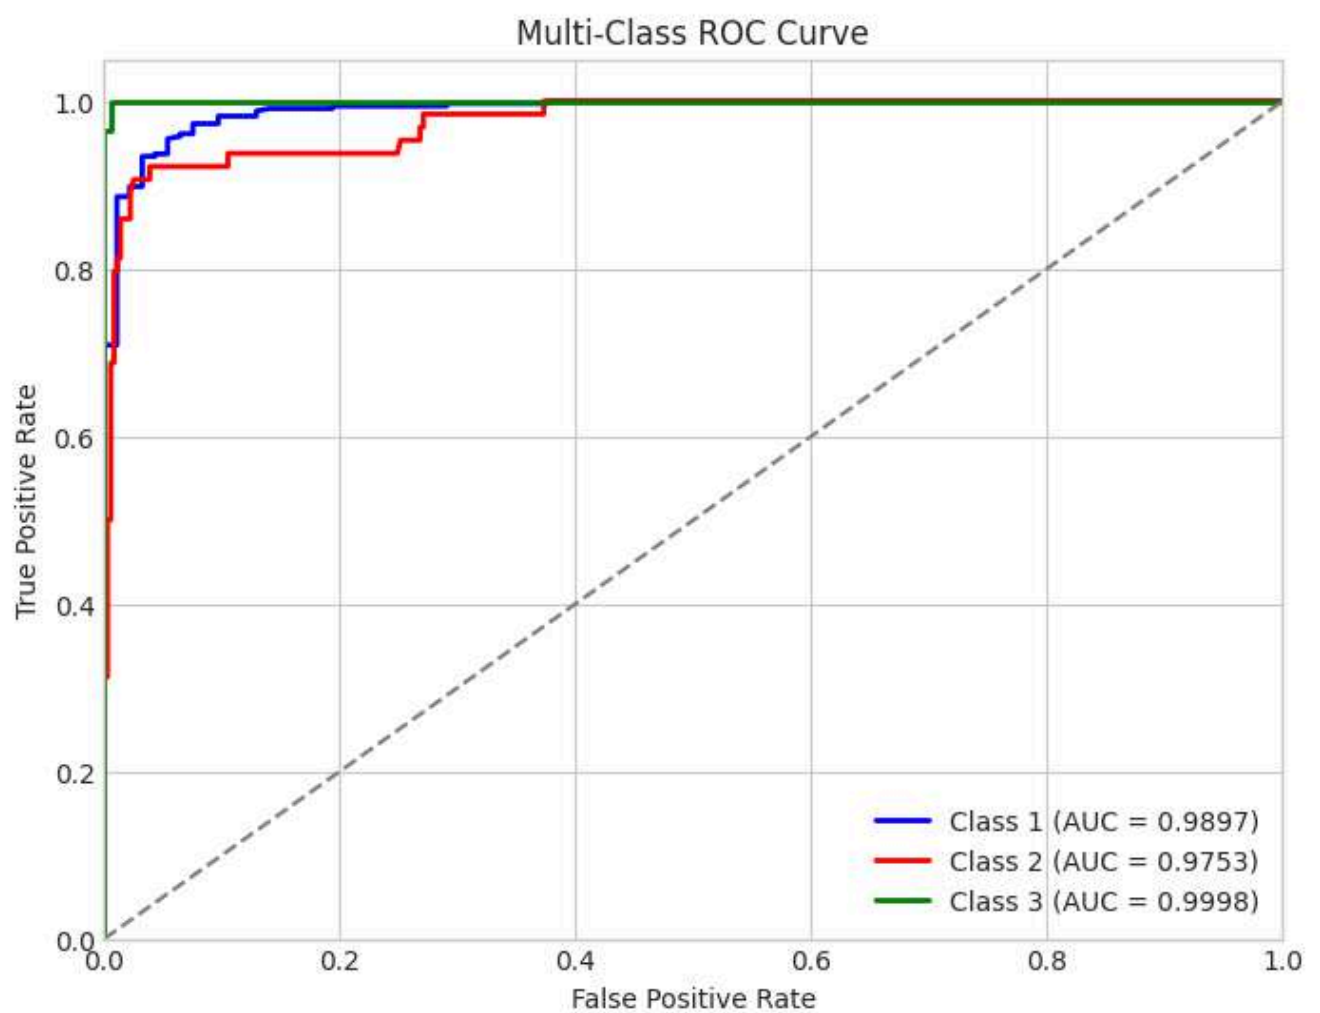
Blending (Voting) Confusion Matrix
[[324  8  1]
 [ 8 55  1]
 [ 1  0 28]]
Blending Accuracy: 0.9554
Matthews Correlation Coefficient (MCC): 0.8768

```

```
In [29]: plot_confusion_matrix_heatmap(y_test, voting_y_pred, 'Blending')
```



```
In [30]: plot_multi_class_roc(voting_model ,X_test, y_test, num_classes=3)
```



● Stacking Classifier with Cross-Validation

```
In [31]: # Stacking Classifier with Cross-Validation
stacking_model = StackingClassifier(
    estimators=[('xgb', xgb_model), ('rf', rf_model)],
    final_estimator=LogisticRegression(), # Meta-model
    cv=5 # Cross-validation folds in Stacking Classifier
)

# Fit Stacking Model
stacking_model.fit(X_train, y_train)

# Cross-Validation Scores
cv_scores = cross_val_score(stacking_model, X_train, y_train, cv=5, scoring='accuracy')

# Predictions on Test Set
stacking_y_pred = stacking_model.predict(X_test)

# Accuracy
stacking_accuracy = accuracy_score(y_test, stacking_y_pred)

# Classification Report & Confusion Matrix for Stacking
print("Stacking Classification Report")
print(classification_report(y_test, stacking_y_pred))
print("Stacking Confusion Matrix")
print(confusion_matrix(y_test, stacking_y_pred))

# Accuracy of Stacking
print(f"Stacking Accuracy: {stacking_accuracy:.4f}")

# Cross-Validation Results
print(f"Cross-Validation Accuracy Scores: {cv_scores}")
print(f"Mean Cross-Validation Accuracy: {cv_scores.mean():.4f}")

# Calculate MCC for the model's predictions
```

```
mcc_stacking = matthews_corrcoef(y_test, stacking_y_pred)
```

```
# Print MCC
```

```
print(f"Matthews Correlation Coefficient (MCC): {mcc_stacking :.4f}")
```

Stacking Classification Report

	precision	recall	f1-score	support
0.0	0.97	0.97	0.97	333
1.0	0.87	0.86	0.87	64
2.0	0.93	0.97	0.95	29
accuracy			0.96	426
macro avg	0.93	0.93	0.93	426
weighted avg	0.96	0.96	0.96	426

Stacking Confusion Matrix

```
[[324  8  1]
 [ 8 55  1]
 [ 1  0 28]]
```

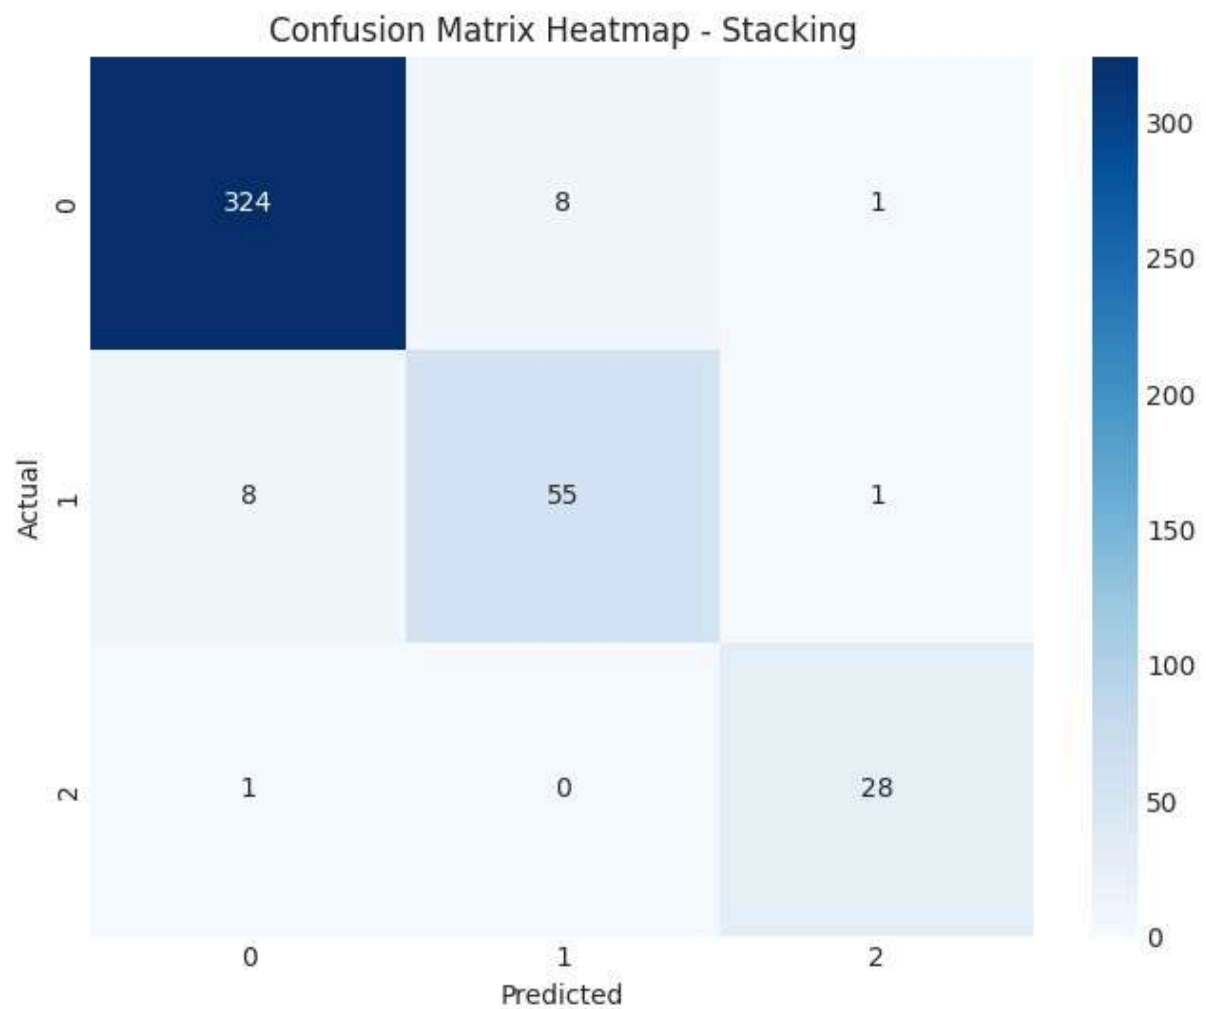
Stacking Accuracy: 0.9554

Cross-Validation Accuracy Scores: [0.94705882 0.93823529 0.93529412 0.95294118 0.95588235]

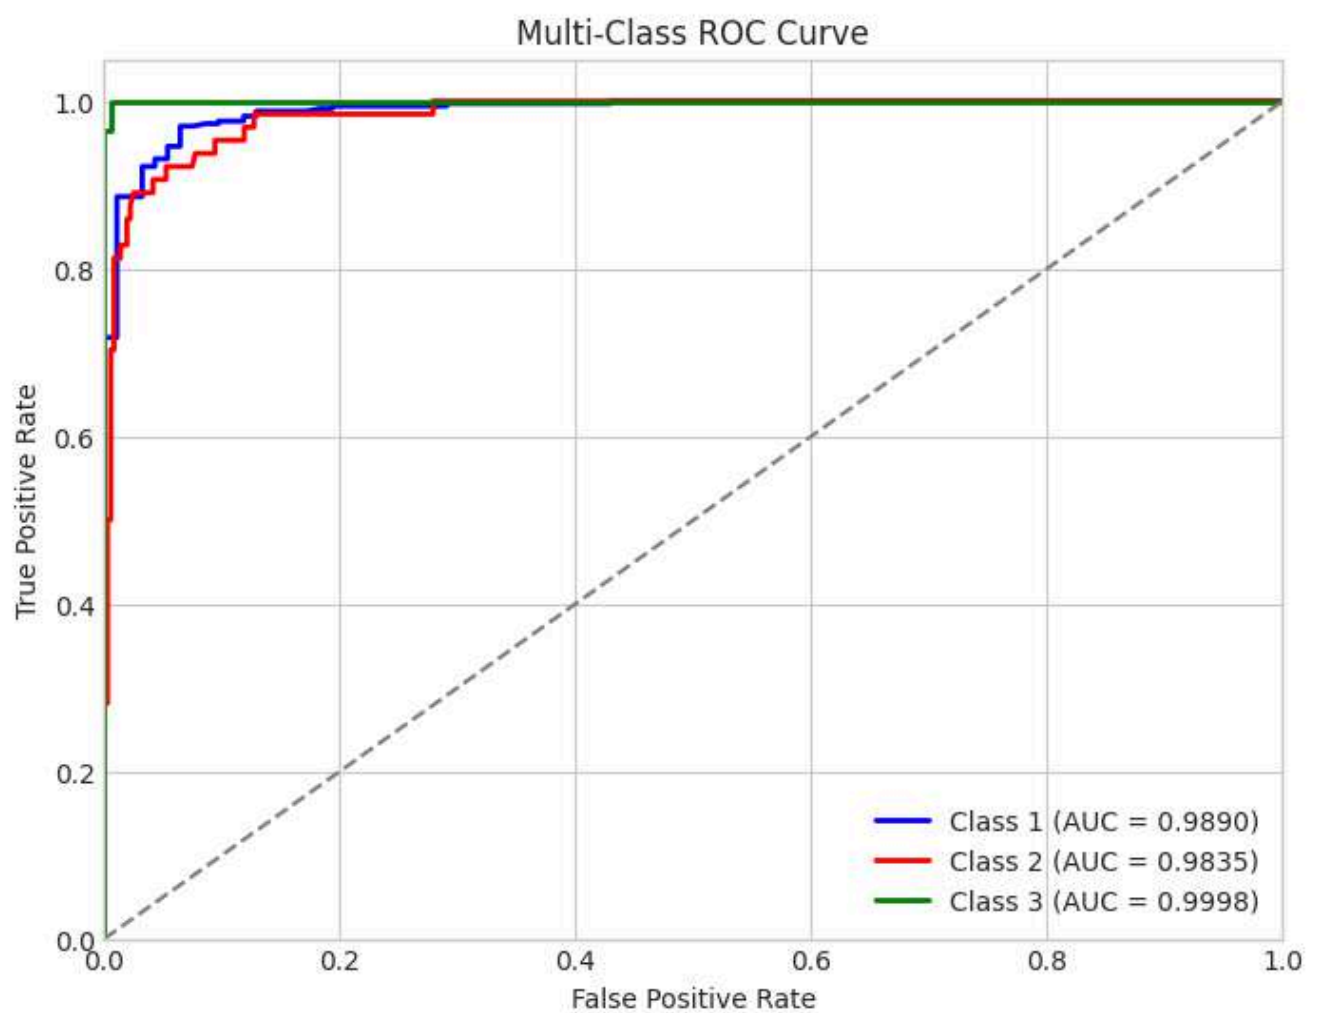
Mean Cross-Validation Accuracy: 0.9459

Matthews Correlation Coefficient (MCC): 0.8768

```
In [32]: plot_confusion_matrix_heatmap(y_test, stacking_y_pred, 'Stacking')
```



```
In [33]: plot_multi_class_roc(stacking_model ,X_test, y_test, num_classes=3)
```



• Model Comparison - Accuracy

```
In [34]: import plotly.graph_objs as go

# Data for the bar plot
models = ['Logistic Regression', 'SVM', 'Random Forest', 'KNN', 'XGBoost', 'Stacking', 'Blending']
accuracies = [log_accuracy, svm_accuracy, rf_accuracy, knn_accuracy, xgb_accuracy, stacking_accuracy, blending_accuracy]

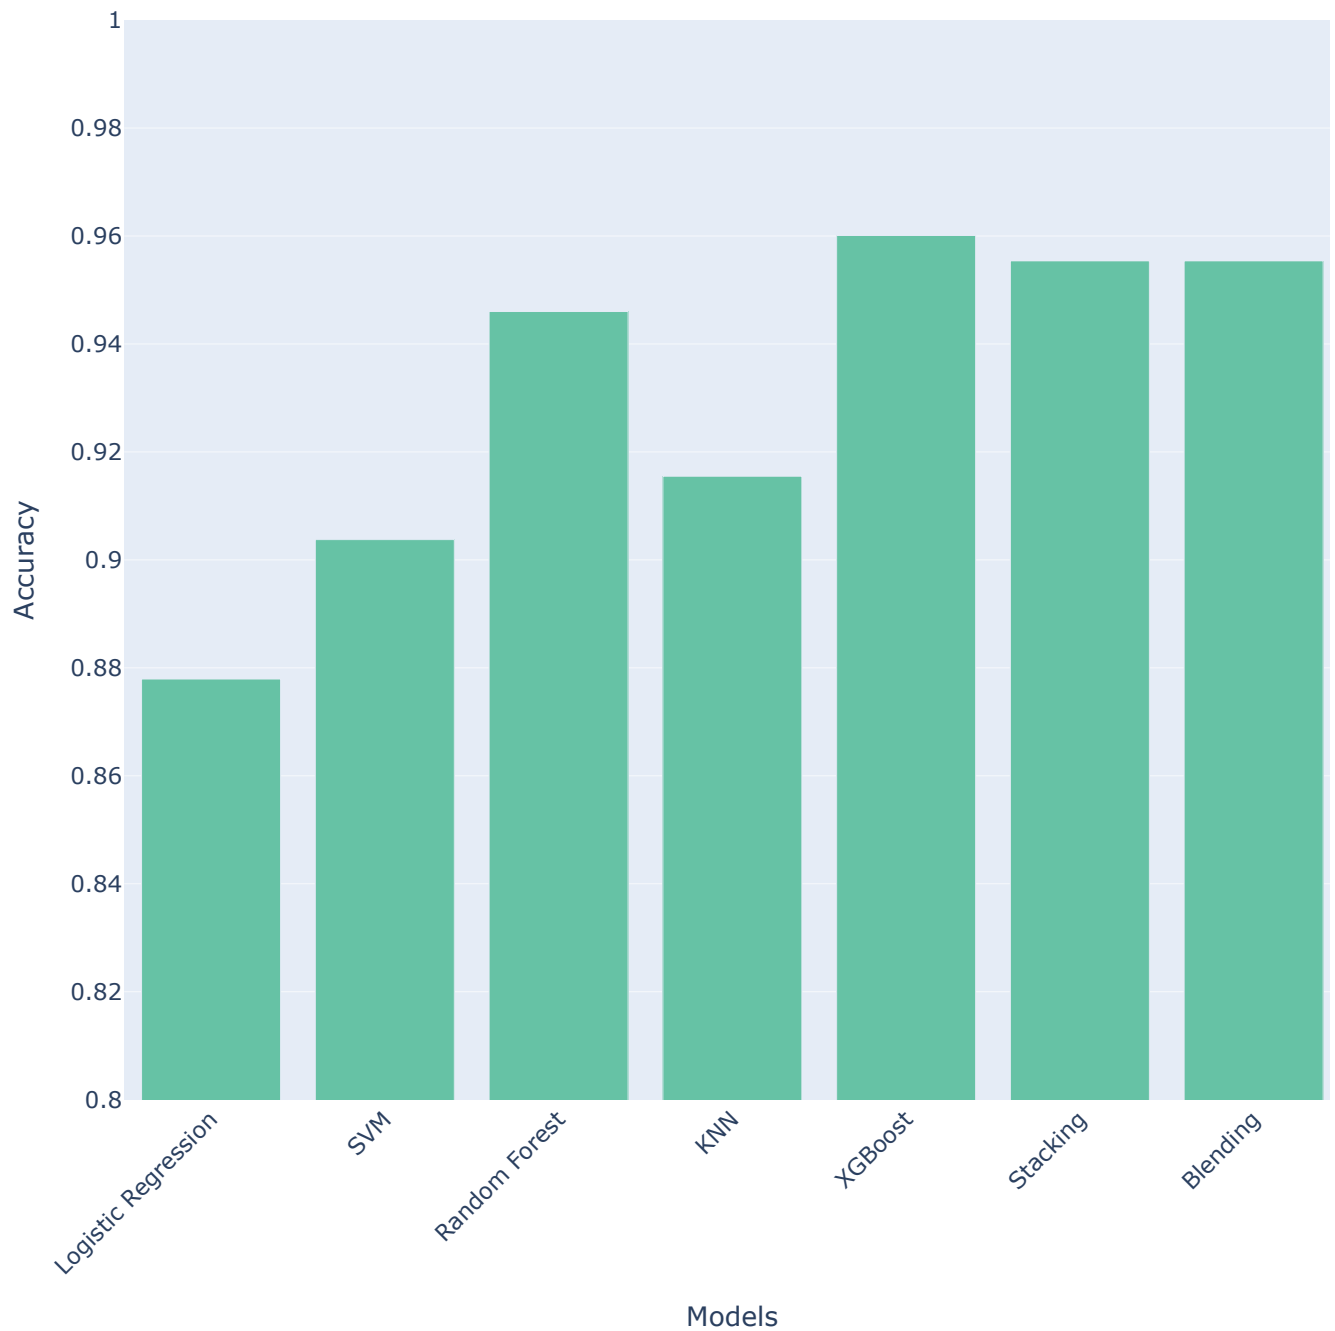
# Create a bar plot
bar_fig = go.Figure()

# Add bar trace
bar_fig.add_trace(go.Bar(
    x=models,
    y=accuracies,
    marker_color='#66C2A5'
))

# Update layout with increased size
bar_fig.update_layout(
    title='Model Comparison - Accuracy',
    xaxis_title='Models',
    yaxis_title='Accuracy',
    yaxis=dict(range=[0.8, 1.0]), # Set y-axis limits
    xaxis_tickangle=-45, # Rotate x-axis labels for better readability
    width=800, # Increase width
    height=800 # Increase height
)

# Show the figure
bar_fig.show()
```


Model Comparison - Accuracy



● Model Comparison - Matthews Correlation Coefficient (MCC)

```
In [35]: # Model names and corresponding MCC values
models = ['Logistic Regression', 'SVM', 'Random Forest', 'KNN', 'XGBoost', 'Stacking', 'Blending']
mcc_values = [mcc_log, mcc_svm, mcc_rf, mcc_knn, mcc_XGb, mcc_stacking, mcc_voting]

# Create a bar plot
mcc_fig = go.Figure()

# Add bar trace
mcc_fig.add_trace(go.Bar(
    x=models,
    y=mcc_values,
    marker_color='#66C2A5'
))

# Update layout with increased size
```

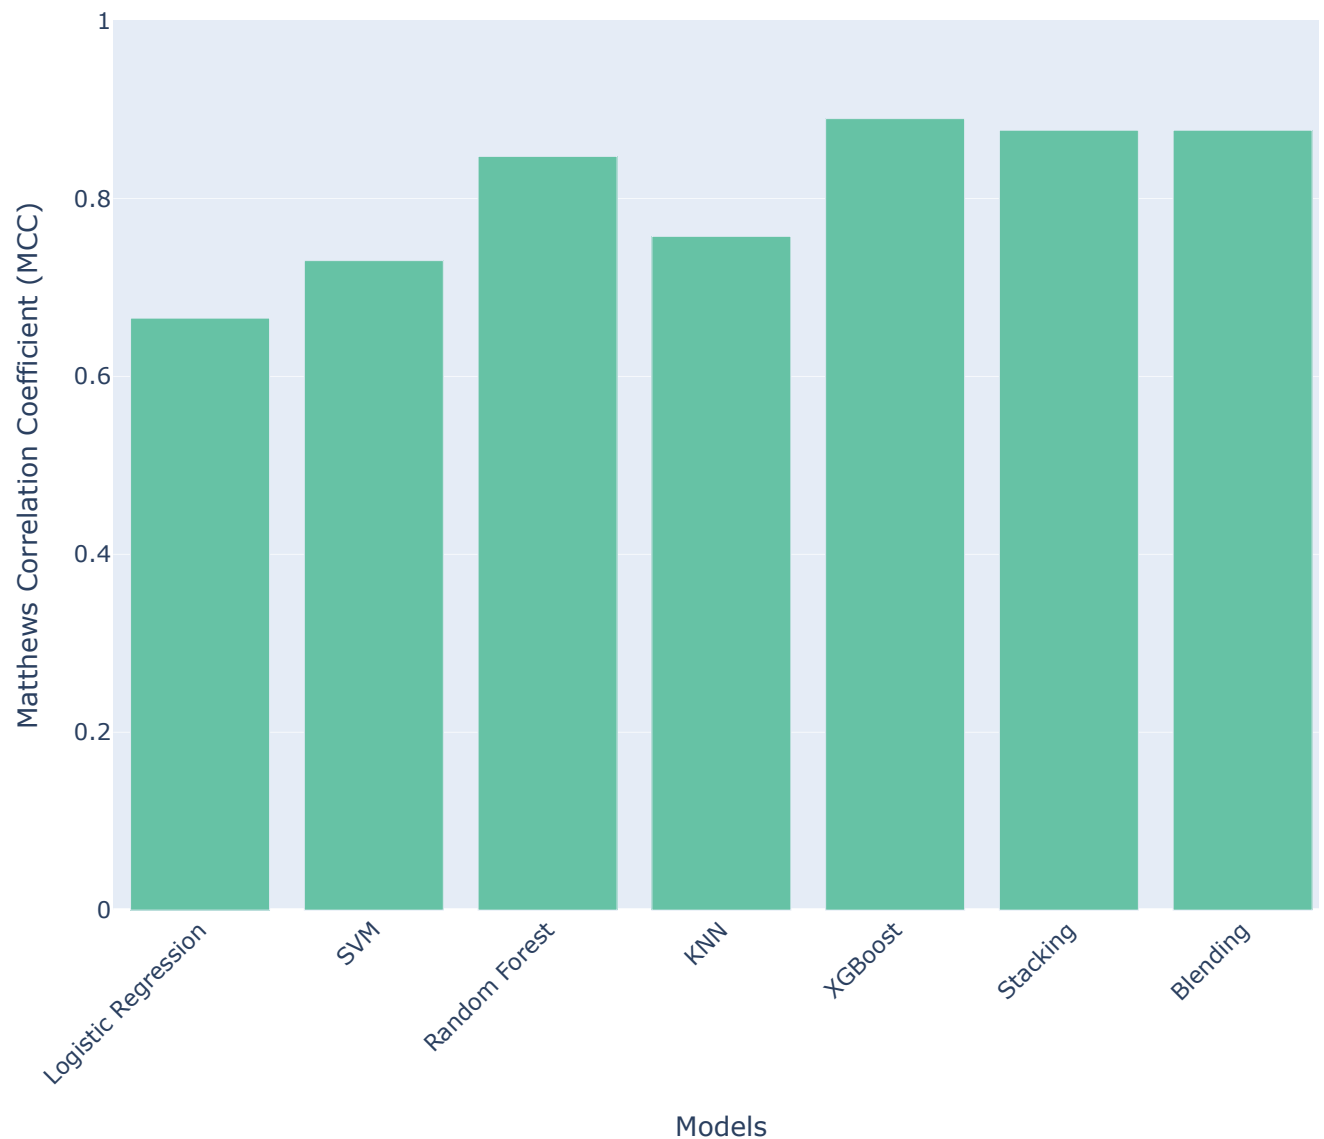
```

mcc_fig.update_layout(
    title='Model Comparison - MCC',
    xaxis_title='Models',
    yaxis_title='Matthews Correlation Coefficient (MCC)',
    yaxis=dict(range=[0.0, 1.0]), # Set y-axis limits
    xaxis_tickangle=-45, # Rotate x-axis labels for better readability
    width=800, # Increase width
    height=700 # Increase height
)

# Show the figure
mcc_fig.show()

```

Model Comparison - MCC



```

In [36]: d={'Models':models,'Accuracies':accuracies,'MCC':mcc_values}
Models_comparison=pd.DataFrame(d)

from IPython.display import display, HTML
# Sort the DataFrame by 'Accuracies' and reset the index
Models_comparison_sorted = Models_comparison.sort_values(by='Accuracies', ascending=False).reset_index
# Define the HTML table string with inline CSS styling
html_table = """
<table style="width:60%; margin:auto; border-collapse:collapse; font-family:Georgia, serif; font-si
<thead>
<tr style="background-color:#4A90E2; color:white; border:2px solid #66C2A5;">
<th style="padding:12px;">Models</th>
<th style="padding:12px;">Accuracies</th>

```

```

        <th style="padding:12px;">MCC</th>
    </tr>
</thead>
<tbody>
"""
# Add table rows dynamically from the DataFrame
for _, row in Models_comparison_sorted.iterrows():
    html_table += f"""
    <tr style="border-bottom:1px solid #ddd;">
        <td style="padding:12px; background-color:#E6F9E6;">{row['Models']}

```

Models	Accuracies	MCC
XGBoost	0.960094	0.889860
Stacking	0.955399	0.876789
Blending	0.955399	0.876789
Random Forest	0.946009	0.847387
KNN	0.915493	0.757509
SVM	0.903756	0.730215
Logistic Regression	0.877934	0.665549

Conclusion

After training and evaluating various machine learning models for fetal health classification, the **XGBoost** model demonstrated superior performance. With an accuracy of **96.01%** and an MCC (Matthews Correlation Coefficient) of **0.8899**, XGBoost stands out as the most reliable option for predicting fetal health status based on CTG data.

Ensemble methods such as **Stacking** and **Blending** also performed well, achieving accuracies of **95.54%** with an MCC of **0.8768**. However, they fell slightly behind XGBoost. The **Random Forest** model achieved an accuracy of **94.60%** and MCC of **0.8474**, positioning it as a strong alternative.

While **K-Nearest Neighbors (KNN)** and **Support Vector Machine (SVM)** provided acceptable results with accuracies of **91.55%** and **90.38%**, their performance was

significantly lower compared to the ensemble models. Lastly, **Logistic Regression** yielded the lowest performance with an accuracy of **87.79%** and MCC of **0.6655**.

In conclusion, the **XGBoost** model is the recommended choice for fetal health classification due to its high accuracy and balanced performance. Its use could assist in early diagnosis, helping healthcare professionals intervene before complications arise, ultimately contributing to reducing child and maternal mortality.

Thank You!

We truly appreciate your time and effort in reviewing this analysis. Your valuable feedback and suggestions are always welcomed! Thank you for being a part of this journey toward learning and growth. Together, we can achieve great things in data science!

Feel free to reach out with your comments and insights. Let's continue learning and building together!

Best regards,
Faiz Siddiqui